ExactFuse: Closed-Form Unfolding with Exact Operators for Sensor-Independent Hyperspectral and Multispectral Image Fusion

Anonymous CVPR submission

Abstract

Deep unfolding networks for hyperspectral and multispectral image fusion alternate data steps with learned priors, but their blur, downsampling and spectral-response operators are usually learned or zero-padded, so they differ from the operators that produced the data. We study what exact operators bring to such a network. With a separable blur and the true border rule, the data step has a closed form in a product eigenbasis, and the null space of fusion, the part of the image that no observation constrains, is explicit. An exact step keeps the ground truth fixed, an inexact one moves it by an amount that we bound, and blur, spectral response and sub-pixel grid offsets can be estimated from the two observations. Our network, ExactFuse, takes this closed-form step with learned spectral gains, gives its prior inputs in the hyperspectral bands only, so that one network can serve any multispectral sensor whose response is known, and is trained over a family of observation models whose operators it estimates at test time. Against twelve published networks retrained from their authors' code under one protocol and budget, ExactFuse has the highest PSNR on Chikusei and Pavia Centre, by 0.79 and 0.55 dB, and in 7 of 7 further settings up to $\times 16$, and it reproduces the low-resolution input 17 dB more accurately than the strongest of them. A prior without data steps already beats every baseline, so much of the margin comes from the prior and its training; the exact operators make the result consistent with its inputs, replaceable at test time and robust to changed blurs, responses and noise.

1. Introduction

Hyperspectral cameras record hundreds of narrow bands at a coarse spatial resolution, and multispectral cameras record a few broad bands at a fine one. Fusing the two images gives a high-resolution hyperspectral image (HR-HSI) [29]. Most recent methods are deep networks trained on simulated pairs: convolutional networks [41], transformers [6, 14, 17], state-space models [19] and diffusion models [5].

Deep unfolding networks [8, 12, 31, 32] unroll an iterative solver of the inverse problem: each stage takes a data step, which reduces the mismatch with the two observations, and applies a learned prior. The data step needs the blur, downsampling and spectral response, which most unfolding networks learn as convolutions or fix with zero padding, although the benchmark data are simulated with a known blur and another border rule; even the ground truth then fails the network's own consistency check. Each network also serves a single sensor pair.

Evaluation is a second problem: splits, responses, normalizations and metric code differ between papers [29], so PSRT [6], for example, is reported at 44.60 dB PSNR on Chikusei $\times 4$ in [19] and at 50.67 dB in [5], and copied numbers cannot rank methods.

This paper asks what exact operators bring to an unfolded fusion network. We write fusion as maximum a posteriori (MAP) estimation with parametric operators, unroll half-quadratic splitting (HQS) [39] and study the resulting data step. Our contributions are:

- Our central contribution is a theory of exact operators in unfolded fusion that links observation noise to reconstruction error. With a separable blur and the true border rule, the data step has a closed form and the fusion null space is explicit (88% of the unknowns on Chikusei). The operators are identifiable from the observation pair over continuous parameters, with an exact condition for when two sensor configurations cannot be told apart and a stability bound under noise, and an operator error moves the reconstruction by an amount bounded independently of the learned weights (Sec. 3). None of this applies to networks with learned or zero-padded operators.
- ExactFuse, an unfolded network with exact and replaceable operators, a learned spectral data step that keeps these guarantees, inputs in the hyperspectral bands, so that one set of weights serves any multispectral sensor whose response is known, and training over a family of observation models with the operators estimated at test time (Sec. 4). Its building blocks are standard and not claimed.
- An equal-budget benchmark of twelve published networks, retrained from their authors' code, and a null-space split of the error that shows where each method fails (Sec. 5). It also shows where ExactFuse's margin comes from: much of it from the prior and its training, which

keep a lead under equal recipes, while the exact operators make the result consistent with its inputs and matter most when the operator changes.

2. Related work

Deep fusion networks map the two observations directly to the HR-HSI with convolutions [41], transformers [6, 10, 14, 17, 33], general fusion backbones [20], state-space models [19] or diffusion [5], without using the observation model. SSA [16] slices its convolution kernels to the number of bands and reconstructs at arbitrary scales, so one model serves several sensors; it does not model their responses.

Unfolding and operator knowledge. Algorithm unrolling [18] turns solver iterations into network stages. MHF-Net [31], MoG-DCN [8] and DHIF-Net [12] unfold fusion solvers with learned or approximate operators, later networks strengthen the prior [32] or the proximal update [36], and IR&ArF [23] unfolds the fusion of unregistered images with a learned registration. USRNet [39] gives the blur kernel of single-image super-resolution to a closed-form data step under periodic borders, and blind methods train on randomized degradations [40]; we carry this to fusion, with two operators and the true border rule.

Closed forms and separable operators. FUSE [30] solves the quadratic fusion problem with a Sylvester equation under periodic borders, and tensor methods [13, 15, 22] use separable degradations to prove recoverability of low-rank models. Our closed form holds for the reflect border of the simulation and is a differentiable layer whose operators change per image; its learned filters resemble deep Wiener deconvolution [7], bounded so that the step stays nonexpansive.

Range and null space. A split of the reconstruction into a part fixed by the measurements and a null-space part underlies range–null-space networks [4], back-projection improves the consistency of pansharpening [27], and an adjoint that does not match the data biases tomographic reconstruction [38]. With exact separable operators the projection has a closed form (Theorem 1), which we use to split every method's error.

3. Exact operators in unfolded fusion

3.1. Observation model and unfolding

Let $\mathbf{X} \in \mathbb{R}^{B\times H\times W}$ be the HR-HSI. The two observations are

$$\mathbf{Y}_{\mathrm{H}} = \mathcal{D}_\theta\mathbf{X} + \mathbf{N}_{\mathrm{H}}, \qquad \mathbf{Y}_{\mathrm{M}} = \mathcal{R}_\phi\mathbf{X} + \mathbf{N}_{\mathrm{M}}, \tag{1}$$

where $\mathcal{D}_\theta$ applies to every band reflect padding, a Gaussian blur of width θ and decimation by s starting at index zero, so $\mathbf{Y}_{\mathrm{H}} \in \mathbb{R}^{B\times h\times w}$ with $h = H/s$, $w = W/s$; $\mathcal{R}_\phi$ applies to every pixel the M spectral responses of the multispectral sensor, shifted along wavelength by ϕ nm; and $\mathbf{N}_{\mathrm{H}}, \mathbf{N}_{\mathrm{M}}$ are white Gaussian noise with standard deviations $\varsigma_{\mathrm{H}}, \varsigma_{\mathrm{M}}$. The benchmarks use $s = 4$, a 7×7 blur with $\theta = \sigma_0 = 2$, $\phi = 0$ and no noise [5]. We implement $\mathcal{D}_\theta$ and its exact adjoint, in which the adjoint of reflect padding adds the padded margins back onto the pixels they were copied from (the supplementary material); the adjoint test $\langle \mathcal{D}\mathbf{X}, \mathbf{Y}\rangle = \langle \mathbf{X}, \mathcal{D}^\top\mathbf{Y}\rangle$ holds to 10^{-13} in double precision.

With a prior $p(\mathbf{X}) \propto \exp(-\lambda\Phi(\mathbf{X}))$, the MAP estimate minimizes $\frac{1}{2\varsigma_{\mathrm{H}}^2}\|\mathcal{D}_\theta\mathbf{X}-\mathbf{Y}_{\mathrm{H}}\|^2 + \frac{1}{2\varsigma_{\mathrm{M}}^2}\|\mathcal{R}_\phi\mathbf{X}-\mathbf{Y}_{\mathrm{M}}\|^2 + \lambda\Phi(\mathbf{X})$. HQS couples $\mathbf{X}$ to an auxiliary image $\mathbf{Z}$ with a penalty $\frac{\mu_k}{2}\|\mathbf{Z}-\mathbf{X}\|^2$ and alternates a data step and a prior step, $\mathbf{X}_k = \mathrm{prox}_{(\lambda/\mu_k)\Phi}(\mathbf{Z}_k)$. With $\alpha_k = 1/(\mu_k\varsigma_{\mathrm{H}}^2)$ and $\beta_k = 1/(\mu_k\varsigma_{\mathrm{M}}^2)$ the data step is the linear system

$$(\mathbf{I}+\alpha_k\mathcal{D}_\theta^\top\mathcal{D}_\theta+\beta_k\mathcal{R}_\phi^\top\mathcal{R}_\phi)\mathbf{Z}_k = \mathbf{X}_{k-1}+\alpha_k\mathcal{D}_\theta^\top\mathbf{Y}_{\mathrm{H}}+\beta_k\mathcal{R}_\phi^\top\mathbf{Y}_{\mathrm{M}}, \tag{2}$$

and one gradient step from $\mathbf{X}_{k-1}$ is its linearization,

$$\mathbf{Z}_k \approx \mathbf{X}_{k-1}-\alpha_k\mathcal{D}_\theta^\top(\mathcal{D}_\theta\mathbf{X}_{k-1}-\mathbf{Y}_{\mathrm{H}})-\beta_k\mathcal{R}_\phi^\top(\mathcal{R}_\phi\mathbf{X}_{k-1}-\mathbf{Y}_{\mathrm{M}}). \tag{3}$$

An unfolded network replaces the proximal map by a learned module. The derivation already says two things about such a network: the ratio $\beta_k/\alpha_k = \varsigma_{\mathrm{H}}^2/\varsigma_{\mathrm{M}}^2$ should follow the noise of the two observations, and the operators enter only through $\mathcal{D}_\theta$, $\mathcal{R}_\phi$ and their adjoints, so they can be replaced at test time without retraining.

3.2. Separable structure and the fusion null space

A Gaussian blur aligned with the image axes is separable, and padding and decimation act on rows and columns independently. The operators therefore factor exactly, border included:

$$\mathcal{D}_\theta\mathbf{X} = \mathbf{X}\times_2\mathbf{D}_{\mathrm{r}}\times_3\mathbf{D}_{\mathrm{c}}, \qquad \mathcal{R}_\phi\mathbf{X} = \mathbf{X}\times_1\mathbf{r}_\phi, \tag{4}$$

where $\times_n$ is the mode-n product, $\mathbf{D}_{\mathrm{r}} \in \mathbb{R}^{h\times H}$ and $\mathbf{D}_{\mathrm{c}} \in \mathbb{R}^{w\times W}$ pad, blur and decimate one axis, and $\mathbf{r}_\phi \in \mathbb{R}^{M\times B}$ holds the spectral responses.

Theorem 1 (Closed-form data step and null space)**.** *Let* $\mathbf{D}_{\mathrm{r}}^\top\mathbf{D}_{\mathrm{r}} = \mathbf{U}_{\mathrm{r}}\mathrm{diag}(\boldsymbol{\lambda}^{\mathrm{r}})\mathbf{U}_{\mathrm{r}}^\top$, $\mathbf{D}_{\mathrm{c}}^\top\mathbf{D}_{\mathrm{c}} = \mathbf{U}_{\mathrm{c}}\mathrm{diag}(\boldsymbol{\lambda}^{\mathrm{c}})\mathbf{U}_{\mathrm{c}}^\top$ *and* $\mathbf{r}_\phi^\top\mathbf{r}_\phi = \mathbf{V}\mathrm{diag}(\boldsymbol{\nu})\mathbf{V}^\top$*, and let* $\mathcal{Q}(\cdot) = (\cdot)\times_1\mathbf{V}\times_2\mathbf{U}_{\mathrm{r}}\times_3\mathbf{U}_{\mathrm{c}}$*. (a)* $\mathcal{D}_\theta^\top\mathcal{D}_\theta$ *and* $\mathcal{R}_\phi^\top\mathcal{R}_\phi$ *are diagonal in the orthonormal basis* $\mathcal{Q}$*, with eigenvalues* $\lambda_i^{\mathrm{r}}\lambda_j^{\mathrm{c}}$ *and* ν_b*. (b) The solution of* (2) *is*

$$\mathbf{Z}_k = \mathcal{Q}\Big(\mathcal{Q}^{-1}(\mathbf{C}_k) \oslash \big[1+\alpha_k\lambda_i^{\mathrm{r}}\lambda_j^{\mathrm{c}}+\beta_k\nu_b\big]_{bij}\Big), \tag{5}$$

with $\mathbf{C}_k$ *the right-hand side of* (2) *and* $\oslash$ *element-wise division. (c) The joint null space* $\mathcal{N}(\mathcal{D}_\theta)\cap\mathcal{N}(\mathcal{R}_\phi)$ *is spanned by the basis tensors with* $\nu_b = 0$ *and* $\lambda_i^{\mathrm{r}}\lambda_j^{\mathrm{c}} = 0$*. If* $\mathbf{r}_\phi$*,* $\mathbf{D}_{\mathrm{r}}$ *and* $\mathbf{D}_{\mathrm{c}}$ *have full row rank, its dimension is* $(B-M)(HW-hw)$*, a fraction* $(1-M/B)(1-1/s^2)$ *of the unknowns.*

The proof (the supplementary material) uses the mixed-product rule: every factor of $\mathcal{Q}$ diagonalizes one mode. The step (5) costs three mode products in each direction, $O(BHW(B+H+W))$ operations, and needs no iteration. It is exact for any separable kernel, not only a Gaussian, with any border rule and sub-pixel offset per axis, and for any linear spectral response, given eigendecompositions of an $H \times H$, a $W \times W$ and a $B \times B$ matrix, cached per image size; FUSE [30] needs periodic borders. A non-separable or spatially varying kernel breaks the Kronecker structure, and the network then takes a conjugate-gradient step with the same operator. Part (c) states what any method must infer from a prior: the high spatial frequencies of the spectral directions the multispectral sensor does not see, 88% of the unknowns on Chikusei and 90% on Pavia Centre at $\times 4$ (all high frequencies, $1 - 1/s^2$ of the unknowns, without the HR-MSI). Because spatial detail is shared across bands, the M observed directions carry 84% of its energy on Chikusei and 77% on Pavia Centre: the minimum-norm image consistent with both observations reaches 41.3 dB on Chikusei, against 33.5 dB from the LR-HSI alone. The split of every method's error in Sec. 5.3 rests on this null space.

3.3. The cost of an inexact operator

Stack the data terms as $\mathbf{A} = [\mathcal{D}_\theta; \mathcal{R}_\phi]$, $\mathbf{Y} = [\mathbf{Y}_\mathrm{H}; \mathbf{Y}_\mathrm{M}]$ and $\mathbf{W}_k = \mathrm{diag}(\alpha_k \mathbf{I}, \beta_k \mathbf{I})$, and let $\tilde{\mathbf{A}}$ be the operator a network actually uses: one with a zero border, a learned one, or one with wrong parameters. Write $T(\mathbf{V}) = (\mathbf{I} + \tilde{\mathbf{A}}^\top \mathbf{W}_k \tilde{\mathbf{A}})^{-1}(\mathbf{V} + \tilde{\mathbf{A}}^\top \mathbf{W}_k \mathbf{Y})$ for the exact data step and $G(\mathbf{V}) = \mathbf{V} - \tilde{\mathbf{A}}^\top \mathbf{W}_k(\tilde{\mathbf{A}}\mathbf{V} - \mathbf{Y})$ for its linearization.

Theorem 2 (Fixed point and operator mismatch)**.** *Let* $\mathbf{Y} = \mathbf{A}\mathbf{X}^*$. *(a) If* $\tilde{\mathbf{A}} = \mathbf{A}$, *then* $T(\mathbf{X}^*) = G(\mathbf{X}^*) = \mathbf{X}^*$ *for all* $\alpha_k, \beta_k > 0$. *(b) For any* $\tilde{\mathbf{A}}$, *with* $\mathbf{e} = \mathbf{W}_k^{1/2}(\mathbf{A} - \tilde{\mathbf{A}})\mathbf{X}^*$,

$$\|T(\mathbf{X}^*) - \mathbf{X}^*\| \le \tfrac{1}{2}\|\mathbf{e}\|, \qquad \|G(\mathbf{X}^*) - \mathbf{X}^*\| \le \|\mathbf{W}_k^{1/2}\tilde{\mathbf{A}}\|\,\|\mathbf{e}\|. \tag{6}$$

(c) If stage k *computes* $\mathbf{X}_k = (\mathcal{I} + \mathcal{P}_k)(T_k(\mathbf{X}_{k-1}))$ *with a residual prior that satisfies* $\mathcal{P}_k(\mathbf{X}^*) = \mathbf{0}$ *and makes* $\mathcal{I} + \mathcal{P}_k$ L_k*-Lipschitz, then a* K*-stage network started at* $\mathbf{X}^*$ *returns* $\|\mathbf{X}_K - \mathbf{X}^*\| \le \frac{1}{2}\sum_{k=1}^{K}\big(\prod_{j=k}^{K} L_j\big)\|\mathbf{e}_k\|$, *where* $\mathbf{e}_k$ *is* $\mathbf{e}$ *for the weights* $\mathbf{W}_k$ *of stage* k.

The bounds follow from the singular values $b_i/(1+b_i^2) \le \frac{1}{2}$ of $(\mathbf{I}+\mathbf{B}^\top\mathbf{B})^{-1}\mathbf{B}^\top$, $\mathbf{B} = \mathbf{W}_k^{1/2}\tilde{\mathbf{A}}$, and from the nonexpansiveness of T_k (the supplementary material). With the exact operator the ground truth is a fixed point of every data step; with an inexact one every data step moves it, by an image-dependent amount that no fixed correction removes. The exact step bounds this displacement whatever the learned weights are, the gradient step does not, and giving the network the true operator of other data restores (a), which a learned operator cannot do.

Corollary 1 (Border rule)**.** *If* $\tilde{\mathbf{A}}$ *differs from* $\mathbf{A}$ *only in the border rule of* $\mathcal{D}_\theta$, *then* $(\mathbf{A} - \tilde{\mathbf{A}})\mathbf{X}^*$ *is nonzero only on the low-resolution pixels whose kernel support reaches the padding, i.e., within* $\lceil p/s \rceil$ *pixels of the border for a* $(2p+1) \times (2p+1)$ *kernel.*

For our benchmarks only the first low-resolution row and column are affected: 3% of a 64×64 test image, but 12% of the 16×16 patches on which networks are trained.

3.4. Estimating the operators

Theorem 3 (Identifiability and stability)**.** *Let* $\mathcal{D}_\theta$ *act identically on every band and* $\mathcal{R}_\phi$ *identically on every pixel, with kernels and responses that sum to one, and let* $\mathbf{r}(\theta, \phi) = \mathcal{D}_\theta \mathbf{Y}_\mathrm{M} - \mathcal{R}_\phi \mathbf{Y}_\mathrm{H}$. *(a)* $\mathcal{D}_\theta \mathcal{R}_{\phi'} = \mathcal{R}_{\phi'} \mathcal{D}_\theta$, *so for noise-free data generated with* (θ^*, ϕ^*) *the criterion*

$$J(\theta, \phi) = \|\mathbf{r}(\theta, \phi)\|^2 = \|\mathcal{D}_\theta \mathbf{Y}_\mathrm{M} - \mathcal{R}_\phi \mathbf{Y}_\mathrm{H}\|^2 \tag{7}$$

vanishes at (θ^*, ϕ^*). *(b)* $J(\theta, \phi) = 0$ *if and only if* $\mathbf{X}$ *lies in the null space of* $\mathcal{C}_{\theta\phi} = \mathcal{D}_\theta \mathcal{R}_{\phi^*} - \mathcal{R}_\phi \mathcal{D}_{\theta^*}$, *and* $\mathcal{C}_{\theta\phi} \neq 0$ *whenever* $(\mathcal{D}_\theta, \mathcal{R}_\phi) \neq (\mathcal{D}_{\theta^*}, \mathcal{R}_{\phi^*})$*: two configurations are confused only on a proper subspace of images, such as spatially constant images (blur and offsets) or images with equal bands (response shift). (c) If the Jacobian* $\mathbf{G}$ *of* $\mathbf{r}$ *with respect to* (θ, ϕ) *has full column rank at* (θ^*, ϕ^*), *this point is an isolated zero of* J *over the continuous parameters, and with observation noise the minimizer of* J *moves, to first order, by* $-\mathbf{G}^+\mathbf{n}$ *with* $\mathbf{n} = \mathcal{D}_{\theta^*}\mathbf{N}_\mathrm{M} - \mathcal{R}_{\phi^*}\mathbf{N}_\mathrm{H}$, *hence by at most* $\|\mathbf{n}\|/\sigma_{\min}(\mathbf{G})$.

Corollary 2 (From noise to reconstruction)**.** *Under (c), the exact data step with the estimated operator moves the ground truth by at most* $\|\mathbf{W}_k^{1/2}\mathbf{A}'\|\,\|\mathbf{n}\|/(2\,\sigma_{\min}(\mathbf{G}))$ *to first order, with* $\mathbf{A}'$ *the Jacobian of* $(\theta, \phi) \mapsto \mathbf{A}_{\theta\phi}\mathbf{X}^*$.

Unlike uniqueness on a finite grid, (b) is exact for continuous families, and (c) makes identifiability quantitative: $\sigma_{\min}(\mathbf{G})$, the conditioning of the estimation problem, is the stability constant, and the corollary carries observation noise through the operator estimate to the reconstruction (proofs in the supplementary material). θ may collect two blur widths and two grid offsets. On the test images $\mathbf{G}$ has full rank for these five parameters (condition number 1.28 with unit-norm columns on Chikusei); at 30 dB SNR, (c) predicts errors of 0.0013 px for the offsets and 0.0017 px for the blur width, against 0.0013 and 0.0042 px measured, the excess being a second-order noise bias (the supplementary material). Part (a) is the relation HySure [25] uses. We minimize J on a grid (steps $0.025\,\sigma_0$ and 1 nm); since $\mathbf{r}$ is linear in the kernel, a blur of any shape can also be estimated by least squares (Sec. 5.5).

4. ExactFuse

Stages. ExactFuse (Fig. 1) unrolls HQS for $K = 3$ stages from the bicubic upsampling $\mathbf{X}_0$ of $\mathbf{Y}_\mathrm{H}$. Stage k first takes

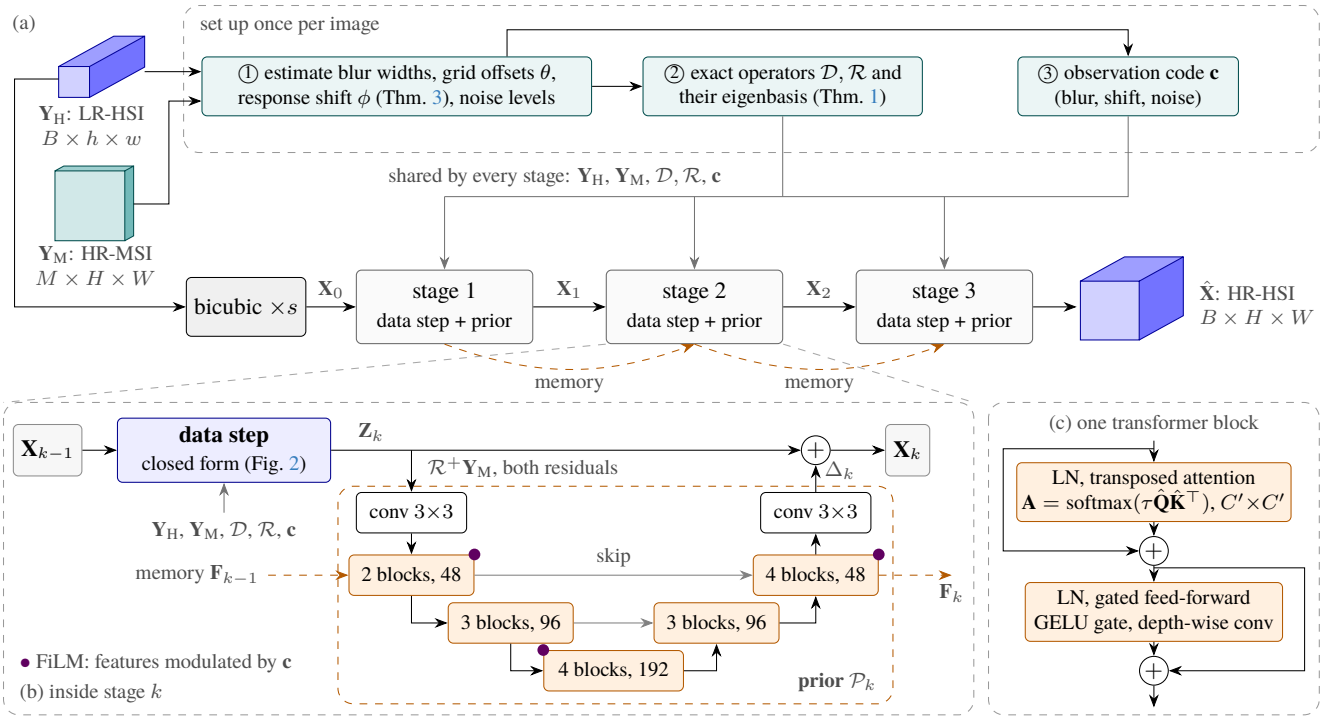


Figure 1. ExactFuse. (a) The HQS iteration of Section 3.1 unrolled for three stages. Once per image, the operator parameters and noise levels are given or estimated from the observation pair (Theorem 3) and the exact operators and their eigenbasis are built (Theorem 1); every stage passes its decoder features to the next (memory). (b) Inside a stage, the data step (Fig. 2) makes the estimate agree with both observations, and the prior adds a residual Δ_k. The prior is a three-level U-Net (transformer blocks and channels per level; strided convolutions halve the resolution, transposed convolutions double it); all its inputs lie in the hyperspectral bands, and its last convolution starts at zero. (c) Every block is a Restormer block [37]: layer normalization (LN), transposed attention between the C' channels of each head, whose cost is linear in the number of pixels, and a gated feed-forward network, each in a residual branch. The supplementary material lists every layer.

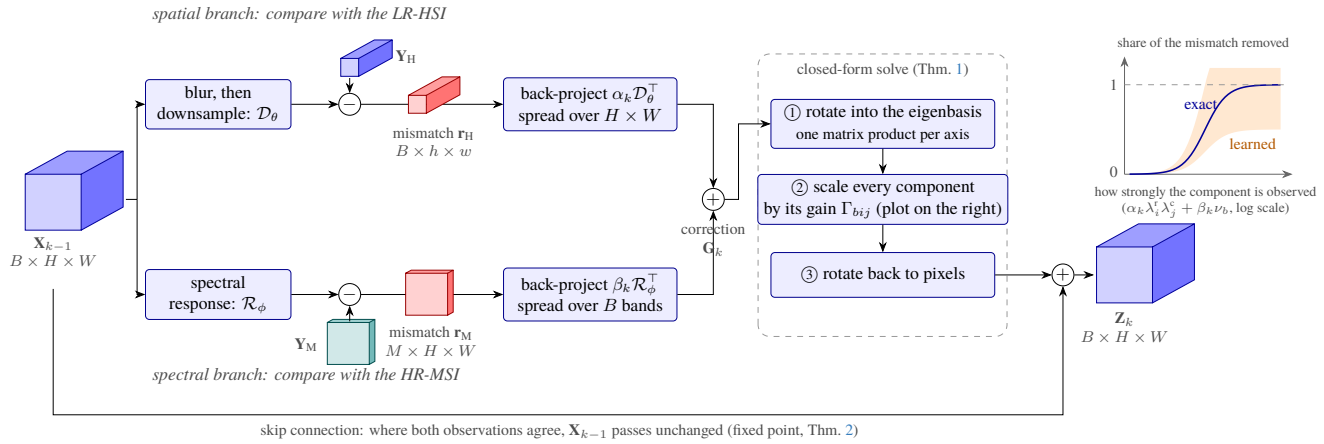


Figure 2. The data step of stage k, the exact solution of (2). Two physics branches compare the estimate with the observations: the spatial branch blurs and downsamples it and compares it with the LR-HSI, the spectral branch applies the multispectral responses and compares it with the HR-MSI; both mismatches are spread back to the full cube by the exact adjoints. The closed-form solve rotates this correction into the eigenbasis of the separable operators, removes in each component a share $\Gamma_{bij}(\alpha_k\lambda_i^{\mathrm{r}}\lambda_j^{\mathrm{c}} + \beta_k\nu_b)$ of the mismatch that grows with how strongly the observations see that component (blue: the exact step (5); orange: the range allowed to the learned gains of Proposition 1), and rotates back. Components that neither observation sees (the null space) are left to the prior.

a data step with the exact operators: the learned spectral step (9) below, or, for non-separable operators and in the base model of Sec. 5.4, the linearized step

$$\mathbf{Z}_k = \mathbf{X}_{k-1} - s^2\alpha_k\mathcal{D}_\theta^\top(\mathcal{D}_\theta\mathbf{X}_{k-1} - \mathbf{Y}_\mathrm{H}) - \beta_k\mathcal{R}_\phi^\top(\mathcal{R}_\phi\mathbf{X}_{k-1} - \mathbf{Y}_\mathrm{M}) \tag{8}$$

The step sizes $\alpha_k = \mathrm{softplus}(a_k)$ and $\beta_k = \mathrm{softplus}(b_k)$ are learned per stage, and the factor s^2 compensates the $1/s^2$ gain of $\mathcal{D}^\top\mathcal{D}$ on smooth images. The prior step is a learned residual, $\mathbf{X}_k = \mathbf{Z}_k + \Delta_k$ with $(\Delta_k, \mathbf{F}_k) = \mathcal{P}_k([\mathbf{Z}_k, \mathcal{R}_\phi^+\mathbf{Y}_\mathrm{M}, s^2\mathcal{D}^\top(\mathcal{D}\mathbf{Z}_k - \mathbf{Y}_\mathrm{H}), \frac{B}{M}\mathcal{R}_\phi^\top(\mathcal{R}_\phi\mathbf{Z}_k - \mathbf{Y}_\mathrm{M})], \mathbf{F}_{k-1})$. Here $\mathcal{R}_\phi^+\mathbf{Y}_\mathrm{M}$ are the minimum-norm spectra in the B hyperspectral bands that reproduce the HR-MSI ($\mathcal{R}_\phi\mathcal{R}_\phi^+ = \mathbf{I}$, so they determine $\mathbf{Y}_\mathrm{M}$ exactly), and the last two inputs tell the prior where the estimate still disagrees with each observation. Every input lives in the hyperspectral bands, so the prior does not depend on the number or the responses of the multispectral bands. $\mathbf{F}_{k-1}$ are the decoder features of the previous stage (cross-stage memory [26]). The last layer of every prior starts at zero, so training starts from plain HQS, for which the condition $\mathcal{P}_k(\mathbf{X}^*) = \mathbf{0}$ of Theorem 2(c) holds. All operator applications run in 32-bit precision.

Learned spectral data step. In the basis of Theorem 1 the exact step multiplies each component of the data residual by $1/(1+\alpha_k\lambda_i^\mathrm{r}\lambda_j^\mathrm{c} + \beta_k\nu_b)$, the filter of a white Gaussian prior. Images are not white and noise does not affect all components equally, so we let two small networks modulate these gains:

$$\mathbf{Z}_k = \mathbf{X}_{k-1} + \mathcal{Q}\big(\mathcal{Q}^{-1}(\mathbf{G}_k) \odot \mathbf{\Gamma}_k\big), \qquad \Gamma_{k,bij} = \frac{\exp\big(\ln 2\, \tanh(u_k(\lambda_i^\mathrm{r}\lambda_j^\mathrm{c}) + v_k(\nu_b))\big)}{1 + \alpha_k\lambda_i^\mathrm{r}\lambda_j^\mathrm{c} + \beta_k\nu_b}, \tag{9}$$

where $\mathbf{G}_k = \alpha_k\mathcal{D}_\theta^\top(\mathbf{Y}_\mathrm{H} - \mathcal{D}_\theta\mathbf{X}_{k-1}) + \beta_k\mathcal{R}_\phi^\top(\mathbf{Y}_\mathrm{M} - \mathcal{R}_\phi\mathbf{X}_{k-1})$ and u_k, v_k are two-layer perceptrons of the log-eigenvalues (and of the observation code below), initialized to zero so that training starts from (5). The gain factors over space and spectrum, so the perceptrons run $HW + B$ times per stage, not BHW times.

Proposition 1. *For any u_k, v_k: (a) with exact operators and noise-free data, $\mathbf{X}^*$ is a fixed point of (9); (b) the step is nonexpansive; (c) with an inexact operator it moves $\mathbf{X}^*$ by at most $\|\mathbf{e}\|$, twice the bound (6) of the exact step.*

The gain lies in $(\frac{1}{2}, 2)$ times the exact one, which keeps every factor of the step in $(-1, 1]$ (the supplementary material). The learned step can thus change how much the network trusts each observed component, but it cannot move a correct estimate or amplify errors.

Prior. Each prior is a three-level U-Net [24] of Restormer blocks [37] (Fig. 1b,c), whose transposed attention forms a channel-by-channel attention map with a cost linear in the number of pixels. ExactFuse holds 10.5 M parameters for Chikusei, almost all in the three priors. Their design is not our contribution, and Sec. 5.4 shows that a convolutional mixer works about as well.

Training over a family of observation models. A network trained with one operator pair on noise-free data adapts its priors to that blur, that response and no noise. Since the operators enter ExactFuse only through its data steps, we can instead draw for every batch a blur width $\sigma \sim \mathcal{U}[0.5, 2]\sigma_0$, a spectral shift $\delta \sim \mathcal{U}[-12, 12]$ nm and, with probability 0.5, noise at an SNR from $\mathcal{U}[30, 50]$ dB, simulate the batch with $(\mathcal{D}_\sigma, \mathcal{R}_\delta)$, and use the same operators in the data steps. At test time the operators are given or estimated with (7). As HQS suggests, we also condition on the noise: a code $\mathbf{c} = (\sigma/\sigma_0 - 1, \delta/12, \log_{10} n_\mathrm{H} + 2, \log_{10} n_\mathrm{M} + 2)$, with the relative noise levels n estimated from each observation by the median absolute deviation of its second differences, sets the step sizes $\alpha_k = \mathrm{softplus}(a_k + \mathbf{u}_k^\top\mathbf{c})$, $\beta_k = \mathrm{softplus}(b_k + \mathbf{v}_k^\top\mathbf{c})$ and modulates the prior features (FiLM [21]), similar to the hyper-parameter module of USRNet [39].

Loss. We train with $\|\mathbf{X}_K - \mathbf{X}\|_1 + 0.1\sum_{k<K}\|\mathbf{X}_k - \mathbf{X}\|_1 + 0.05\,\mathrm{SAM}(\mathbf{X}_K, \mathbf{X}) + 0.1\,(1 - \mathrm{SSIM}(\mathbf{X}_K, \mathbf{X}))$ and AdamW, and test the exponential moving average of the weights.

5. Experiments

5.1. Setup

Data. Chikusei [35] has 128 bands (363–1018 nm) and is paired with the eight WorldView-2 responses; Pavia Centre [11] has 92 bands and is paired with the four IKONOS responses. We use the split sizes of [5] (8 and 4 test images of 256×256 pixels) with fixed, disjoint locations given in the supplementary material. CAVE [34] (20 training and 12 test scenes) and Harvard [3] (30 and 20 images) have 31 bands and are paired with a Nikon RGB response. All pairs follow Wald's protocol [28]: both observations are simulated from the reference with (1), every training crop on its own. At $\times8$ and $\times16$ the blur is a $(2s-1)\times(2s-1)$ Gaussian with $\sigma = s/2$.

Baselines and budget. We compare with GSA [1] and twelve networks: SSR-NET [41], PSRT [6], DCT [17], MIMFormer [14], DHIF-Net [12], U2Net [20], MIMO-SST [10], OTPNet [36], SMGU-Net [32], CLSNet [33], CaFormer [2] and M2DTN [9]. Each is built from its authors' code; only the numbers of bands and the scale factor are changed (the supplementary material lists the commits).

All methods, ExactFuse included, are trained for 4 h on one Tesla T4 GPU with the same patches (64×64), batch size (8), augmentation, EMA weights and validation-only model selection; SSR-NET, which converges sooner, gets 1.5 h. The published networks keep their published learning rates with an L1 loss. The metrics are computed by the same code for all methods: PSNR (peak equal to the dataset maximum), SSIM, the spectral angle SAM (degrees) and ERGAS; Q2n, CC, SCC and RMSE are in the supplementary material.

5.2. Comparison

Table 1 shows the $\times 4$ results. ExactFuse reaches 58.05 dB on Chikusei and 50.63 dB on Pavia Centre, 0.79 and 0.55 dB above the best retrained network, and it has the lowest SAM and ERGAS on both datasets. Over three seeds, the PSNR of ExactFuse and of MIMFormer varies by 0.01 and 0.18 dB on Chikusei, and the worst ExactFuse seed is still 1.09 dB above the best MIMFormer seed. Pooled over the 96 test images of all 9 settings below, ExactFuse is more accurate than the strongest retrained network of the setting on 95 of them (one-sided Wilcoxon signed-rank test, $p = 9.2 \times 10^{-18}$). Table 3 extends the comparison to CAVE and Harvard, whose test scenes are unseen in training, and to $\times 8$ and $\times 16$: ExactFuse has the highest PSNR in 7 of the 7 settings, by 0.21–2.29 dB over the strongest network run in each.

A fixed budget raises two questions. First, does our pipeline weaken the baselines? It does not: the retrained networks match or exceed the values reported for them, e.g., MIMFormer reaches 56.96 dB on Chikusei against 53.36 dB in [5], and their authors' own training recipes do not improve them (the supplementary material). Second, does the budget favor fast networks? Partly: DCT completes only 4,800 iterations in 4 h on Pavia Centre, against 33,850 for ExactFuse. With 8 h of training, ExactFuse and MIMFormer reach 58.18 and 57.37 dB on Chikusei; the gap narrows from 1.10 to 0.81 dB but remains, although ExactFuse completes 63,850 iterations against 186,800.

For context, the supplementary material lists ExactFuse trained for 8 h next to the 17 methods tabulated in [5]; that table cannot rank methods, since our retrained baselines exceed their reported values by up to 3.6 dB, and its strongest entry, a diffusion model, has no public code.

5.3. Where the margin comes from

We reduce ExactFuse step by step (the supplementary material). A single U-Net prior on the bicubic image and the HR-MSI, without data steps or stages, already reaches 57.67 dB on Chikusei and 50.39 dB on Pavia Centre, above every retrained network. Most of the margin therefore comes from this residual multiscale prior and its training recipe, not from the transformer blocks, and the exact data steps add 0.38 and 0.23 dB on top. The recipe is not the whole story: with the same recipe for both, ExactFuse leads MIMFormer by 0.53 and 0.40 dB under ours and by 0.72 and 0.44 dB under the baselines'.

The null space of Theorem 1(c) gives a second view: we split the error of every method into its part in the observed subspace $\mathrm{row}(\mathcal{D}) + \mathrm{row}(\mathcal{R})$, which contradicts the observations, and its null-space part, which only a prior can reduce. Bicubic interpolation and GSA err mostly in the observed subspace (only 8.6% and 13.7% of their squared error lies in the null space on Chikusei). ExactFuse's observed error (column Obs. of Tab. 1) is 27 times smaller than that of the best retrained network on Chikusei (MIMO-SST) and 9 times on Pavia Centre. Removing MIMFormer's observed error alone would raise its PSNR by 0.43 dB on Chikusei; the remaining 0.67 dB of ExactFuse's margin come from a better null-space estimate. The exact data steps account for the first part, the prior for the second.

5.4. What the exact operator does

Table 4 retrains variants from scratch with the same budget. Its top block goes from the base model to ExactFuse. The closed-form step with learned gains and the HR-MSI residual input each raise the agreement with the LR-HSI (89.4 and 86.5 against 83.9 dB on Chikusei) at nearly equal PSNR; combined, with every input in the hyperspectral bands, they change the PSNR by +0.10 and 0.00 dB (Chikusei, Pavia Centre) and raise the agreement to 93.4 dB. A learned gate that confines the prior to the null space lowers the PSNR (-0.08 and -0.04 dB); ExactFuse does not use it.

Without data steps ExactFuse loses 0.63 and 0.33 dB and its agreement with the LR-HSI falls from 93.4 to 67.5 dB; a gradient step instead of the closed form keeps the PSNR but lowers the agreement to 86.7 dB. The learned gains, the residual input and $\mathcal{R}^{+}\mathbf{Y}_{\mathrm{M}}$ each change the PSNR by at most 0.05 dB, a zero border by 0.17 and 0.06 dB. In the base model, a zero border or a learned blur and response cost 0.22 and 0.22 dB on Chikusei and reproduce the LR-HSI about 3 dB worse, as Theorem 2 predicts; over three seeds the zero border costs 0.22 and 0.08 dB, against a seed spread of 0.004 and 0.02 dB.

The difference becomes large when the blur changes (the supplementary material): given the true blur, the exact base model loses at most 0.25 dB on Chikusei, the zero-border model 3.41 dB and the learned operator, which cannot be replaced, 3.45 dB. A convolutional mixer instead of attention changes the PSNR by only +0.07 and -0.02 dB (Chikusei, Pavia Centre).

5.5. One network for a family of observation models

Table 2 tests networks trained over the operator family on conditions that leave the training protocol. Within the family the estimate of Theorem 3 is accurate: over all Chikusei test images and conditions, 30 dB SNR included, its largest errors are 0.05 px in blur width and 0.0 nm in spectral shift.

Table 1. Equal-budget comparison at $\times 4$. Every network is retrained by us from its authors' code with the same data, patches, batch size, EMA weights, validation-only model selection and 4 h on one T4 GPU (single pass). SSIM with the data range set to each image's own maximum. Obs.: mean squared error in the observed subspace ($\times 10^{-6}$), the part of the error that contradicts the observations (Sec. 5.3). Parameters for 128 bands. Best learned method in bold.

		Chikusei $\times 4$					Pavia Centre $\times 4$				
Method	Params (M)	PSNR↑	SSIM↑	SAM↓	ERGAS↓	Obs.↓	PSNR↑	SSIM↑	SAM↓	ERGAS↓	Obs.↓
Bicubic	–	33.63	0.7404	4.406	9.011	431	25.52	0.6282	6.554	9.852	2746
GSA [1]	–	43.20	0.9778	2.730	3.279	46.9	34.78	0.9438	5.229	3.494	370
SSR-NET [41]	0.44	55.55	0.9958	0.954	1.707	0.703	48.45	0.9954	2.443	0.758	2.049
PSRT [6]	0.30	53.99	0.9948	1.057	1.794	1.645	47.87	0.9958	2.212	0.801	2.483
DCT [17]	8.46	53.18	0.9945	1.133	2.066	2.240	46.94	0.9953	2.442	0.888	4.312
MIMFormer [14]	16.46	56.96	0.9964	0.801	1.488	0.186	49.63	0.9967	1.982	0.674	0.430
DHIF-Net [12]	48.67	54.39	0.9953	1.020	1.660	1.546	47.67	0.9956	2.295	0.840	3.980
U2Net [20]	3.06	53.35	0.9943	1.076	1.948	–	47.29	0.9955	2.327	0.858	–
MIMO-SST [10]	5.32	57.27	0.9965	0.780	1.419	0.066	49.69	0.9967	1.961	0.670	0.257
OTPNet [36]	2.51	57.06	0.9965	0.786	1.448	0.218	50.08	0.9968	1.951	0.653	0.677
SMGU-Net [32]	9.50	55.14	0.9957	0.949	1.705	0.842	48.03	0.9958	2.331	0.783	2.293
CLSNet [33]	4.17	56.65	0.9963	0.811	1.514	0.242	49.50	0.9967	1.993	0.684	0.639
CaFormer [2]	7.83	53.56	0.9949	1.054	2.138	2.095	48.43	0.9961	2.145	0.757	1.713
M2DTN [9]	23.79	54.87	0.9952	1.003	1.977	1.410	48.65	0.9963	2.156	0.729	2.103
ExactFuse	10.45	**58.05**	**0.9969**	**0.724**	**1.334**	**0.002**	**50.63**	**0.9970**	**1.892**	**0.619**	**0.028**

Table 2. One network for a family of observation models: mean PSNR (dB) over groups of test conditions. Randomized: blur width (0.5–$2\sigma_0$), SRF shift (± 12 nm) and noise (30–50 dB) drawn anew for every training batch. The operator is the training one (MIMFormer, randomized data), or estimated from the test pair with (7); the last row uses the true operator (reference, not marked). †Anisotropic and defocus blurs, outside the family. §LR-HSI sampled on a grid offset by 0.5–1 HR pixel (output on the HR-MSI grid), alone and with an axis-aligned anisotropic blur. ‡HR-MSI shifted by 0.5–1 pixel against the reference. Best per column in bold.

	Chikusei $\times 4$								Pavia Centre $\times 4$							
Model	Nom.	Blur	SRF	Noise	All	Kernel†	Offset§	Misreg.‡	Nom.	Blur	SRF	Noise	All	Kernel†	Offset§	Misreg.‡
MIMFormer	56.96	54.18	53.61	50.33	48.98	54.23	50.60	38.00	49.63	47.05	48.66	45.24	44.15	46.89	43.83	29.86
MIMFormer, randomized	56.50	56.38	55.66	51.67	51.30	56.43	54.90	37.46	49.37	49.28	48.31	45.87	45.22	49.32	48.21	29.33
ExactFuse	**58.05**	**57.79**	55.13	47.60	46.71	50.79	44.82	**38.38**	**50.63**	**50.50**	49.02	43.15	42.32	42.37	37.40	**30.31**
base model, randomized data	57.43	57.33	57.35	52.35	52.30	57.18	55.40	37.48	50.40	50.33	50.32	46.40	46.35	50.01	47.79	29.50
+ estimated operator, conditioning	57.61	57.51	57.54	**52.92**	**52.83**	54.78	50.87	38.12	50.42	50.35	50.38	**46.84**	**46.80**	47.02	43.20	30.01
ExactFuse, randomized + conditioning	57.77	57.68	**57.73**	52.64	52.59	55.34	51.90	38.06	50.43	50.38	**50.39**	46.55	46.53	46.76	43.01	30.06
wide estimate	57.77	57.68	**57.73**	52.64	52.59	56.84	56.07	37.48	50.43	50.38	**50.39**	46.55	46.53	48.71	49.12	29.32
any-shape kernel (least squares)	57.76	–	–	–	–	**57.60**	–	–	50.38	–	–	–	–	**50.30**	–	–
ExactFuse, wider family (+ aniso, offsets)	57.75	57.65	57.70	52.55	52.50	56.85	**56.10**	37.48	50.40	50.36	50.36	46.55	46.54	48.84	**49.17**	29.32
conditioning, true operator	57.77	57.68	57.73	52.64	52.59	57.59	56.08	–	50.43	50.38	50.39	46.55	46.53	50.33	49.23	–

Training over the family makes the network robust to noise and combined changes: at 40–30 dB SNR the mean PSNR of ExactFuse rises from 47.60 to 52.64 dB on Chikusei, against 51.67 dB for MIMFormer trained on the same randomized data. The price is a lower PSNR at the training operator (57.77 against 58.05 dB). One network trained on random responses also serves seven real sensors it never saw, 3.2–11.3 dB above MIMFormer trained the same way on Chikusei and 1.79 dB below ExactFuse trained for WorldView-2, and one network trained at $\times 4$ and $\times 8$ leads the best retrained network at the unseen $\times 16$ by 1.76 dB (the supplementary material).

Outside the family a parametric estimate fails: with an anisotropic or defocus blur ExactFuse drops to 55.34 dB on Chikusei. Because $\mathbf{r}$ is linear in the kernel, a kernel of any shape can instead be estimated by least squares; with it ExactFuse reaches 57.60 and 50.30 dB under these blurs, as with the true kernel (57.59 and 50.33 dB) and ahead of MIMFormer trained on the same randomized data (56.43 and 49.32 dB). An HR-MSI misregistered by half a pixel

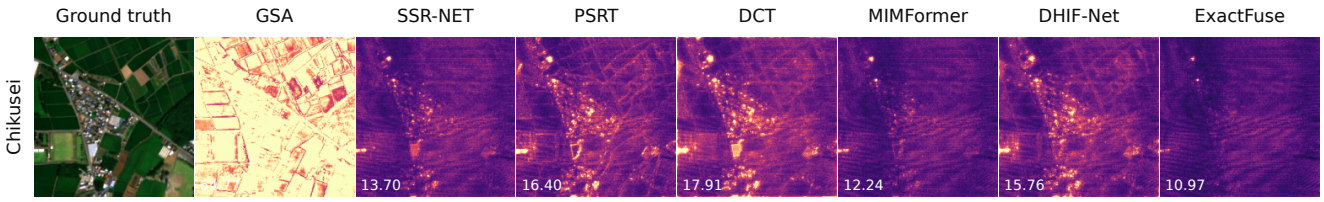


Figure 3. Mean absolute error over all bands (DN) on the Chikusei test image where ExactFuse has its lowest PSNR (Pavia Centre, CAVE and Harvard in the supplementary material). The color scale is set by the retrained networks (clipped at the 99.5th percentile), so GSA saturates; the number in each map is its mean.

Table 3. PSNR (dB) on CAVE and Harvard (unseen test scenes) and at larger scale factors, same protocol and budget as Tab. 1. –: not run (SMGU-Net's code supports $\times 4$ only).

Method	CAVE $\times 4$	Harv. $\times 4$	CAVE $\times 8$	Harv. $\times 8$	Chik. $\times 8$	Pavia $\times 8$	Chik. $\times 16$
GSA	42.30	46.30	38.01	43.13	38.69	30.24	35.20
SSR-NET	46.51	49.18	44.25	48.61	54.57	47.71	53.98
PSRT	48.51	49.29	45.21	–	53.05	47.13	–
DCT	46.36	48.92	43.17	–	51.77	45.51	–
MIMFormer	50.93	49.72	48.10	49.12	55.57	49.09	54.21
DHIF-Net	49.82	49.58	–	–	–	–	–
U2Net	48.04	48.70	–	–	–	–	–
MIMO-SST	50.39	49.83	–	–	55.88	48.90	–
OTPNet	49.45	49.53	–	–	56.11	49.51	–
SMGU-Net	48.54	49.34	–	–	–	–	–
CLSNet	–	–	–	–	55.59	49.02	–
CaFormer	47.91	49.09	–	–	–	–	–
ExactFuse	**52.46**	**50.03**	**50.39**	**49.45**	**56.95**	**50.05**	**56.11**

Table 4. Ablation of ExactFuse, each variant trained from scratch for 4 h: one change per row, then one input removed, and the base model for reference (all variants in the supplementary material). Cons.: PSNR$(\mathcal{D}\hat{\mathbf{X}}, \mathbf{Y}_{\mathrm{H}})$. Swap: worst PSNR over test blurs $\sigma \in \{0.75, 1.25, 1.5\}\sigma_0$ when the true blur is given to the data steps (n/a: no operator to replace; –: not evaluated).

	Chikusei			Pavia Centre		
Variant	PSNR	Cons.	Swap	PSNR	Cons.	Swap
ExactFuse	58.05	93.4	57.93	50.63	82.5	50.57
closed form without learned gains	58.00	90.4	57.88	50.67	80.6	50.62
gradient step instead of closed form	58.03	86.7	57.84	50.63	79.5	50.56
no HR-MSI residual input	58.03	92.7	57.91	50.60	81.7	50.54
$\mathbf{Y}_{\mathrm{M}}$ instead of $\mathcal{R}^{+}\mathbf{Y}_{\mathrm{M}}$ as input	58.07	93.9	57.96	50.58	82.6	50.52
no data steps	57.42	67.5	n/a	50.30	68.1	n/a
zero border	57.88	80.9	57.60	50.56	75.2	50.41
+ null-space gate	57.97	96.9	57.82	50.58	87.0	50.48
+ gate, EON prior instead of U-Net	57.83	96.5	57.74	50.37	84.9	50.32
LR-HSI only (no HR-MSI)	39.62	62.1	n/a	30.26	59.3	n/a
HR-MSI only (no LR-HSI)	54.88	59.7	n/a	48.23	53.5	n/a
base model (gradient step, $\mathbf{Y}_{\mathrm{M}}$ input)	57.95	83.9	57.70	50.63	77.6	50.50

costs every method, GSA included, more than 15 dB. Registering it before fusion with the offset estimated from the pair recovers most of this loss (the supplementary material): for a sharp half-pixel shift, as a sensor records it, ExactFuse reaches 54.56 dB on Chikusei (MIMFormer 54.05, GSA 43.06 dB), and with 35 dB noise added, trained over the wider family, 51.57 dB. The linear estimate thus extends the operator family to any spatially invariant kernel.

6. Limitations and conclusion

Limitations. (i) Results with a reference rest on simulated observations (Wald's protocol), as in every benchmark of this field; on real NEON data (spectrometer and RGB camera, the supplementary material), ExactFuse is again the most accurate at reduced resolution (38.70 against 38.45 dB for MIMFormer), and has the best no-reference score at full resolution (HQNR 0.9203 against 0.8923), and leads by 1.11 dB on an unseen site; its operator variants differ by less than 0.1 dB there, because the fitted camera response is only approximate; the sensor tests use real responses on simulated pairs. (ii) Exact operators help only when the operator family contains the true operator: spatially varying blurs fall outside even the least-squares kernel estimate, which costs 0.10–0.16 dB at 30 dB SNR, registration leaves several decibels of a misregistration loss (registration-aware networks such as IR&ArF [23] were not retrained), non-separable kernels lose the closed form, and corrupted pixels in a four-band HR-MSI are tolerated but localized only by a model-free detector (the supplementary material). (iii) All methods train for 4 h; the strongest baselines are close to converged (the supplementary material), and 8 h narrows the gap but keeps it. (iv) The Chikusei and Pavia Centre test sets are small; three seeds, bootstrap intervals (the supplementary material), the pooled test and CAVE and Harvard back the comparison. (v) ExactFuse takes 0.39 s per 256×256 pair on a T4, operator estimate included, about twice as long as MIMFormer (the supplementary material).

Exact separable operators with the true border give unfolded fusion a closed-form data step, an explicit null space, a bounded effect of operator errors and operators identifiable from the two observations. Under one protocol and budget, ExactFuse is more accurate than twelve retrained networks, mainly through its prior and training, and its exact operators let one network adapt to new blurs, responses, noise, sensors and scale factors. Spatially varying operators, band-dependent noise and real paired references remain open. Code will be made publicly available.

LLM Usage Statement

A large language model assistant was used extensively in this work: to write and debug the training and evaluation code, to run and analyze the experiments on cloud GPUs, and to draft and edit the text. The authors directed the work, reviewed the code, results and claims, and take full responsibility for the content.

References

[1] B. Aiazzi, S. Baronti, and M. Selva. Improving component substitution pansharpening through multivariate regression of MS +Pan data. *IEEE Transactions on Geoscience and Remote Sensing*, 45(10):3230–3239, 2007. 5, 7

[2] Xuheng Cao, Xuquan Wang, Xiong Dun, Yusheng Lian, Xinbin Cheng, and Xiaopeng Hao. Cross-domain-aware deep unfolding transformer for hyperspectral image super-resolution. *Pattern Recognition*, 172:112374, 2026. 5, 7

[3] Ayan Chakrabarti and Todd Zickler. Statistics of real-world hyperspectral images. In *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, pages 193–200, 2011. 5

[4] Dongdong Chen and Mike E. Davies. Deep decomposition learning for inverse imaging problems. In *European Conference on Computer Vision (ECCV)*, pages 510–526, 2020. 2

[5] Yingxia Chen, Ruijie Liu, Wai Keung Wong, and Jie Wen. A two-stage conditional diffusion model with differential attention for hyperspectral and multispectral image fusion. *IEEE Transactions on Image Processing*, 35:8980–8995, 2026. 1, 2, 5, 6

[6] Shang-Qi Deng, Liang-Jian Deng, Xiao Wu, Ran Ran, Danfeng Hong, and Gemine Vivone. PSRT: Pyramid shuffle-and-reshuffle transformer for multispectral and hyperspectral image fusion. *IEEE Transactions on Geoscience and Remote Sensing*, 61:1–15, 2023. 1, 2, 5, 7

[7] Jiangxin Dong, Stefan Roth, and Bernt Schiele. Deep Wiener deconvolution: Wiener meets deep learning for image deblurring. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. 2

[8] Weisheng Dong, Chen Zhou, Fangfang Wu, Jinjian Wu, Guangming Shi, and Xin Li. Model-guided deep hyperspectral image super-resolution. *IEEE Transactions on Image Processing*, 30:5754–5768, 2021. 1, 2

[9] Wenqian Dong, Yang Xu, Jiahui Qu, and Shaoxiong Hou. Learning multi-modal cross-scale deformable transformer network for unregistered hyperspectral image super-resolution. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 1573–1581, 2024. 5, 7

[10] Jian Fang, Jingxiang Yang, Abdolraheem Khader, and Liang Xiao. MIMO-SST: Multi-input multi-output spatial-spectral transformer for hyperspectral and multispectral image fusion. *IEEE Transactions on Geoscience and Remote Sensing*, 62: 1–20, 2024. 2, 5, 7

[11] P. Gamba. A collection of data for urban area characterization. In *Proc. IEEE Int. Geosci. Remote Sens. Symp. (IGARSS)*, pages 69–72, 2004. 5

[12] Tao Huang, Weisheng Dong, Jinjian Wu, Leida Li, Xin Li, and Guangming Shi. Deep hyperspectral image fusion network with iterative spatio-spectral regularization. *IEEE Transactions on Computational Imaging*, 8:201–214, 2022. 1, 2, 5, 7

[13] Charilaos I. Kanatsoulis, Xiao Fu, Nicholas D. Sidiropoulos, and Wing-Kin Ma. Hyperspectral super-resolution: A coupled tensor factorization approach. *IEEE Transactions on Signal Processing*, 66(24):6503–6517, 2018. 2

[14] Rumei Li, Liyan Zhang, Zun Wang, and Xiaojuan Li. MIMFormer: Multiscale inception mixer transformer for hyperspectral and multispectral image fusion. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 17:15122–15135, 2024. 1, 2, 5, 7

[15] Shutao Li, Renwei Dian, Leyuan Fang, and Jose M. Bioucas-Dias. Fusing hyperspectral and multispectral images via coupled sparse tensor factorization. *IEEE Transactions on Image Processing*, 27(8):4118–4130, 2018. 2

[16] Yujie Liang, Zihan Cao, Liang-Jian Deng, Yang Yang, and Malu Zhang. Hyperspectral image fusion with spectral-band and fusion-scale agnosticism. In *International Conference on Machine Learning (ICML)*, 2026. 2

[17] Qing Ma, Junjun Jiang, Xianming Liu, and Jiayi Ma. Reciprocal transformer for hyperspectral and multispectral image fusion. *Information Fusion*, 104:102148, 2024. 1, 2, 5, 7

[18] Vishal Monga, Yuelong Li, and Yonina C. Eldar. Algorithm unrolling: Interpretable, efficient deep learning for signal and image processing. *IEEE Signal Processing Magazine*, 38(2): 18–44, 2021. 2

[19] Jieyuan Pei, Wei Li, Zhuoxuan Li, Junwei Zhu, Meiyi Lu, Jiawei Jiang, Chenyu Wang, and Jianwei Zheng. S^3O: Selective spatial-spectral operator for cross-scale fusion. In *Findings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2026. 1, 2

[20] Siran Peng, Chenhao Guo, Xiao Wu, and Liang-Jian Deng. U2Net: A general framework with spatial-spectral-integrated double U-Net for image fusion. In *Proceedings of the 31st ACM International Conference on Multimedia*, pages 3219–3227, 2023. 2, 5, 7

[21] Ethan Perez, Florian Strub, Harm De Vries, Vincent Dumoulin, and Aaron Courville. FiLM: Visual reasoning with a general conditioning layer. In *Proc. AAAI Conf. Artif. Intell.*, 2018. 5

[22] Clémence Prévost, Konstantin Usevich, Pierre Comon, and David Brie. Hyperspectral super-resolution with coupled Tucker approximation: Recoverability and SVD-based algorithms. *IEEE Transactions on Signal Processing*, 68:931–946, 2020. 2

[23] Jiahui Qu, Xiaoyang Wu, Wenqian Dong, Jizhou Cui, and Yunsong Li. IR&ArF: Toward deep interpretable arbitrary resolution fusion of unregistered hyperspectral and multispectral images. *IEEE Transactions on Image Processing*, 34:1934–1949, 2025. 2, 8

[24] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-Net: Convolutional networks for biomedical image segmentation. In *Proc. Med. Image Comput. Comput.-Assist. Interv. (MICCAI)*, pages 234–241, 2015. 5

[25] Miguel Simoes, Jose Bioucas-Dias, Luis B. Almeida, and Jocelyn Chanussot. A convex formulation for hyperspectral image superresolution via subspace-based regularization. *IEEE Transactions on Geoscience and Remote Sensing*, 53 (6):3373–3388, 2015. 3

[26] Jiechong Song, Bin Chen, and Jian Zhang. Memory-augmented deep unfolding network for compressive sensing. In *Proc. 29th ACM Int. Conf. Multimedia*, pages 4249–4258, 2021. 5

[27] M. R. Vicinanza, R. Restaino, G. Vivone, M. Dalla Mura, G. Licciardi, and J. Chanussot. A method for improving the consistency property of pansharpening algorithms. In *IEEE International Geoscience and Remote Sensing Symposium (IGARSS)*, pages 2534–2537, 2014. 2

[28] Gemine Vivone, Luciano Alparone, Jocelyn Chanussot, Mauro Dalla Mura, Andrea Garzelli, Giorgio A. Licciardi, Rocco Restaino, and Lucien Wald. A critical comparison among pansharpening algorithms. *IEEE Transactions on Geoscience and Remote Sensing*, 53(5):2565–2586, 2015. 5

[29] Gemine Vivone, Liang-Jian Deng, Shangqi Deng, Danfeng Hong, Menghui Jiang, Chenyu Li, Wei Li, Huanfeng Shen, Xiao Wu, Jin-Liang Xiao, Jing Yao, Mengmeng Zhang, Jocelyn Chanussot, Salvador García, and Antonio Plaza. Deep learning in remote sensing image fusion: Methods, protocols, data, and future perspectives. *IEEE Geoscience and Remote Sensing Magazine*, 13(1):269–310, 2025. 1

[30] Qi Wei, Nicolas Dobigeon, and Jean-Yves Tourneret. Fast fusion of multi-band images based on solving a Sylvester equation. *IEEE Transactions on Image Processing*, 24(11):4109–4121, 2015. 2, 3

[31] Qi Xie, Minghao Zhou, Qian Zhao, Zongben Xu, and Deyu Meng. MHF-Net: An interpretable deep network for multispectral and hyperspectral image fusion. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 44(3):1457–1473, 2022. 1, 2

[32] Jun Yan, Kai Zhang, Qinzhu Sun, Chiru Ge, Wenbo Wan, Jiande Sun, and Huaxiang Zhang. Spatial–spectral unfolding network with mutual guidance for multispectral and hyperspectral image fusion. *Pattern Recognition*, 161:111277, 2025. 1, 2, 5, 7

[33] Heng Yang, Jinyang Liu, Renwei Dian, and Shutao Li. Correlation-aware learning self-guiding network for hyperspectral fusion imaging. *Information Fusion*, 133:104263, 2026. 2, 5, 7

[34] Fumihito Yasuma, Tomoo Mitsunaga, Daisuke Iso, and Shree K. Nayar. Generalized assorted pixel camera: Post-capture control of resolution, dynamic range, and spectrum. *IEEE Transactions on Image Processing*, 19(9):2241–2253, 2010. 5

[35] Naoto Yokoya and Akira Iwasaki. Airborne hyperspectral data over Chikusei. Technical Report SAL-2016-05-27, Space Application Laboratory, University of Tokyo, Tokyo, Japan, 2016. 5

[36] Wei Yu, Zonglin Li, Qinglin Liu, and Xin Sun. OTPNet: ODE-inspired tuning-free proximal network for remote sensing image fusion. *Proceedings of the AAAI Conference on Artificial Intelligence*, 39(9):9671–9679, 2025. 2, 5, 7

[37] Syed Waqas Zamir, Aditya Arora, Salman Khan, Munawar Hayat, Fahad Shahbaz Khan, and Ming-Hsuan Yang. Restormer: Efficient transformer for high-resolution image restoration. In *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 5718–5729, 2022. 4, 5

[38] G. L. Zeng and G. T. Gullberg. Unmatched projector/backprojector pairs in an iterative reconstruction algorithm. *IEEE Transactions on Medical Imaging*, 19(5):548–555, 2000. 2

[39] Kai Zhang, Luc Van Gool, and Radu Timofte. Deep unfolding network for image super-resolution. In *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, pages 3214–3223, 2020. 1, 2, 5

[40] Kai Zhang, Jingyun Liang, Luc Van Gool, and Radu Timofte. Designing a practical degradation model for deep blind image super-resolution. In *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, pages 4771–4780, 2021. 2

[41] Xueting Zhang, Wei Huang, Qi Wang, and Xuelong Li. SSR-NET: Spatial–spectral reconstruction network for hyperspectral and multispectral image fusion. *IEEE Transactions on Geoscience and Remote Sensing*, 59(7):5953–5965, 2021. 1, 2, 5, 7