ExactFuse: Closed-Form Unfolding With Exact Operators for Sensor-Independent Hyperspectral and Multispectral Image Fusion

Hruthwik, Amarnath M, Navadeep, and Srilatha Chirbolu

***Abstract*—A hyperspectral camera records hundreds of spectral bands but loses fine spatial detail, while a multispectral camera does the opposite. Deep unfolding networks are a leading way to fuse the two. Two problems hold them back: they learn or approximate the blur and spectral response instead of using the true ones, and they must be retrained for each sensor pair. We present ExactFuse, an unfolding network built around exact operators. With a separable blur and the true border rule of the simulation, the data step has a closed form with learned per-component gains, and the part of the image that neither observation constrains, 88% of the unknowns on Chikusei, is explicit. The two observations alone identify the true operators among candidates for almost every scene, so they are estimated at test time. An exact data step leaves the true image fixed, and a wrong operator can move it only by a bounded amount, whatever the network has learned. The transformer prior sees the multispectral image only as spectra in the hyperspectral bands, and the network is trained over a family of blurs, responses and noise levels, so one network serves different sensors. Against twelve published networks retrained from their authors' code under one budget, ExactFuse is the most accurate on Chikusei and Pavia Centre, by 0.79 and 0.55 dB, and in all eight further settings, up to $\times 16$, and reproduces the low-resolution input 17 dB more accurately than the strongest of them. Much of that lead comes from the prior and its training; the exact operators matter most when the blur, sensor or noise changes. A half-pixel misregistration still costs every network trained on registered pairs, ExactFuse included, more than 15 dB; registering the images first recovers most of that, but several decibels remain.**

***Index Terms*—Deep unfolding, hyperspectral image fusion, identifiability, Kronecker structure, null space, observation model, super-resolution.**

I. Introduction

HYPERSPECTRAL imagers record hundreds of narrow spectral bands but, for a fixed photon budget, at a coarse spatial resolution; multispectral imagers do the opposite. Hyperspectral and multispectral image fusion recovers a high-resolution hyperspectral image (HR-HSI) from one observation of each kind. Classical methods formulate fusion as a regularized inverse problem, using component substitution [1], coupled matrix factorization, subspace regularization or fast Sylvester-equation solvers [2]. Deep networks have since become the dominant approach, from convolutional reconstruction networks [3] to transformers [4], [5] and diffusion models [6].

Model-guided and unfolded networks [7], [8] bring the observation model back into the network: each stage takes a step that reduces the mismatch between the current estimate and the two observations, and a learned module then refines the estimate. How faithfully the network reproduces the observation model has received less attention than the design of the learned module. In most of these networks the blur–decimation and spectral operators are learned convolution layers, or are implemented with the zero padding that convolution layers apply at the image border. The benchmark data, however, are simulated with a fixed blur and a different border rule. The network's operator is then not the one that produced the data, and even the ground truth fails the network's own consistency check. In simulated benchmarks this mismatch is avoidable, because the forward model is known exactly. A second difficulty concerns evaluation: reported numbers on the same dataset vary widely across papers because of differences in sensor response, blur, data split, normalization and even metric definitions, so improvements are only meaningful within one fully specified protocol in which all methods are run.

This paper treats the observation model inside an unfolded network as an object to be analyzed, estimated and exploited rather than approximated. We write fusion as maximum a posteriori estimation with parametric operators, unroll its half-quadratic splitting, and derive the properties that the operators of such a network have; the network, ExactFuse, is built from them. Our contributions are:

1) **Theory of the observation operators in unfolded fusion.** We prove that the blur–decimation and spectral operators commute and that their parameters are identified by the two observations alone for almost every scene (Theorem 1); that with a separable blur and the true border rule they factor along the three axes of the image cube, which gives the data step in closed form and the joint null space explicitly (Theorem 2); and that an exact data step leaves the ground truth fixed, while an inexact operator moves it by an amount that the exact proximal step bounds independently of the learned weights, with a corollary for border rules (Theorem 3).
2) **An unfolded network built on these results.** ExactFuse uses the closed-form exact data step, a learned spectral data step that reweights its eigencomponents while provably keeping the fixed point and nonexpansiveness

Hruthwik, Amarnath M and Navadeep carried out this work under the mentorship of Srilatha Chirbolu. Corresponding author: Amarnath M (e-mail: amarnath10chinu@gmail.com).

Code, training logs and all result files are available at https://github.com/Amarnath10i/Research-On-Hyperspectral-ImageFussion. Supplementary material (training stages, robustness, metric analyses, per-image results and complete benchmark metrics) accompanies this paper.

(Proposition 1), inputs of the prior in the hyperspectral bands only, so that one network serves any multispectral sensor, observation-model conditioning of the step sizes and the prior, and training over a family of blurs, spectral responses, noise levels and grid offsets with the operators estimated at test time.

3) **An evaluation that explains the results.** Twelve published networks are retrained from their authors' code under one budget and two of them also with their own recipes; the null-space split of the error shows how much of each method's error contradicts the observations and attributes ExactFuse's margin; ablations, tests inside and outside the operator family, and further datasets and scale factors delimit what the exact operators do and do not achieve.

II. Related Work

This section cites only the methods compared in our experiments and our closest references; Section S-VII of the supplementary material repeats it with all references.

A. Model-based fusion

Component substitution methods such as GSA [1] inject multispectral detail into the upsampled hyperspectral bands. Unmixing-based methods such as CNMF factor both observations with shared endmembers, and subspace methods such as HySure regularize the HR-HSI within a low-dimensional spectral subspace, whose dimension can be estimated with HySime. FUSE [2] solves the resulting quadratic problem in closed form through a Sylvester equation. Pineau et al. [9] obtain an explicit solution by inverting the Hessian of a quadratic criterion under a linear mixing model, with circulant blurs. These methods use the forward model explicitly but rely on hand-crafted priors.

B. Supervised deep fusion

SSR-NET [3] and GuidedNet learn a direct mapping from the two observations. Transformers model long-range spatial and spectral dependencies: PSRT [4] uses pyramid shuffle-and-reshuffle attention, MIMO-SST [5] multi-input multi-output supervision, LGCT local–global collaboration, DSPNet a dual spatial–spectral pyramid, SSDT dilated attention, and MIMFormer [10], AMSF-Net, ASMNet and DPFormer multiscale spectral–spatial designs. Recent work also studies state-space models, mixture-of-experts routing and modality interaction. General image-fusion frameworks such as U2Net [11] and OTPNet [12] have also been applied to this task. AMHF-Net [13] handles an arbitrary number of spectral bands. SA-RFM [14] performs non-rigid registration and fusion jointly with Riemannian flow matching.

C. Model-guided and unfolded networks

Algorithm unrolling turns the iterations of an optimization method into network stages. MHF-Net [7], MoG-DCN and DHIF-Net [8] unfold fusion-specific solvers, MHF-Net and DHIF-Net with operators learned as convolution layers, DUNET adds transformer priors for unregistered inputs, IR&ArF unfolds the fusion of unregistered images at arbitrary resolution with a learned registration, HMPNet extends unfolding to three modalities, and an information-aware transformer-based unfolding network (ITU-Net), a mutual-guidance unfolding network (SMGU) [15], the perceptive spectral transformer unfolding network (PSTUN), which also handles arbitrary scales, and the cross-domain-aware unfolding transformer CaFormer [16] strengthen the learned proximal steps with transformer priors. ExactFuse belongs to this family, and none of its prior's components is new: passing features from one stage to the next follows memory-augmented unfolding, spectral-wise attention inside an unfolded network appears in DAUHST, and the transposed attention is that of Restormer [17]. A single unfolded network that takes the blur kernel as an input, and therefore handles kernels it was not trained on, was shown for single-image super-resolution by USRNet [18].

D. Boundary conditions

In classical deconvolution, an image border handled differently by the model and by the data causes ringing that spreads inward, and reflexive (Neumann) boundary conditions, i.e., reflect padding, are a standard remedy. In iterative reconstruction, a back-projection that is not the exact adjoint of the forward model changes the point to which the iteration converges, and in convolutional networks zero padding is itself a source of spatial bias near the border. In unfolded fusion networks the operators are either learned or implemented with the zero padding of convolution layers, while benchmark data are simulated with other border rules; to our knowledge, the effect of this mismatch has not been examined. Our contribution is on this side: an unfolded network whose operators, border included, are exactly those of the simulation, and an ablation that separates the exact operator from a zero-border and a learned operator under the same budget.

E. Generative fusion

Diffusion models have been conditioned on the two observations (DDIF), constrained by the observation model during sampling (CDPS), restricted to a learned subspace (ARGS-Diff) or combined with differential attention in a two-stage conditioning scheme [6]. Null-space formulations make the data-consistency part of the reconstruction exact. The two-stage conditional diffusion model of Chen *et al.* [6] holds the best published Chikusei $\times 4$ results in our survey and is our main reference; its comparison also includes the correlation-aware network CLSNet [19] and the dual cross transformer (DCTransformer) [20], denoted DCT in their table.

III. Problem Formulation and Analysis

A. Observation model

Let $\mathbf{X} \in \mathbb{R}^{B\times H\times W}$ be the HR-HSI with B bands ($B = 128$ for Chikusei and 92 for Pavia Centre). The two observations are

$$\mathbf{Y}_{\mathrm{H}} = \mathcal{D}_\theta \mathbf{X} + \mathbf{N}_{\mathrm{H}}, \qquad \mathbf{Y}_{\mathrm{M}} = \mathcal{R}_\phi \mathbf{X} + \mathbf{N}_{\mathrm{M}}, \tag{1}$$

where $\mathcal{D}_\theta = \mathcal{S}\,\mathcal{K}_\theta\,\mathcal{P}$ applies, band by band, reflect padding $\mathcal{P}$ of width p, a valid correlation $\mathcal{K}_\theta$ with a normalized $(2p+1)\times(2p+1)$ Gaussian kernel of standard deviation θ, and decimation $\mathcal{S}$ that keeps every s-th pixel starting at index zero, so $\mathbf{Y}_\mathrm{H} \in \mathbb{R}^{B\times H/s\times W/s}$; $\mathcal{R}_\phi$ applies to every pixel the M spectral responses of the multispectral sensor, shifted along wavelength by ϕ nm; and $\mathbf{N}_\mathrm{H}$, $\mathbf{N}_\mathrm{M}$ are independent white Gaussian noise of standard deviation ς_H and ς_M. The benchmark protocol is the special case $\theta=\sigma_0$, $\phi=0$ without noise, with $s=4$, $p=3$ and $\sigma_0=2$, the 7×7 blur of [6]; we write $\mathcal{D}=\mathcal{D}_{\sigma_0}$ and $\mathcal{R}=\mathcal{R}_0$ for the nominal operators and keep s and p general where the analysis does not depend on them. The spectral operator $\mathcal{R}\in\mathbb{R}^{M\times B}$ integrates the spectrum with the $M=8$ WorldView-2 band responses for Chikusei and with the $M=4$ IKONOS band responses for Pavia Centre. Because the manufacturers publish the responses only as plots together with band edges, we model each band as a flat top with logistic flanks (4 nm roll-off) that pass exactly through the published edges, at 5% response for WorldView-2, whose edges are tabulated at that level, and at 50% for the IKONOS band limits (blue 445–516, green 506–595, red 632–698 and near infrared 757–853 nm), and normalize every band to unit sum.

Exact adjoint. Every gradient step on $\|\mathcal{D}\mathbf{X}-\mathbf{Y}_\mathrm{H}\|^2$ requires $\mathcal{D}^\top=\mathcal{P}^\top\mathcal{K}^\top\mathcal{S}^\top$. $\mathcal{S}^\top$ inserts zeros and $\mathcal{K}^\top$ is a full convolution; together they are a strided transposed convolution with an output padding of $s-1$. The adjoint of reflect padding folds the padded margins back onto the pixels they were copied from. For a padded row index $j\in[0,H+2p)$ the source row is $p-j$ for $j<p$, $j-p$ in the interior and $2H-2+p-j$ in the bottom margin, hence

$$(\mathcal{P}^\top\mathbf{g})_i = \mathbf{g}_{i+p} + [1\le i\le p]\,\mathbf{g}_{p-i} + [H-1-p\le i\le H-2]\,\mathbf{g}_{2H-2+p-i}, \tag{2}$$

applied to rows and then to columns. With (2), $\mathcal{D}$ in the network is exactly the simulation operator, and the inner-product test $\langle\mathcal{D}\mathbf{X},\mathbf{Y}\rangle=\langle\mathbf{X},\mathcal{D}^\top\mathbf{Y}\rangle$ holds to a relative error below 10^{-13} in double precision at every scale factor. $\mathcal{R}^\top$ is the transpose of the response matrix.

B. MAP estimation and its unfolding

For given operator parameters and noise levels, and a prior $p(\mathbf{X})\propto\exp(-\lambda\Phi(\mathbf{X}))$, the maximum a posteriori estimate is

$$\hat{\mathbf{X}}=\arg\min_{\mathbf{X}}\ \tfrac{1}{2\varsigma_\mathrm{H}^2}\|\mathcal{D}_\theta\mathbf{X}-\mathbf{Y}_\mathrm{H}\|^2+\tfrac{1}{2\varsigma_\mathrm{M}^2}\|\mathcal{R}_\phi\mathbf{X}-\mathbf{Y}_\mathrm{M}\|^2+\lambda\Phi(\mathbf{X}). \tag{3}$$

Half-quadratic splitting [18] introduces an auxiliary image $\mathbf{Z}$ coupled to $\mathbf{X}$ by a penalty $\frac{\mu_k}{2}\|\mathbf{Z}-\mathbf{X}\|^2$ and alternates, for $k=1,\dots,K$, a data step and a prior step:

$$\mathbf{Z}_k=\arg\min_{\mathbf{Z}}\ \tfrac{1}{2\varsigma_\mathrm{H}^2}\|\mathcal{D}_\theta\mathbf{Z}-\mathbf{Y}_\mathrm{H}\|^2+\tfrac{1}{2\varsigma_\mathrm{M}^2}\|\mathcal{R}_\phi\mathbf{Z}-\mathbf{Y}_\mathrm{M}\|^2 + \tfrac{\mu_k}{2}\|\mathbf{Z}-\mathbf{X}_{k-1}\|^2, \tag{4}$$

$$\mathbf{X}_k=\mathrm{prox}_{(\lambda/\mu_k)\Phi}(\mathbf{Z}_k). \tag{5}$$

With $\alpha_k=1/(\mu_k\varsigma_\mathrm{H}^2)$ and $\beta_k=1/(\mu_k\varsigma_\mathrm{M}^2)$, the data step (4) is the symmetric positive definite linear system

$$(\mathbf{I}+\alpha_k\mathcal{D}_\theta^\top\mathcal{D}_\theta+\beta_k\mathcal{R}_\phi^\top\mathcal{R}_\phi)\,\mathbf{Z}_k=\mathbf{X}_{k-1}+\alpha_k\mathcal{D}_\theta^\top\mathbf{Y}_\mathrm{H}+\beta_k\mathcal{R}_\phi^\top\mathbf{Y}_\mathrm{M}, \tag{6}$$

which conjugate gradients solve with the exact operators and adjoints alone, and a single gradient step from $\mathbf{X}_{k-1}$ gives its linearization

$$\mathbf{Z}_k\approx\mathbf{X}_{k-1}-\alpha_k\mathcal{D}_\theta^\top(\mathcal{D}_\theta\mathbf{X}_{k-1}-\mathbf{Y}_\mathrm{H})-\beta_k\mathcal{R}_\phi^\top(\mathcal{R}_\phi\mathbf{X}_{k-1}-\mathbf{Y}_\mathrm{M}). \tag{7}$$

An unfolded network replaces the proximal map (5) by a learned module. Three properties of this derivation shape the network of Section IV. First, α_k and β_k are not independent design parameters: their ratio $\beta_k/\alpha_k=\varsigma_\mathrm{H}^2/\varsigma_\mathrm{M}^2$ weights the two observations by their noise. Second, the proximal map depends on λ/μ_k and hence on the noise level, so a learned prior should be told it. Third, the operators enter only through $\mathcal{D}_\theta$, $\mathcal{R}_\phi$ and their adjoints in (6), so operators known or estimated at test time can be used without retraining. Theorems 1 and 3 state when the operators can be identified from the observations and what an inexact operator costs.

C. Identifiability of the operators

Theorem 1 (Commutation, identifiability and stability): Let $\mathcal{D}_\theta$ apply the same padding, blur and decimation to every band and $\mathcal{R}_\phi$ the same matrix $\mathbf{r}_\phi\in\mathbb{R}^{M\times B}$ to every pixel, with kernels and responses that sum to one, and let $\mathbf{r}(\theta,\phi)=\mathcal{D}_\theta\mathbf{Y}_\mathrm{M}-\mathcal{R}_\phi\mathbf{Y}_\mathrm{H}$. (a) $\mathcal{D}_\theta\mathcal{R}_{\phi'}=\mathcal{R}_{\phi'}\mathcal{D}_\theta$ for all θ,ϕ', with $\mathcal{D}_\theta$ acting on M or B bands. Hence, for noise-free observations generated with (θ^*,ϕ^*), the criterion

$$J(\theta,\phi)=\|\mathbf{r}(\theta,\phi)\|^2=\|\mathcal{D}_\theta\mathbf{Y}_\mathrm{M}-\mathcal{R}_\phi\mathbf{Y}_\mathrm{H}\|^2 \tag{8}$$

vanishes at (θ^*,ϕ^*). (b) $J(\theta,\phi)=0$ if and only if $\mathbf{X}$ lies in the null space of $\mathcal{C}_{\theta\phi}=\mathcal{D}_\theta\mathcal{R}_{\phi^*}-\mathcal{R}_\phi\mathcal{D}_{\theta^*}$, and $\mathcal{C}_{\theta\phi}\neq 0$ unless $\mathcal{D}_\theta=\mathcal{D}_{\theta^*}$ and $\mathcal{R}_\phi=\mathcal{R}_{\phi^*}$. Two different configurations are therefore indistinguishable only for images in a proper subspace, for example spatially constant images (any blur and grid offset) and images whose bands are all equal (any response shift); on a finite set of injectively parametrized candidates, (θ^*,ϕ^*) is the unique zero of J for almost every $\mathbf{X}$. (c) If the Jacobian $\mathbf{G}$ of $\mathbf{r}$ with respect to (θ,ϕ) has full column rank at (θ^*,ϕ^*), this point is an isolated zero of J over the continuous parameters. With observation noise $\mathbf{N}_\mathrm{H},\mathbf{N}_\mathrm{M}$, the minimizer of J near (θ^*,ϕ^*) is, to first order, $(\theta^*,\phi^*)-\mathbf{G}^+\mathbf{n}$ with $\mathbf{n}=\mathcal{D}_{\theta^*}\mathbf{N}_\mathrm{M}-\mathcal{R}_{\phi^*}\mathbf{N}_\mathrm{H}$, so its error is at most $\|\mathbf{n}\|/\sigma_{\min}(\mathbf{G})$ and its covariance is $\mathbf{G}^+\mathbf{\Sigma}_\mathbf{n}\mathbf{G}^{+\top}$.

Proof: (a) On vectorized images, $\mathcal{D}_\theta=\mathbf{I}_B\otimes\bar{\mathcal{D}}_\theta$ and $\mathcal{R}_\phi=\mathbf{r}_\phi\otimes\mathbf{I}_N$, with $\bar{\mathcal{D}}_\theta$ the single-band spatial operator and $N=HW$; by the mixed-product rule, $(\mathbf{I}_M\otimes\bar{\mathcal{D}}_\theta)(\mathbf{r}_\phi\otimes\mathbf{I}_N)=\mathbf{r}_\phi\otimes\bar{\mathcal{D}}_\theta=(\mathbf{r}_\phi\otimes\mathbf{I}_{N/s^2})(\mathbf{I}_B\otimes\bar{\mathcal{D}}_\theta)$. With $\mathbf{Y}_\mathrm{M}=\mathcal{R}_{\phi^*}\mathbf{X}$ and $\mathbf{Y}_\mathrm{H}=\mathcal{D}_{\theta^*}\mathbf{X}$, $\mathcal{D}_{\theta^*}\mathbf{Y}_\mathrm{M}=\mathcal{R}_{\phi^*}\mathcal{D}_{\theta^*}\mathbf{X}=\mathcal{R}_{\phi^*}\mathbf{Y}_\mathrm{H}$. (b) By (a), $\mathbf{r}(\theta,\phi)=\mathcal{C}_{\theta\phi}\mathbf{X}$ with $\mathcal{C}_{\theta\phi}=\mathbf{r}_{\phi^*}\otimes\bar{\mathcal{D}}_\theta-\mathbf{r}_\phi\otimes\bar{\mathcal{D}}_{\theta^*}$. Suppose $\mathcal{C}_{\theta\phi}=\mathbf{0}$ and apply it to $\mathbf{X}=\mathbf{e}_b\otimes\mathbf{u}$, a single band b with spatial pattern $\mathbf{u}$: $(\mathbf{r}_{\phi^*})_{:b}\otimes\bar{\mathcal{D}}_\theta\mathbf{u}=(\mathbf{r}_\phi)_{:b}\otimes\bar{\mathcal{D}}_{\theta^*}\mathbf{u}$ for all $\mathbf{u}$. If $\bar{\mathcal{D}}_\theta\neq\bar{\mathcal{D}}_{\theta^*}$, then $\bar{\mathcal{D}}_\theta\neq c\,\bar{\mathcal{D}}_{\theta^*}$ for every scalar c: both kernels sum to one and reflect padding preserves constants, so both operators map a constant image to the same constant, which forces $c=1$. Two linear maps of rank greater than one that are not scalar multiples of each other map some $\mathbf{u}$ to linearly independent images, and for this $\mathbf{u}$ the identity forces $(\mathbf{r}_{\phi^*})_{:b}=(\mathbf{r}_\phi)_{:b}=\mathbf{0}$, impossible for responses that sum to one. If $\bar{\mathcal{D}}_\theta=\bar{\mathcal{D}}_{\theta^*}$, the identity gives $(\mathbf{r}_{\phi^*})_{:b}=(\mathbf{r}_\phi)_{:b}$ for every b. Hence $\mathcal{C}_{\theta\phi}\neq\mathbf{0}$

unless both operators coincide, and its null space is a proper subspace of measure zero; so is the finite union of these null spaces over $\mathcal{G}$. For a spatially constant image $\mathbf{X} = \mathbf{s} \otimes \mathbf{1}$, $\bar{\mathcal{D}}_\theta \mathbf{1} = \bar{\mathcal{D}}_{\theta^*} \mathbf{1}$ gives $\mathcal{C}_{\theta\phi^*}\mathbf{X} = \mathbf{0}$ for every θ; for an image with equal bands $\mathbf{X} = \mathbf{1} \otimes \mathbf{u}$, $\mathbf{r}_\phi \mathbf{1} = \mathbf{1}$ gives $\mathcal{C}_{\theta^*\phi}\mathbf{X} = \mathbf{0}$ for every ϕ. (c) Write $\vartheta = (\theta, \phi)$ and $F(\vartheta, \mathbf{N}) = \mathbf{G}(\vartheta, \mathbf{N})^\top \mathbf{r}(\vartheta, \mathbf{N})$, the gradient of $J/2$, where the noise enters as $\mathbf{r}(\vartheta, \mathbf{N}) = \mathbf{r}(\vartheta, \mathbf{0}) + \mathcal{D}_\theta \mathbf{N}_\mathrm{M} - \mathcal{R}_\phi \mathbf{N}_\mathrm{H}$. At $(\vartheta^*, \mathbf{0})$, $\mathbf{r} = \mathbf{0}$, so $\partial F / \partial \vartheta = \mathbf{G}^\top \mathbf{G}$, which is invertible. The implicit function theorem gives a unique smooth stationary point $\hat{\vartheta}(\mathbf{N})$ near ϑ^* with $\hat{\vartheta} - \vartheta^* = -\mathbf{G}^+ \mathbf{n} + O(\|\mathbf{N}\|^2)$, and $\|\mathbf{G}^+ \mathbf{n}\| \le \|\mathbf{n}\| / \sigma_{\min}(\mathbf{G})$. The zero is isolated because $\|\mathbf{r}(\vartheta, \mathbf{0})\| \ge \sigma_{\min}(\mathbf{G}) \|\vartheta - \vartheta^*\| - O(\|\vartheta - \vartheta^*\|^2)$. ∎

We therefore estimate (θ, ϕ) for each test pair by minimizing (8) on a grid (steps of $0.025\,\sigma_0$ and 1 nm), using only the two observations; HySure uses the same relation to estimate the responses of a model-based method. With noise, $\mathbb{E}\, J(\theta, \phi) = \|\mathbf{E}_{\theta\phi}\mathbf{X}\|^2 + \frac{N}{s^2}\big(M\varsigma_\mathrm{M}^2 \|\mathbf{k}_\theta\|^2 + \varsigma_\mathrm{H}^2 \|\mathbf{r}_\phi\|_F^2\big)$ up to border terms, where $\mathbf{k}_\theta$ is the blur kernel. The noise term decreases with θ and barely depends on ϕ, so noise biases $\hat{\theta}$ upward by an amount that vanishes as the SNR grows. Theorem 1 also requires the family to contain the true operator; Section VI-H tests anisotropic and defocus blurs and misregistered observations, which it does not.

D. Separable structure: a closed-form data step and the fusion null space

A Gaussian kernel, isotropic or aligned with the image axes, is separable, and reflect and zero padding treat rows and columns independently; decimation does too. The blur–decimation operator therefore factors exactly, border included: for the cube $\mathbf{X} \in \mathbb{R}^{B \times H \times W}$,

$$\mathcal{D}_\theta \mathbf{X} = \mathbf{X} \times_2 \mathbf{D}_\mathrm{r} \times_3 \mathbf{D}_\mathrm{c}, \qquad \mathcal{R}_\phi \mathbf{X} = \mathbf{X} \times_1 \mathbf{r}_\phi, \tag{9}$$

where $\times_n$ is the mode-n product and $\mathbf{D}_\mathrm{r} \in \mathbb{R}^{h \times H}$ and $\mathbf{D}_\mathrm{c} \in \mathbb{R}^{w \times W}$ pad, blur and decimate a single axis. Separable operators of this form are used by tensor-factorization fusion methods to identify low-rank models, under periodic borders; here they yield an exact, closed-form data step and an explicit description of what fusion cannot observe.

Theorem 2 (Diagonalization and null space): Let $\mathbf{D}_\mathrm{r}^\top \mathbf{D}_\mathrm{r} = \mathbf{U}_\mathrm{r} \mathrm{diag}(\boldsymbol{\lambda}^\mathrm{r}) \mathbf{U}_\mathrm{r}^\top$, $\mathbf{D}_\mathrm{c}^\top \mathbf{D}_\mathrm{c} = \mathbf{U}_\mathrm{c} \mathrm{diag}(\boldsymbol{\lambda}^\mathrm{c}) \mathbf{U}_\mathrm{c}^\top$ and $\mathbf{r}_\phi^\top \mathbf{r}_\phi = \mathbf{V} \mathrm{diag}(\boldsymbol{\nu}) \mathbf{V}^\top$ be eigendecompositions, and $\mathcal{Q}(\cdot) = (\cdot) \times_1 \mathbf{V} \times_2 \mathbf{U}_\mathrm{r} \times_3 \mathbf{U}_\mathrm{c}$. (a) $\mathcal{D}_\theta^\top \mathcal{D}_\theta$ and $\mathcal{R}_\phi^\top \mathcal{R}_\phi$ are diagonal in the orthonormal basis $\mathcal{Q}$, with eigenvalues $\lambda_i^\mathrm{r} \lambda_j^\mathrm{c}$ for every band and ν_b for every pixel. (b) The exact data step (6) is

$$\mathbf{Z}_k = \mathcal{Q}\Big(\mathcal{Q}^{-1}(\mathbf{C}_k) \oslash \big[1 + \alpha_k \lambda_i^\mathrm{r} \lambda_j^\mathrm{c} + \beta_k \nu_b\big]_{b,i,j}\Big), \tag{10}$$

with $\mathbf{C}_k$ the right-hand side of (6) and $\oslash$ the element-wise division: three mode products in each direction, $O(BHW(B + H + W))$ operations, no iteration and no periodic-border assumption. (c) The null space of the joint operator, $\mathcal{N}(\mathcal{D}_\theta) \cap \mathcal{N}(\mathcal{R}_\phi)$, is spanned by the basis tensors with $\nu_b = 0$ and $\lambda_i^\mathrm{r} \lambda_j^\mathrm{c} = 0$; its orthogonal projector is $(\mathbf{I} - \mathbf{P}_\mathrm{s}) \otimes (\mathbf{I} - \mathbf{P}_\mathrm{r} \otimes \mathbf{P}_\mathrm{c})$ in mode order, with $\mathbf{P}_\mathrm{s}$, $\mathbf{P}_\mathrm{r}$, $\mathbf{P}_\mathrm{c}$ the projectors onto the row spaces of $\mathbf{r}_\phi$, $\mathbf{D}_\mathrm{r}$, $\mathbf{D}_\mathrm{c}$. If these have full row rank, its dimension is $(B - M)(HW - hw)$, a fraction $(1 - M/B)(1 - 1/s^2)$ of all unknowns.

Proof: (a) By (9), $\mathcal{D}_\theta^\top \mathcal{D}_\theta$ applies $\mathbf{D}_\mathrm{r}^\top \mathbf{D}_\mathrm{r}$ and $\mathbf{D}_\mathrm{c}^\top \mathbf{D}_\mathrm{c}$ to the two spatial modes and the identity to the spectral mode, and $\mathcal{R}_\phi^\top \mathcal{R}_\phi$ applies $\mathbf{r}_\phi^\top \mathbf{r}_\phi$ to the spectral mode alone; each factor is diagonalized by the corresponding factor of $\mathcal{Q}$, and an identity factor by any orthonormal basis. (b) In this basis the matrix of (6) is diagonal with entries $1 + \alpha_k \lambda_i^\mathrm{r} \lambda_j^\mathrm{c} + \beta_k \nu_b > 0$. (c) $\mathcal{N}(\mathcal{D}_\theta) \cap \mathcal{N}(\mathcal{R}_\phi) = \mathcal{N}(\mathcal{D}_\theta^\top \mathcal{D}_\theta + \mathcal{R}_\phi^\top \mathcal{R}_\phi)$, whose eigenvalues $\lambda_i^\mathrm{r} \lambda_j^\mathrm{c} + \nu_b$ are sums of nonnegative terms and vanish only when both terms do. The basis tensors with $\lambda_i^\mathrm{r} \lambda_j^\mathrm{c} > 0$ span the row space of $\mathbf{D}_\mathrm{r} \otimes \mathbf{D}_\mathrm{c}$, of dimension hw, and those with $\nu_b > 0$ the row space of $\mathbf{r}_\phi$, of dimension M, which gives the projector and the dimension. ∎

The closest closed form, that of FUSE [2], solves the same system by Fourier transforms under periodic borders; (10) holds for the border rule of the simulation, which Theorem 3 shows to matter, and is a differentiable layer whose operators may change from one image to the next, at the cost of eigendecompositions of an $H \times H$, a $W \times W$ and a $B \times B$ matrix. Part (c) states what any fusion method has to infer from a prior: the high spatial frequencies of the spectral directions that the multispectral sensor does not see, 88% of the unknowns for Chikusei and 90% for Pavia Centre at $\times 4$; every other component is determined by the observations. Without the HR-MSI the null space would hold all high spatial frequencies, $1 - 1/s^2$ of the unknowns, as in single-image super-resolution. The HR-MSI observes them in only M spectral directions, but spatial detail is so strongly shared across bands that these directions carry 84% of their energy on Chikusei and 77% on Pavia Centre: the minimum-norm image consistent with both observations reaches 41.3 dB on Chikusei, against 33.5 dB for the LR-HSI alone. Section VI-C uses this to split the error of every method into a part that contradicts the observations and a part in this null space.

E. The cost of an inexact operator

Stack the two data terms as $\mathbf{A} = [\mathcal{D}_\theta; \mathcal{R}_\phi]$, $\mathbf{Y} = [\mathbf{Y}_\mathrm{H}; \mathbf{Y}_\mathrm{M}]$ and $\mathbf{W}_k = \mathrm{diag}(\alpha_k \mathbf{I}, \beta_k \mathbf{I})$, and let $\tilde{\mathbf{A}}$ be the operator a network uses in place of $\mathbf{A}$: one with a zero border, a learned one, or one with wrong parameters. Write $T(\mathbf{V}) = (\mathbf{I} + \tilde{\mathbf{A}}^\top \mathbf{W}_k \tilde{\mathbf{A}})^{-1} (\mathbf{V} + \tilde{\mathbf{A}}^\top \mathbf{W}_k \mathbf{Y})$ for the exact data step (6) applied to an estimate $\mathbf{V}$, and $G(\mathbf{V}) = \mathbf{V} - \tilde{\mathbf{A}}^\top \mathbf{W}_k (\tilde{\mathbf{A}} \mathbf{V} - \mathbf{Y})$ for its linearization (7).

Theorem 3 (Fixed point and operator mismatch): Let $\mathbf{Y} = \mathbf{A}\mathbf{X}^*$. (a) If $\tilde{\mathbf{A}} = \mathbf{A}$, then $T(\mathbf{X}^*) = G(\mathbf{X}^*) = \mathbf{X}^*$ for all $\alpha_k, \beta_k > 0$. (b) For every $\tilde{\mathbf{A}}$,

$$\|T(\mathbf{X}^*) - \mathbf{X}^*\| \le \tfrac{1}{2} \|\mathbf{W}_k^{1/2} (\mathbf{A} - \tilde{\mathbf{A}}) \mathbf{X}^*\|, \tag{11}$$

$$\|G(\mathbf{X}^*) - \mathbf{X}^*\| \le \|\mathbf{W}_k^{1/2} \tilde{\mathbf{A}}\| \, \|\mathbf{W}_k^{1/2} (\mathbf{A} - \tilde{\mathbf{A}}) \mathbf{X}^*\|. \tag{12}$$

(c) Let stage k of a network be $\mathbf{X}_k = (\mathcal{I} + \mathcal{P}_k)(T_k(\mathbf{X}_{k-1}))$, where the exact data step T_k uses $\mathbf{W}_k$ and the residual prior satisfies $\mathcal{P}_k(\mathbf{X}^*) = \mathbf{0}$ and makes $\mathcal{I} + \mathcal{P}_k$ L_k-Lipschitz. Started at $\mathbf{X}_0 = \mathbf{X}^*$, the network returns

$$\|\mathbf{X}_K - \mathbf{X}^*\| \le \tfrac{1}{2} \sum_{k=1}^{K} \Big(\prod_{j=k}^{K} L_j \Big) \|\mathbf{W}_k^{1/2} (\mathbf{A} - \tilde{\mathbf{A}}) \mathbf{X}^*\|. \tag{13}$$

Proof: (a) $\tilde{\mathbf{A}}\mathbf{X}^* - \mathbf{Y} = \mathbf{0}$ makes both updates vanish. (b) With $\mathbf{B} = \mathbf{W}_k^{1/2} \tilde{\mathbf{A}}$, a direct computation gives

$T(\mathbf{X}^*) - \mathbf{X}^* = (\mathbf{I} + \mathbf{B}^\top\mathbf{B})^{-1}\mathbf{B}^\top\mathbf{W}_k^{1/2}(\mathbf{A} - \tilde{\mathbf{A}})\mathbf{X}^*$; in the singular value decomposition of $\mathbf{B}$, $(\mathbf{I} + \mathbf{B}^\top\mathbf{B})^{-1}\mathbf{B}^\top$ has singular values $b_i/(1 + b_i^2) \le \frac{1}{2}$. Similarly $G(\mathbf{X}^*) - \mathbf{X}^* = \mathbf{B}^\top\mathbf{W}_k^{1/2}(\mathbf{A} - \tilde{\mathbf{A}})\mathbf{X}^*$, whose norm is at most $\|\mathbf{B}\|$ times the last factor. (c) T_k is nonexpansive, since $\|T_k(\mathbf{V}) - T_k(\mathbf{V}')\| = \|(\mathbf{I} + \mathbf{B}^\top\mathbf{B})^{-1}(\mathbf{V} - \mathbf{V}')\| \le \|\mathbf{V} - \mathbf{V}'\|$. With $e_k = \|\mathbf{X}_k - \mathbf{X}^*\|$ and $(\mathcal{I} + \mathcal{P}_k)(\mathbf{X}^*) = \mathbf{X}^*$, $e_k \le L_k\|T_k(\mathbf{X}_{k-1}) - \mathbf{X}^*\| \le L_k\big(\|T_k(\mathbf{X}_{k-1}) - T_k(\mathbf{X}^*)\| + \|T_k(\mathbf{X}^*) - \mathbf{X}^*\|\big) \le L_k(e_{k-1} + b_k)$, where b_k is the bound (11); unrolling from $e_0 = 0$ gives (13). ∎

Corollary 1 (From observation noise to reconstruction): Under Theorem 1(c), the exact data step with the estimated operator $\tilde{\mathbf{A}} = \mathbf{A}_{\hat{\theta}\hat{\phi}}$ moves the ground truth by at most $\|\mathbf{W}_k^{1/2}\mathbf{A}'\|\,\|\mathbf{n}\|/(2\,\sigma_{\min}(\mathbf{G}))$ to first order, where $\mathbf{A}'$ is the Jacobian of $(\theta, \phi) \mapsto \mathbf{A}_{\theta\phi}\mathbf{X}^*$.

Proof: $(\mathbf{A} - \tilde{\mathbf{A}})\mathbf{X}^* = -\mathbf{A}'(\hat{\vartheta} - \vartheta^*) + O(\|\hat{\vartheta} - \vartheta^*\|^2)$; combine (11) with Theorem 1(c). ∎

Together, Theorems 1 and 3 carry the observation noise through the operator estimate to the reconstruction: the conditioning $\sigma_{\min}(\mathbf{G})$ of the estimation problem decides how much noise moves the operator, and the exact data step bounds how much the operator error moves the reconstruction, whatever the learned weights. Section VI-H measures both on the test images. Theorem 3 separates the two roles of the operator. With the exact operator, the ground truth is a fixed point of every data step, so the priors only have to remove the reconstruction error. With an inexact one, every data step moves the ground truth, and the priors must also undo this displacement; because it depends on the image, it cannot be removed by a fixed correction. The exact proximal step bounds the displacement by $\frac{1}{2}\|\mathbf{W}_k^{1/2}(\mathbf{A} - \tilde{\mathbf{A}})\mathbf{X}^*\|$ whatever the learned weights, whereas for the gradient step the bound grows with $\|\mathbf{W}_k^{1/2}\tilde{\mathbf{A}}\|$, which training may make large, and the step is not nonexpansive, so (13) has no counterpart for it. The theorem also explains why an exact operator can be replaced at test time: for data generated by $\mathbf{A}'$, giving the network $\tilde{\mathbf{A}} = \mathbf{A}'$ restores (a), which a learned operator cannot do.

Corollary 2 (Border rule): If $\tilde{\mathbf{A}}$ differs from $\mathbf{A}$ only in the border rule of $\mathcal{D}_\theta$, e.g., zero instead of reflect padding, $(\mathbf{A} - \tilde{\mathbf{A}})\mathbf{X}^*$ is nonzero only on the LR pixels whose $(2p + 1) \times (2p+1)$ kernel support reaches the padding, i.e., within $\lceil p/s \rceil$ pixels of the LR border.

With decimation from index zero and $p < s$, only the first LR row and column are affected, a fraction $(h + w - 1)/(hw)$ of an $h \times w$ LR image: 3% of a 64×64 image but 12% of the 16×16 LR patches on which networks are trained. The displacement (11) is then small, which matches the small accuracy gap at the training blur (Section VI-A), whereas a learned operator differs from $\mathbf{A}$ everywhere. In real images the true border rule is unknown, and Corollary 2 bounds what any fixed choice can cost; the main role of exactness is the one stated by Theorem 3: an operator that can be supplied, estimated and replaced. Section VI-H measures the displacement and the test PSNR as functions of the operator residual $\|(\tilde{\mathbf{A}} - \mathbf{A})\mathbf{X}^*\|$.

IV. ExactFuse

A. Unfolded iterations

ExactFuse (Fig. 1) is the HQS iteration of Section III-B unrolled for $K = 3$ stages, starting from the bicubic upsampling $\mathbf{X}_0$ of $\mathbf{Y}_\mathrm{H}$. Stage k (Fig. 1, bottom) first performs the data step (4) with the exact operators. ExactFuse takes the learned spectral step (16) below, the closed form (10) with learned gains; the base model of Section VI-A takes the linearized step

$$\begin{aligned}\mathbf{Z}_k = \mathbf{X}_{k-1} &- s^2\alpha_k\,\mathcal{D}_\theta^\top(\mathcal{D}_\theta\mathbf{X}_{k-1} - \mathbf{Y}_\mathrm{H}) \\ &- \beta_k\,\mathcal{R}_\phi^\top(\mathcal{R}_\phi\mathbf{X}_{k-1} - \mathbf{Y}_\mathrm{M}),\end{aligned} \tag{14}$$

and a kernel that is not separable would take conjugate gradients on (6) (Fig. 2). The weights $\alpha_k = \mathrm{softplus}(a_k)$ and $\beta_k = \mathrm{softplus}(b_k)$ stand for $1/(\mu_k\varsigma_\mathrm{H}^2)$ and $1/(\mu_k\varsigma_\mathrm{M}^2)$ in (6) and are learned per stage; the factor s^2 compensates the $1/s^2$ gain of $\mathcal{D}^\top\mathcal{D}$ on smooth content under $\times s$ decimation, so that both terms start at a comparable scale. The prior step (5) is a learned residual,

$$\begin{aligned}(\Delta_k, \mathbf{F}_k) &= \mathcal{P}_k\big([\mathbf{Z}_k, \mathcal{R}_\phi^+\mathbf{Y}_\mathrm{M}, s^2\mathcal{D}^\top(\mathcal{D}\mathbf{Z}_k - \mathbf{Y}_\mathrm{H}), \tfrac{B}{M}\mathcal{R}_\phi^\top(\mathcal{R}_\phi\mathbf{Z}_k - \mathbf{Y}_\mathrm{M})], \mathbf{F}_{k-1}\big), \\ \mathbf{X}_k &= \mathbf{Z}_k + \Delta_k,\end{aligned} \tag{15}$$

where $\mathcal{R}_\phi^+\mathbf{Y}_\mathrm{M}$ are the minimum-norm spectra in the B hyperspectral bands that reproduce the HR-MSI ($\mathcal{R}_\phi\mathcal{R}_\phi^+ = \mathbf{I}$, so they determine $\mathbf{Y}_\mathrm{M}$ exactly), the last two inputs, the gradients of both data terms at $\mathbf{Z}_k$, tell the prior where the updated estimate still disagrees with each observation, and $\mathbf{F}_{k-1}$ are the final decoder features of the previous stage (cross-stage memory, merged by a 1×1 convolution). Every input lies in the hyperspectral bands, so the prior does not depend on the number or the responses of the multispectral bands; the base model gives it $[\mathbf{Z}_k, \mathbf{Y}_\mathrm{M}, s^2\mathcal{D}^\top(\mathcal{D}\mathbf{Z}_k - \mathbf{Y}_\mathrm{H})]$. The last convolution of every prior is initialized to zero, so training starts from the pure HQS iteration with $\mathcal{P}_k = \mathbf{0}$, which leaves the ground truth unchanged as required by Theorem 3(c). The image path, including all operator applications, stays in 32-bit floating point; only the priors run in 16-bit mixed precision.

Learned spectral data step. In the basis of Theorem 2, the exact step multiplies every component of the data residual by the gain $1/(1 + \alpha_k\lambda_i^\mathrm{r}\lambda_j^\mathrm{c} + \beta_k\nu_b)$, the filter of a white Gaussian prior. Natural images are not white, and noise does not harm all components equally, so we let two small networks modulate these gains:

$$\begin{aligned}\mathbf{Z}_k &= \mathbf{X}_{k-1} + \mathcal{Q}\big(\mathcal{Q}^{-1}(\mathbf{G}_k) \odot \mathbf{\Gamma}_k\big), \\ \Gamma_{k,bij} &= \frac{\exp\big(\ln 2\,\tanh(u_k(\lambda_i^\mathrm{r}\lambda_j^\mathrm{c}, \mathbf{c}) + v_k(\nu_b, \mathbf{c}))\big)}{1 + \alpha_k\lambda_i^\mathrm{r}\lambda_j^\mathrm{c} + \beta_k\nu_b},\end{aligned} \tag{16}$$

where $\mathbf{G}_k = \alpha_k\mathcal{D}_\theta^\top(\mathbf{Y}_\mathrm{H} - \mathcal{D}_\theta\mathbf{X}_{k-1}) + \beta_k\mathcal{R}_\phi^\top(\mathbf{Y}_\mathrm{M} - \mathcal{R}_\phi\mathbf{X}_{k-1})$ is the data residual and u_k, v_k are two-layer perceptrons of the log-eigenvalue and of the observation code $\mathbf{c}$ introduced below, initialized to zero so that training starts from the exact step (10). The gain is separable, so a stage evaluates the networks $HW + B$ times rather than BHW times.

Proposition 1: For all u_k, v_k: (a) with exact operators and noise-free data, $\mathbf{X}^*$ is a fixed point of (16); (b) (16) is

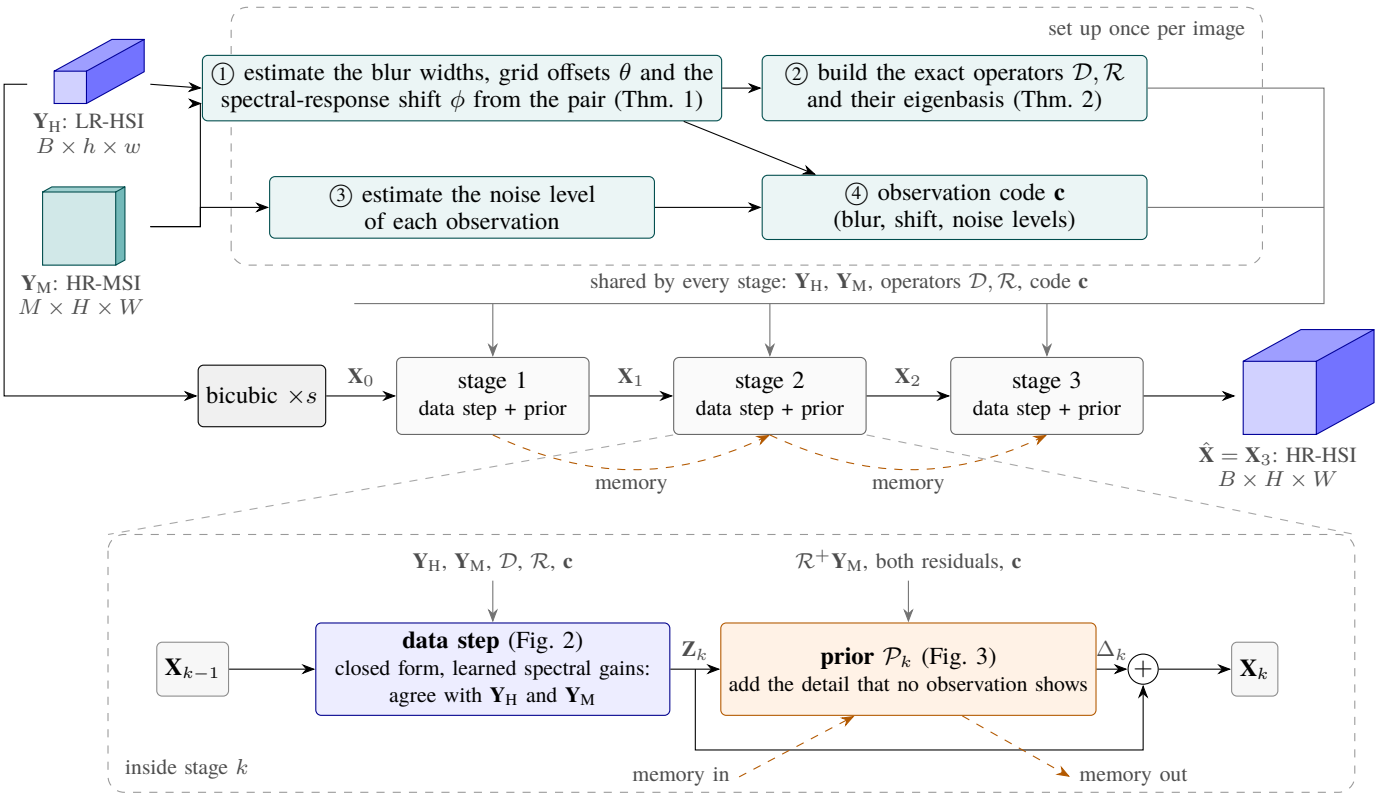


Fig. 1. ExactFuse, the HQS iteration of Section III-B unrolled for three stages. Top: the observation model is set up once per image: the operator parameters are given or estimated from the observation pair (Theorem 1), the exact operators and their eigenbasis are built (Theorem 2), and the noise levels form the observation code **c**. Middle: every stage refines the estimate and passes its decoder features to the next (memory). Bottom: inside a stage, a data step makes the estimate agree with both observations, and a learned prior adds the detail that the observations cannot provide.

nonexpansive; (c) with an inexact operator $\tilde{\mathbf{A}}$, it moves $\mathbf{X}^*$ by at most $\|\mathbf{W}_k^{1/2}(\mathbf{A}-\tilde{\mathbf{A}})\mathbf{X}^*\|$, twice the bound (11) and independent of the weights.

Proof: Let $\varsigma_{bij} = \alpha_k \lambda_i^{\mathrm{r}} \lambda_j^{\mathrm{c}} + \beta_k \nu_b$, the eigenvalues of $\tilde{\mathbf{A}}^\top \mathbf{W}_k \tilde{\mathbf{A}}$, and e^f the numerator of Γ, so $\frac{1}{2} < e^f < 2$. (a) $\mathbf{G}_k = \mathbf{0}$ at $\mathbf{X}^*$. (b) The step maps a difference $\mathbf{V} - \mathbf{V}'$ to $\mathcal{Q}\big((1-\Gamma\varsigma) \odot \mathcal{Q}^{-1}(\mathbf{V}-\mathbf{V}')\big)$, and $0 \le \Gamma\varsigma = e^f\varsigma/(1+\varsigma) < 2$, so every factor lies in $(-1, 1]$. (c) The displacement is $\mathcal{Q}\big(\Gamma \odot \mathcal{Q}^{-1}(\mathbf{B}^\top \mathbf{e})\big)$ with $\mathbf{B} = \mathbf{W}_k^{1/2}\tilde{\mathbf{A}}$ and $\mathbf{e} = \mathbf{W}_k^{1/2}(\mathbf{A}-\tilde{\mathbf{A}})\mathbf{X}^*$; its operator norm is $\max e^f\sqrt{\varsigma}/(1+\varsigma) \le 2 \cdot \frac{1}{2} = 1$. ∎

Theorem 3(a) and (c) therefore hold for the learned step too: it can reweight how much the network trusts each observed component, but it cannot move a correct estimate or amplify errors.

Algorithm 1 ExactFuse inference (unrolled HQS)

Require: $\mathbf{Y}_{\mathrm{H}}$, $\mathbf{Y}_{\mathrm{M}}$; operator parameters (θ, ϕ), given or $\arg\min_{\mathcal{G}} J$ of (8)

build the exact $\mathcal{D}_\theta, \mathcal{D}_\theta^\top, \mathcal{R}_\phi, \mathcal{R}_\phi^\top$ (reflect border, (2))
$\mathbf{X}_0 \leftarrow \mathrm{bicubic}_{\times s}(\mathbf{Y}_{\mathrm{H}})$, $\mathbf{F}_0 \leftarrow \varnothing$
for $k = 1, \dots, K$ **do**
  $\mathbf{Z}_k \leftarrow$ data step: (14), or CG on (6) ▷ fp32
  $\mathbf{r}_k \leftarrow s^2\,\mathcal{D}^\top(\mathcal{D}\mathbf{Z}_k - \mathbf{Y}_{\mathrm{H}})$
  $(\Delta_k, \mathbf{F}_k) \leftarrow \mathcal{P}_k([\mathbf{Z}_k, \mathbf{Y}_{\mathrm{M}}, \mathbf{r}_k], \mathbf{F}_{k-1})$ ▷ fp16
  $\mathbf{X}_k \leftarrow \mathbf{Z}_k + \Delta_k$
end for
return $\mathbf{X}_K$

B. Transposed-attention prior

Each prior (Fig. 3) is a three-level U-Net with widths 48, 96 and 192, two, three and four transformer blocks on the encoder levels and bottleneck, mirrored decoder levels, and two refinement blocks. A block (Fig. 4(a)) follows Restormer [17]: multi-Dconv-head transposed attention (Fig. 5) computes queries, keys and values with 1×1 and depth-wise 3×3 convolutions, normalizes each feature map over space, and forms a $c' \times c'$ attention matrix across channels per head (1, 2 and 4 heads on the three levels) with a learned temperature. Attention is thus computed between spectral feature maps rather than between spatial positions, its cost is linear in the number of pixels, and the receptive field of every block covers the whole image. A gated depth-wise feed-forward network with expansion 2.66 follows (Fig. 4(b)). The three priors hold 3.37–3.38 M parameters each; the whole network has 10.13 M parameters for Chikusei and 9.99 M for Pavia Centre, whose input and output convolutions are smaller.

C. Training objective

With ground truth $\mathbf{X}$, the loss is

$$\mathcal{L} = \|\mathbf{X}_K - \mathbf{X}\|_1 + 0.1 \sum_{k<K} \|\mathbf{X}_k - \mathbf{X}\|_1 + \lambda_{\mathrm{SAM}}\,\mathrm{SAM}(\mathbf{X}_K, \mathbf{X}) + \lambda_{\mathrm{S}}\big(1 - \mathrm{SSIM}(\mathbf{X}_K, \mathbf{X})\big), \tag{17}$$

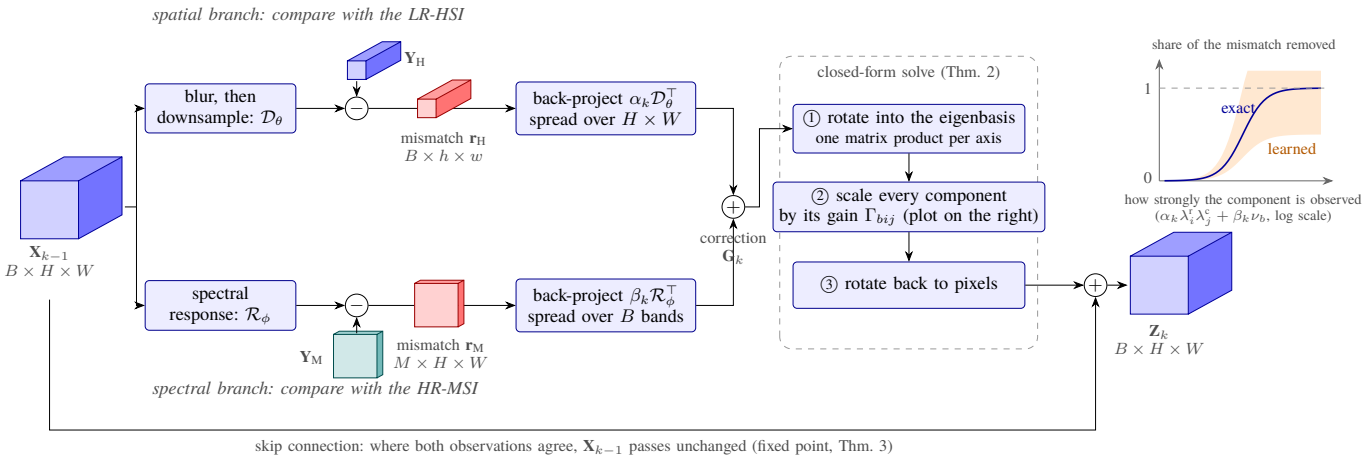


Fig. 2. The data step of stage k, the exact solution of (6). Two physics branches compare the estimate with the observations: the spatial branch blurs and downsamples it and compares it with the LR-HSI, the spectral branch applies the multispectral responses and compares it with the HR-MSI; both mismatches are spread back to the full cube by the exact adjoints. The closed-form solve rotates this correction into the eigenbasis of the separable operators, removes in each component a share $\Gamma_{bij}(\alpha_k \lambda_i^\mathrm{r} \lambda_j^\mathrm{c} + \beta_k \nu_b)$ of the mismatch that grows with how strongly the observations see that component (blue: the exact step (10); orange: the range allowed to the learned gains of Proposition 1), and rotates back. Components that neither observation sees (the null space) are left to the prior.

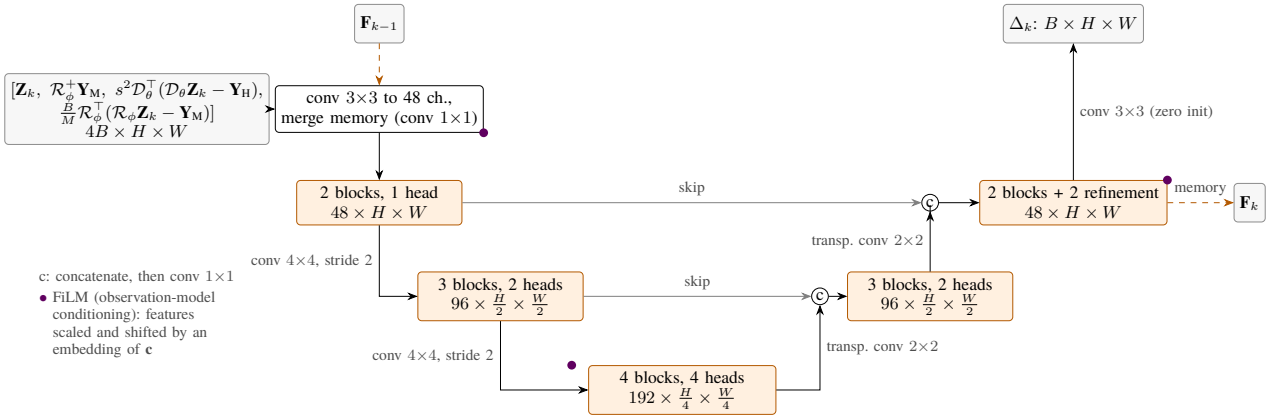


Fig. 3. The prior $\mathcal{P}_k$ of stage k (widths for Chikusei, $B = 128$): a three-level U-Net of transformer blocks (Fig. 4). Its four inputs all lie in the B hyperspectral bands: the data-step output, the minimum-norm spectra $\mathcal{R}_\phi^+ \mathbf{Y}_\mathrm{M}$ that reproduce the HR-MSI, and the residuals of both observations mapped back by the adjoints, so the prior does not depend on the multispectral sensor. The encoder halves the resolution twice with strided convolutions, the decoder restores it with transposed convolutions and merges the encoder features through skip connections; the previous stage's decoder features enter after the input convolution, and the last convolution starts at zero so that training starts from the HQS iteration.

with mean absolute errors, deep supervision of the intermediate stages, the mean spectral angle in radians, and a band-wise Gaussian SSIM whose data range is the maximum of each training patch. We use $\lambda_\mathrm{SAM} = 0.05$ and $\lambda_\mathrm{S} = 0.1$; both terms are computed in 32-bit precision. An exponential moving average of the weights (decay 0.999) is used for validation and testing.

D. Training for a family of observation models

Trained with one fixed operator pair on noise-free data, the priors adapt to one blur, one spectral response and no noise, and Section VI-G shows the consequences: a shifted spectral response or noise is not compensated by giving the network the true operator. Because the operators enter ExactFuse only through its data steps, the same network can instead be trained for a family of observation models. For every batch we draw a blur width $\sigma \sim \mathcal{U}[0.5, 2]\,\sigma_0$, a shift $\delta \sim \mathcal{U}[-12, 12]$ nm of the spectral responses and, with probability 0.5, Gaussian noise on both observations at an SNR drawn from $\mathcal{U}[30, 50]$ dB; the batch is simulated with $(\mathcal{D}_\sigma, \mathcal{R}_\delta)$, and the same operators are used in the data steps. At test time the operators are given, or estimated from the test pair by minimizing (8). Training on randomized degradations is a standard remedy in blind super-resolution, and USRNet [18] passes the blur kernel to the data step of an unfolded single-image super-resolution network. We extend this to the two operators of fusion, with exact borders and with both operators estimated from the observation pair, and compare it with randomized training alone, in which the data steps keep the nominal operator.

Observation-model conditioning. The data steps then know the operator, but the prior does not, and neither knows the noise level, although the HQS derivation (Section III-B) shows that the weights α_k, β_k should follow the noise of the two observations and that the proximal step depends

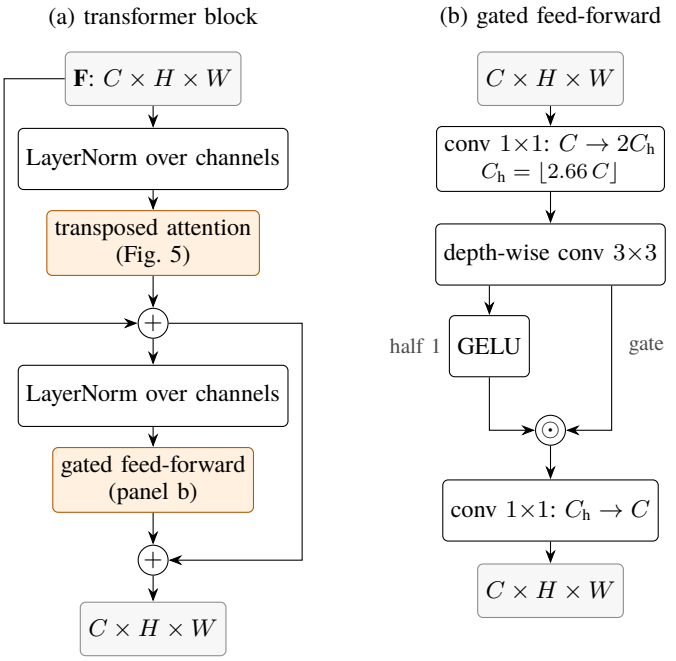


Fig. 4. (a) A transformer block of the prior: layer normalization over channels, transposed attention and a gated feed-forward network, each in a residual branch. (b) The gated feed-forward network: the expanded features are split in two halves, and the GELU of one half gates the other.

on it. We therefore describe each input pair by a code $\mathbf{c} = (\sigma/\sigma_0 - 1,\ \delta/12,\ \log_{10} n_{\mathrm{H}} + 2,\ \log_{10} n_{\mathrm{M}} + 2)$, where (σ, δ) are the parameters of the operator used in the data steps and n is the noise standard deviation of an observation relative to its RMS, estimated from the observation itself as the median absolute diagonal second difference over all bands. The step sizes become $\alpha_k = \mathrm{softplus}(a_k + \mathbf{u}_k^{\top}\mathbf{c})$ and $\beta_k = \mathrm{softplus}(b_k + \mathbf{v}_k^{\top}\mathbf{c})$, so that the two data terms can be weighted by their noise as in a maximum a posteriori estimate, and an embedding of $\mathbf{c}$ (two linear layers) scales and shifts the features of every prior (FiLM) after its input convolution, at the input of its bottleneck and before its output convolution. All added parameters start at zero, so training starts from the unconditioned network; at the nominal operator without noise, $\mathbf{c}$ is nearly constant and the network reduces to ExactFuse.

E. Self-ensemble

At test time we average the prediction for the input with the transposed prediction for the transposed input. Transposition preserves the decimation phase, so both passes see observations generated by exactly the same operator. Flips and rotations do not: flipping a 256-pixel row maps sampled position $4m$ to $255 - 4m$, which has phase three, so we do not use them.

V. Experimental Setup

A. Data and protocol

Chikusei [21] is an airborne scene of 2517×2335 pixels with 128 bands from 363 to 1018 nm. The split sizes follow [6]: eight 256×256 test images, 64 validation patches and 832 training patches of 64×64. Because their locations are not specified there, we fix them as follows. We take the central 2048×2048 crop, which contains no no-data border, and divide it by its global maximum (15,133 digital numbers, DN). The test images are the eight 256×256 tiles of the top 256 rows, the validation patches tile rows 272–400, and training draws random 64×64 crops with dihedral augmentation from rows 416–2048 (800 non-overlapping patch equivalents). Sixteen-row gaps keep the three regions disjoint.

Pavia Centre [22] was recorded by the ROSIS sensor over the centre of Pavia, Italy; the distributed scene has 1096×715 pixels at 1.3 m and 102 bands between 430 and 860 nm. For this dataset, [6] uses the benchmark files released with PSRT [4]: 92 bands, a four-band multispectral image, four 256×256 test images and 142 validation and 1284 training patches of 64×64. These files are distributed through a file-sharing service that we could not access, so we rebuild the benchmark from the original scene to the same specification. We keep bands 11–102 (473–860 nm if the 102 bands are taken as evenly spaced over 430–860 nm), divide by the maximum of the retained bands (8,000 DN), apply the same blur and decimation as for Chikusei, and simulate the HR-MSI with the IKONOS responses. The test images are the four 256×256 tiles of the top-left 512×512 region. The 63 validation patches lie on a 32-pixel grid in rows 528–656, and training draws random 64×64 crops with dihedral augmentation from rows 672–1096 and from the strip to the right of the test region (rows 0–512, columns 528–715), 97 non-overlapping patch equivalents in total. Gaps of 16 pixels keep the regions disjoint. As a check on the rebuilt data, our GSA implementation reaches 34.78 dB on our test images, against 35.35 dB reported for GSA in [6].

For both datasets we follow Wald's protocol: the original image serves as ground truth and both observations are simulated from it, and every image, including every training crop, is degraded on its own with (1). Table I summarizes the two protocols.

Equal-budget benchmark. Reported numbers come from different splits, normalizations and metric code (Section V-B), so for the main comparison we retrain every learned method ourselves under one pipeline: SSR-NET [3], PSRT [4], DCT [20], MIMFormer [10], DHIF-Net [8], U2Net [11], MIMO-SST [5], OTPNet [12], SMGU-Net [15], CLSNet [19], CaFormer [16] and M2DTN [23], built from their authors' code[1]. They cover a CNN, window, dual-cross and multiscale transformers, among them the strongest entries of [6], and model-guided unfolding networks whose blur, decimation and

[1]PSRT: github.com/Deng-shangqi/PSRT (commit 9c2be90); DCT: github.com/qingma2016/DCTransformer (bd75b15); MIMFormer: github.com/meiruni/MIMFormer (0904158); DHIF-Net: github.com/TaoHuang95/DHIF-Net (6434366); SMGU-Net: github.com/yansql/SMGU-Net (115ce77); CLSNet: github.com/HengYang01/CLSNet (6b13371); OTPNet: github.com/tsharry2201/OTPNet (7c17964); U2Net: github.com/PSRben/U2Net (32e214c), tested in overlapping 128×128 tiles because its attention over all pixels of a 256×256 image does not fit a 16 GB GPU; MIMO-SST: github.com/Freelancefangjian/MIMO-SST (06aacc1). Only the numbers of input and output bands and the scale factor are changed; DHIF-Net's code covers $\times8$ and $\times16$, and at $\times4$ we use one of the two stride-4 layers of its $\times16$ operator; OTPNet and U2Net, written for pansharpening, take the HR-MSI in place of the panchromatic image.

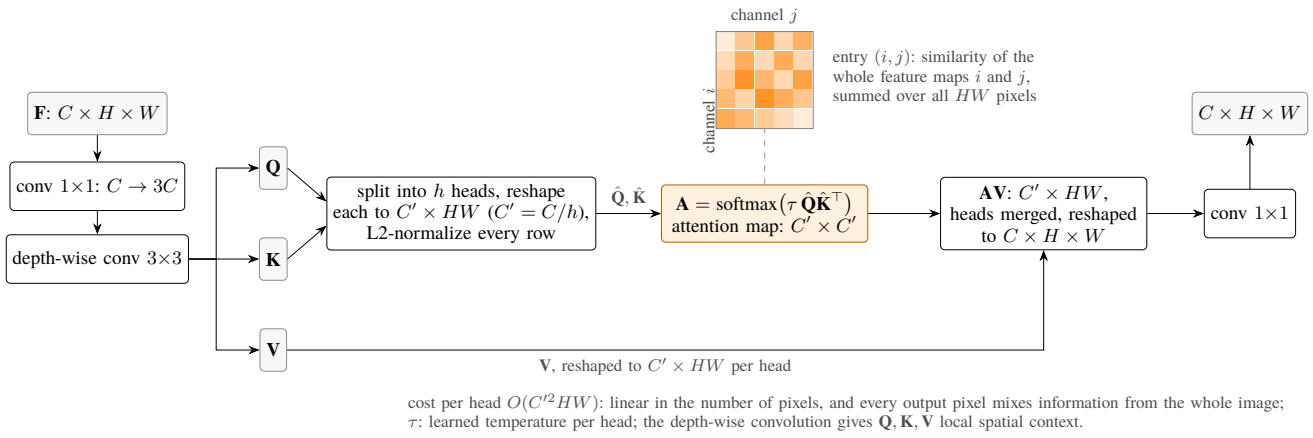


Fig. 5. Multi-Dconv-head transposed attention [17], the token mixer of the prior. Queries, keys and values are computed by point-wise and depth-wise convolutions, and the attention map is formed between the C' feature channels of each head rather than between the HW pixels, so its size does not depend on the image size. Replacing it by a depth-wise convolution changes the PSNR by less than 0.1 dB in the ablation.

TABLE I
THE TWO ×4 PROTOCOLS (DEGRADATION AND SPLIT SIZES OF [6]; IMAGE LOCATIONS OURS)

Item	Chikusei	Pavia Centre
HR-HSI	central 2048^2 crop, 128 bands	1096×715, bands 11–102 (92)
Normalization	÷ 15,133 DN	÷ 8,000 DN
LR-HSI	Gaussian 7×7, $\sigma = 2$; ×4 decimation; reflect border	
HR-MSI	WorldView-2, 8 bands	IKONOS, 4 bands
Test	8 of 256^2 (rows 0–256)	4 of 256^2 (rows, columns 0–512)
Validation	64 of 64^2 (rows 272–400)	63 of 64^2 (rows 528–656)
Training	random 64^2 crops, rows 416–2048	random 64^2 crops, see text
Selection	validation PSNR only	

spectral operators are learned layers. Every method, ExactFuse included, is trained for the same 4 h on one Tesla T4, with the same random patches and augmentation, batch size, EMA weights and validation-only model selection. The published networks keep their official optimizer settings (Adam, L1 loss, learning rate 10^{-4}, or 3×10^{-4} for DHIF-Net, no weight decay or gradient clipping) and are trained in 32-bit precision as released; DCT and MIMFormer use gradient checkpointing to fit into the 16 GB of a T4, and at ×16 CLSNet, OTPNet, MIMO-SST and CaFormer on both datasets and DCT and DHIF-Net on Chikusei split each batch into two micro-batches with summed gradients, U2Net on Pavia Centre into four, which does not change what they compute. The test metrics are computed by the same code for all methods.

B. Metrics

We report the eight metrics of [6], averaged over the test images: PSNR; SSIM; the spectral angle mapper (SAM, degrees, averaged over the pixels whose ground-truth and estimated spectra are both nonzero); ERGAS with ratio 4; Q2n on 32×32 blocks, computed on the image in DN exactly as in the reference implementation of the pansharpening toolbox; the band-averaged correlation coefficient (CC); the spatial correlation coefficient (SCC) of Sobel gradient magnitudes as in the same toolbox; and RMSE in DN. The headline PSNR is $10\log_{10}(1/\text{MSE})$ over the whole normalized image, i.e., its peak is the dataset maximum. The PSNR values of [6] also use a fixed peak: for every row of their Chikusei table, $\text{RMSE} \cdot 10^{\text{PSNR}/20}$ is between 16.6k and 17.1k DN. Since our peak is lower (15,133 DN), our headline PSNR is conservative; with a 16.7k DN peak it would be 0.86 dB higher. For Pavia Centre the same product lies between 7.7k and 8.4k DN (median 8.07k), close to our peak of 8,000 DN, so there the PSNR values are on the same scale. The SSIM definition of [6] is not stated. We report SSIM as in the code released with PSRT [4], whose benchmark data [6] adopts (Gaussian 11×11 window, $\sigma = 1.5$, data range 1, i.e., the dataset peak), and additionally a stricter variant whose data range is each image's own maximum. We also report observation consistency, $\text{PSNR}(\mathcal{D}\hat{\mathbf{X}}, \mathbf{Y}_\text{H})$ and $\text{PSNR}(\mathcal{R}\hat{\mathbf{X}}, \mathbf{Y}_\text{M})$.

C. Implementation

ExactFuse is implemented in PyTorch and trained on Kaggle NVIDIA Tesla T4 GPUs (Table II). On Chikusei we train in two runs. The first run trains from scratch on one GPU with zero-padded operators and the L1 losses only. The final model fine-tunes it with the exact reflect-border operators and the full loss (17) on two GPUs with distributed data parallelism. On Pavia Centre, a single run trains the model from scratch with the exact operators and the full loss on two GPUs, within one 12-hour session that also runs all tests. All runs use AdamW ($\beta = (0.9, 0.99)$, weight decay 10^{-4}), a 1000-iteration linear warm-up followed by a cosine decay to 10^{-3} of the peak rate that ends at the time budget, and gradient-norm clipping at 0.1. The weights with the best validation PSNR are tested once. The ablation and benchmark runs (Sections VI-A and VI-B) train from scratch with these settings, the full loss and a single GPU for 4 h each; their models are evaluated without self-ensemble.

VI. RESULTS

A. Ablation study

Table IV ablates ExactFuse, and Table III the base model module by module; every variant is retrained from scratch with the same budget (4 h on one T4, about 40,000 iterations

TABLE II
TRAINING RUNS

	Chikusei		Pavia Centre
	First run	Fine-tuned (final)	final
Initialization	random	first run	random
Border of $\mathcal{D}$	zero padding	reflect (exact)	reflect (exact)
Loss	L1 + 0.1 deep sup.	(17)	(17)
GPUs $\times$ batch	1×8	2×8	2×8
Peak learning rate	3×10^{-4}	1.5×10^{-4}	3×10^{-4}
Iterations	88,946	94,450	93,850
Training patches	0.71 M	1.51 M	1.50 M
Wall-clock time	9.5 h	10.9 h	10.2 h
Selected at	best val. (58.60 dB)	it. 44,000 (58.72 dB)	it. 43,000 (48.58 dB)

for the base model). From the base model, which takes one gradient data step and gives the prior the HR-MSI itself, the closed-form step with learned gains and the HR-MSI residual input each raise the agreement with the LR-HSI (89.4 and 86.5 against 83.9 dB on Chikusei) at nearly equal PSNR; combined, with every input in the hyperspectral bands, they form ExactFuse, which raises the PSNR by +0.10 dB on Chikusei and 0.00 dB on Pavia Centre and the agreement to 93.4 dB. Confining its prior to the null space with a learned gate lowers the PSNR (-0.08 and -0.04 dB), and an eigen-operator prior with 2.26 M parameters instead of the U-Net costs 0.23 and 0.26 dB while staying above every retrained network. Without data steps ExactFuse loses 0.63 and 0.33 dB; the learned gains, the HR-MSI residual input, $\mathcal{R}^{+}\mathbf{Y}_{\mathrm{M}}$ instead of $\mathbf{Y}_{\mathrm{M}}$ and the closed form instead of a gradient step each change its PSNR by at most 0.05 dB, but without the gains, the residual input or the closed form the agreement with the LR-HSI drops from 93.4 to 90.4, 92.7 and 86.7 dB on Chikusei, and $\mathcal{R}^{+}\mathbf{Y}_{\mathrm{M}}$ makes the prior independent of the MSI sensor. A zero border costs ExactFuse 0.17 and 0.06 dB and learned operators 0.35 and 0.16 dB; with only the LR-HSI it reaches 39.62 dB on Chikusei, with only the HR-MSI 54.88 dB. One stage instead of three, a convolutional mixer instead of attention, or no cross-stage memory change its PSNR by at most 0.09 dB; one stage lowers the agreement with the LR-HSI to 83.4 dB on Chikusei. The rest of this subsection concerns the base model. On Pavia Centre the full model and the zero-border variant were trained with three seeds; their PSNR varies by only 0.02 and 0.03 dB across seeds, which sets the scale for the other differences. Chikusei repeats every variant, again with three seeds for these two rows.

Data step and operator. Removing the data-consistency step together with the residual input (row (2)) costs 0.17 dB and lowers the agreement of the reconstruction with the LR-HSI from 77.6 to 68.9 dB. Keeping the exact kernel but replacing the reflect border by zero padding (row (4)) costs 0.08 dB on average over the three seeds, and each zero-border seed is below each full-model seed. A blur kernel and spectral response learned with the network (row (5)), the operator model of DHIF-Net and MHF-Net, cost 0.06 dB. Both variants reproduce the LR-HSI about 3 dB worse than the exact model (74.5 and 73.6 against 77.6 dB). This is what Theorem 3 predicts: their data steps move the ground truth, and the prior has to compensate. The compensation is tied to the operator used in training: Chikusei weights trained with a zero border fall from 58.60 to 49.52 dB validation PSNR when the exact border is substituted without retraining (Section S-I of the supplementary material). On the test images, the displacement of the ground truth by one LR-HSI data step, $\|\tilde{\mathcal{D}}^{\top}(\tilde{\mathcal{D}} - \mathcal{D})\mathbf{X}\|$ relative to $\|\tilde{\mathcal{D}}^{\top}\tilde{\mathcal{D}}\mathbf{X}\|$, is 4.66% for the zero border and 4.69% for the learned operator, and on 64×64 crops, where border pixels are a larger share, 11.11% and 11.17%. Chikusei (rows (2)–(17)) shows the same ordering with larger gaps: without the data step the PSNR is 0.30 dB lower, with a zero border 0.22 dB and with a learned operator 0.22 dB, and the agreement with the LR-HSI drops from 83.9 to 73.3, 75.9 and 77.1 dB.

The accuracy gap at the training blur is small; it becomes large when the blur changes (last column of Table III and Fig. 6). Given the true blur at test time, the exact model loses at most 0.13 dB for $\sigma \in \{1.5, 2.5, 3.0\}$, the zero-border model with the same swap 0.88 dB, and the learned operator, which cannot be replaced, 1.24 dB. Without the swap the exact model loses 7.50 dB: it is the most sensitive variant to a wrong operator, and the only one that fully recovers when given the right one. On Chikusei the difference is larger still: with the true blur the exact model loses at most 0.25 dB, the zero-border model 3.41 dB and the learned operator 3.45 dB. A zero border thus prevents the swap from working, even though it differs from the exact operator only at the image border, because the prior has learned to compensate a border error tied to the training blur.

Prior. Cross-stage memory (row (10)), the residual input (row (3)), the type of token mixer (rows (14) and (15): a depth-wise convolution or self-attention within 8×8 windows instead of transposed attention) and the number of stages (rows (11)–(13)) change the PSNR on Pavia Centre by at most 0.06 dB, comparable to the spread across seeds. A one-stage model with 3.33 M parameters is within 0.02 dB of the full model with 9.99 M; part of this parity comes from the fixed time budget, in which the smaller model completes almost three times as many iterations. The SAM and SSIM loss terms (row (16)) do not improve the PSNR, whereas on Pavia Centre deep supervision of the intermediate stages does (row (17); 0.12 dB lower without it). On Chikusei, where the three seeds of the full model differ by only 0.004 dB, the same modules change the PSNR by between -0.02 and +0.07 dB: memory +0.01, residual input +0.06, window attention +0.03, convolutional mixer +0.07, two and four stages -0.02 and +0.01, SAM and SSIM loss terms +0.02 and deep supervision +0.04 dB. Several of these simplifications are slightly better than the full model, so its prior is not optimal for this dataset, and deep supervision, which helps on Pavia Centre, does not help here. Only a single stage costs noticeably (0.20 dB): on the larger, 128-band problem at least two stages are needed. We keep the full model for all other experiments, but the ablation shows that ExactFuse's accuracy depends more on its exact data-consistency steps than on the particular design of its transformer prior.

B. Equal-budget benchmark

Table V compares ExactFuse with the twelve retrained networks; CC and SCC are listed in Table S-VI of the supple-

TABLE III

ABLATION: EACH VARIANT CHANGES ONE PART OF THE FULL MODEL AND IS TRAINED FROM SCRATCH WITH THE SAME 4 H BUDGET ON ONE T4 (TEST SET, SINGLE PASS). Δ: PSNR CHANGE AGAINST THE FULL MODEL. CONS.: PSNR$(\mathcal{D}\hat{\mathbf{X}}, \mathbf{Y}_\mathrm{H})$ IN DB. SWAP: WORST TEST PSNR OVER THE BLURS $\sigma \in \{0.75, 1.25, 1.5\}\sigma_0$ WITH THE TRUE BLUR GIVEN AT TEST TIME; N/A: NO OPERATOR TO REPLACE. $\pm$: MEAN AND STANDARD DEVIATION OVER THREE SEEDS.

			Pavia Centre ×4				Chikusei ×4			
#	Variant	Params (M)	PSNR	Δ	Cons.	Swap	PSNR	Δ	Cons.	Swap
(1)	base model (full)	9.99	50.62±0.02	–	77.6	50.50	57.95±0.004	–	83.9	57.70
(2)	no data step, no residual input	9.99	50.46	-0.17	68.9	n/a	57.65	-0.30	73.3	n/a
(3)	no residual input	9.99	50.62	-0.01	75.7	–	58.01	+0.06	82.8	–
(4)	exact kernel, zero border	9.99	50.54±0.03	-0.08	74.5	49.64	57.73±0.05	-0.22	75.9	54.38
(5)	learned blur kernel and SRF	9.99	50.57	-0.06	73.6	n/a	57.73	-0.22	77.1	n/a
(6)	exact proximal step, conjugate gradients	9.99	50.55	-0.07	79.4	50.44	57.74	-0.21	85.2	57.54
(7)	exact proximal step, closed form	9.99	50.65	+0.02	80.9	50.58	57.93	-0.02	87.3	57.76
(8)	learned spectral data step	9.99	50.63	+0.01	82.9	50.50	58.00	+0.05	89.4	57.83
(9)	with the HR-MSI residual input	10.11	50.64	+0.01	79.0	50.58	58.04	+0.09	86.5	57.85
(10)	no cross-stage memory	9.98	50.67	+0.05	77.2	–	57.96	+0.01	84.1	–
(11)	$K = 1$ stage	3.33	50.60	-0.02	71.3	–	57.75	-0.20	77.4	–
(12)	$K = 2$ stages	6.66	50.64	+0.02	75.0	–	57.92	-0.02	81.0	–
(13)	$K = 4$ stages	13.32	50.68	+0.06	80.2	–	57.96	+0.01	85.0	–
(14)	convolutional mixer	9.94	50.60	-0.02	78.5	–	58.02	+0.07	85.2	–
(15)	8×8 window attention	9.99	50.63	+0.01	78.6	–	57.98	+0.03	84.4	–
(16)	L1 + deep supervision only	9.99	50.66	+0.03	77.5	–	57.97	+0.02	82.2	–
(17)	no deep supervision	9.99	50.50	-0.12	77.7	–	57.99	+0.04	84.1	–

TABLE IV

EVERY ABLATION RUN (4 H EACH): EXACTFUSE WITH ONE CHANGE PER ROW, ONE INPUT REMOVED, AND THE BASE MODEL WITH ITS OWN STEPS TOWARDS EXACTFUSE AND ITS OPERATOR ABLATION. CONS.: PSNR$(\mathcal{D}\hat{\mathbf{X}}, \mathbf{Y}_\mathrm{H})$. SWAP: WORST PSNR OVER TEST BLURS $\sigma \in \{0.75, 1.25, 1.5\}\sigma_0$ WHEN THE TRUE BLUR IS GIVEN TO THE DATA STEPS (N/A: NO OPERATOR TO REPLACE; –: NOT EVALUATED).

	Chikusei			Pavia Centre		
Variant	PSNR	Cons.	Swap	PSNR	Cons.	Swap
ExactFuse	58.05	93.4	57.93	50.63	82.5	50.57
closed form without learned gains	58.00	90.4	57.88	50.67	80.6	50.62
gradient step instead of closed form	58.03	86.7	57.84	50.63	79.5	50.56
no HR-MSI residual input	58.03	92.7	57.91	50.60	81.7	50.54
$\mathbf{Y}_\mathrm{M}$ instead of $\mathcal{R}^{+}\mathbf{Y}_\mathrm{M}$ as input	58.07	93.9	57.96	50.58	82.6	50.52
no data steps	57.42	67.5	n/a	50.30	68.1	n/a
zero border	57.88	80.9	57.60	50.56	75.2	50.41
learned blur and response	57.70	77.7	n/a	50.47	71.6	n/a
+ null-space gate	57.97	96.9	57.82	50.58	87.0	50.48
+ gate, EON prior instead of U-Net	57.83	96.5	57.74	50.37	84.9	50.32
+ gate, transfer branch	58.07	96.2	57.93	50.55	85.0	50.39
convolutional mixer	58.01	95.9	57.89	50.57	82.8	50.51
one stage	57.96	83.4	57.62	50.55	75.9	50.43
no cross-stage memory	58.03	92.7	57.92	50.63	81.6	50.58
LR-HSI only (no HR-MSI)	39.62	62.1	n/a	30.26	59.3	n/a
HR-MSI only (no LR-HSI)	54.88	59.7	n/a	48.23	53.5	n/a
base model (gradient step, $\mathbf{Y}_\mathrm{M}$ input)	57.95	83.9	57.70	50.63	77.6	50.50
+ closed-form step, learned gains	58.00	89.4	57.83	50.63	82.9	50.50
+ HR-MSI residual input	58.04	86.5	57.85	50.64	79.0	50.58
without data steps	57.65	73.3	n/a	50.46	68.9	n/a
zero border	57.79	75.9	54.38	50.52	74.5	49.64
learned operator	57.73	77.1	n/a	50.57	73.6	n/a
convolutional mixer	58.02	85.2	–	50.60	78.5	–
LR-HSI only	39.59	61.9	n/a	30.38	58.4	n/a
HR-MSI only	54.86	59.6	n/a	48.22	53.6	n/a

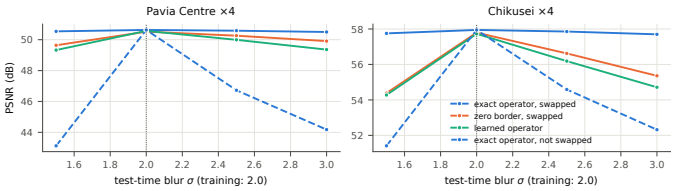


Fig. 6. Test PSNR when the blur of the test data differs from the training blur ($\sigma = 2$, dotted line), for the exact operator with and without the true blur supplied at test time, the zero-border operator with the true blur supplied, and the learned operator, which cannot be replaced.

mentary material. On Chikusei ExactFuse reaches 58.05 dB, 0.79 dB above the best retrained network, and on Pavia Centre 50.63 dB (+0.55 dB), with 8% and 6% lower RMSE; it also has the lowest SAM and ERGAS on both datasets. The strongest competitors are MIMO-SST on Chikusei and OTPNet on Pavia Centre. The unfolded DHIF-Net, whose operators are learned convolutions, reaches 54.39 and 47.67 dB; with 3,250 and 3,950 iterations in 4 h it is also the slowest network to train (Table XIX).

Two checks address the fairness of a fixed budget. First, the retrained baselines reach or exceed the values reported for them in [6] although our data and splits differ: on Chikusei ×4, for example, MIMFormer reaches 56.96 dB against 53.36 dB reported and PSRT 53.99 dB against 50.67 dB, so our pipeline does not weaken them. Second, the budget favors fast networks: DCT completes only 4,800 iterations in 4 h on Pavia Centre, against 33,850 for ExactFuse and 97,400 for MIMFormer. Doubling the budget to 8 h for ExactFuse and MIMFormer gives 58.18 and 57.37 dB on Chikusei ×4: the gap narrows from 1.10 to 0.81 dB but remains, although ExactFuse completes 63,850 iterations in 8 h against 186,800 for MIMFormer. On Pavia Centre ×4 it narrows from 1.00 to 0.67 dB (50.64 against 49.97 dB, after 65,000 and 194,950 iterations). The fixed budget also explains why the small SSR-

NET, which completes its 300,000 iterations in about 1.5 h, ranks ahead of the slower PSRT, DHIF-Net and DCT on both datasets. On Chikusei $\times 4$, ExactFuse is more accurate than the best retrained network on all test images (one-sided Wilcoxon signed-rank test, $p = 0.0039$). Fig. 7 shows the error of every method on the test image of each dataset on which ExactFuse is least accurate. GSA's error follows every edge; the errors of PSRT, DCT and DHIF-Net concentrate on the built-up areas and along building edges, MIMFormer's are smaller, and ExactFuse's are the smallest and the most uniform.

Where the margin comes from. Table VI retrains two of the baselines, MIMFormer and PSRT, with their authors' own recipes and reduces ExactFuse step by step. The authors' recipes do not help under the common budget: MIMFormer reaches 56.00 instead of 56.96 dB on Chikusei and PSRT 49.51 instead of 53.99 dB, so the common pipeline does not put them at a disadvantage. Reduced to its prior alone, a single U-Net on the bicubic image and the HR-MSI without data step or stages, ExactFuse still reaches 57.67 dB on Chikusei and 50.39 dB on Pavia Centre, above every retrained network. Most of the margin over them therefore comes from this residual multiscale prior and its training, and the unrolled exact data steps add 0.38 and 0.23 dB on top; the error split of Section VI-C attributes the margin in the same way.

C. Where the error lies

Theorem 2(c) splits the error of any fused image into a component in the observed subspace $\mathrm{row}(\mathcal{D}) + \mathrm{row}(\mathcal{R})$, which contradicts the observations, and a component in the fusion null space, which no data step can correct and only a prior can reduce. Table VII applies this split to the test predictions of every method of the benchmark. Bicubic interpolation and GSA err mostly in the observed subspace: only 8.6% and 13.7% of their squared error lies in the null space on Chikusei, so they do not even reproduce their inputs. The learned methods differ mainly in how much of the observed part they get wrong. ExactFuse's error in the observed subspace is 27 times smaller than that of the best retrained network on Chikusei (MIMO-SST) and 9 times on Pavia Centre, and 99.9% of its remaining error lies in the null space. The split also attributes ExactFuse's margin: correcting MIMFormer's observed part alone would raise its PSNR by 0.43 dB on Chikusei, and the remaining 0.67 dB come from a better estimate of the null-space component. The exact data steps account for the first part of the margin and the prior for the second; neither the attention mechanism nor any other single module of the prior is responsible for it (Section VI-A).

D. Other datasets and scale factors

Table VIII repeats the equal-budget benchmark on the indoor CAVE [24] and Harvard [25] datasets at $\times 4$, whose test scenes are unseen during training, and at $\times 8$ and $\times 16$, with the blur scaled to a $(2s-1) \times (2s-1)$ Gaussian with $\sigma = s/2$; all metrics are in the supplementary material. Of these 8 settings, ExactFuse has the highest PSNR in 8 and the lowest SAM in 8, with PSNR margins of 0.21–2.29 dB over the strongest network run in each; at $\times 16$, DCT and DHIF-Net complete only 2,150 and 1,050 iterations in 4 h on Chikusei and 2,500 and 1,650 on Pavia Centre, and U2Net 6,100 on Pavia Centre, so these entries are far from converged, and on Chikusei U2Net does not fit into the memory of a T4. Table IX repeats the operator ablation in these settings: the exact operator with the true blur supplied loses at most 0.28 dB over the changed test blurs, against 0.3–3.4 dB for the zero-border and learned operators.

E. Comparison with reported results on Chikusei

The comparisons in this and the next subsection place ExactFuse trained for 8 h (Chikusei) and the long-trained base model (Pavia Centre) against the 17 methods tabulated in [6]. They are less controlled than Section VI-B: the other rows are copied from [6], our test images follow its protocol but are not necessarily the same pixels, and the ExactFuse models here were trained longer (Table II). Table X lists all Chikusei results of [6] together with ExactFuse. The compared methods are GSA [1], CNMF, FUSE [2], U2Net [11], PSRT [4], DCT [20], DDIF, MIMFormer [10], MIMO-SST [5], OTP-Net [12], SMGU [15], AMSF, SSDT, ASMNet, DPFormer, CLSNet [19] and the two-stage conditional diffusion model [6] (references of all of them in Section S-VII of the supplementary material); their values are copied from [6].

The base model ranks first of 18 on PSNR, SSIM (PSRT definition), SAM, ERGAS, CC, SCC and RMSE, with or without the self-ensemble. Relative to the best reported method, PSNR improves by 2.01 dB, RMSE by 26.4%, ERGAS by 12.2% and SAM by 1.1%. It ranks tenth on Q2n. Section S-III of the supplementary material shows that the Q2n deficit comes mainly from the eight ultraviolet bands (363–404 nm), which lie below the first WorldView-2 band and receive no high-resolution detail, and that the reported Q2n values do not follow PSNR the way our implementation does, so they were most likely computed differently. Because the RMSE is expressed in sensor digital numbers, it does not depend on any choice of peak, and it gives the most direct comparison across papers. Table S-IV of the supplementary material lists the long-trained models and the classical baselines under every PSNR and SSIM definition. Our GSA implementation scores 4.4 dB below the value reported in [6], so our protocol does not favor classical methods.

F. Comparison with reported results on Pavia Centre

Table XI compares the base model with the 17 methods of [6] on Pavia Centre. Trained from scratch in one session, ExactFuse reaches 50.69 dB PSNR, 3.53 dB above the best reported value, and lowers the RMSE by 34.1% and ERGAS by 25.3%. It ranks first on six of the eight metrics (PSNR, SSIM in the PSRT definition, ERGAS, CC, SCC and RMSE), but 13th on SAM (1.883° against 1.4542°) and 16th on Q2n (0.9753 against 0.9987). With the stricter per-image data range its SSIM is 0.9970, below the reported 0.9972. The PSNR gain is larger than on Chikusei although the training region is eight times smaller (97 against 800 non-overlapping patch equivalents).

TABLE V

EQUAL-BUDGET BENCHMARK: EVERY LEARNED METHOD RETRAINED BY US UNDER THE IDENTICAL PROTOCOL AND BUDGET (4 H ON ONE T4, THE SAME DATA, PATCHES, BATCH, EMA WEIGHTS AND VALIDATION-ONLY MODEL SELECTION; SINGLE PASS). GSA AND BICUBIC INTERPOLATION NEED NO TRAINING. BEST LEARNED METHOD IN BOLD; SSIM WITH THE DATA RANGE SET TO EACH IMAGE'S MAXIMUM; RMSE IN DN.

	Chikusei ×4						Pavia Centre ×4					
Method	PSNR↑	SSIM↑	SAM↓	ERGAS↓	Q2n↑	RMSE↓	PSNR↑	SSIM↑	SAM↓	ERGAS↓	Q2n↑	RMSE↓
Bicubic	33.63	0.7404	4.406	9.011	0.6914	321.81	25.52	0.6282	6.554	9.852	0.6806	436.32
GSA [1]	43.20	0.9778	2.730	3.279	0.9596	107.69	34.78	0.9438	5.229	3.494	0.9276	157.93
SSR-NET [3]	55.55	0.9958	0.954	1.707	0.9901	25.30	48.45	0.9954	2.443	0.758	0.9719	30.46
PSRT [4]	53.99	0.9948	1.057	1.794	0.9887	30.57	47.87	0.9958	2.212	0.801	0.9716	32.60
DCT [20]	53.18	0.9945	1.133	2.066	0.9871	33.58	46.94	0.9953	2.442	0.888	0.9712	36.26
MIMFormer [10]	56.96	0.9964	0.801	1.488	0.9918	21.65	49.63	0.9967	1.982	0.674	0.9738	26.66
DHIF-Net [8]	54.39	0.9953	1.020	1.660	0.9900	29.37	47.67	0.9956	2.295	0.840	0.9704	33.42
U2Net [11]	53.35	0.9943	1.076	1.948	0.9863	32.92	47.29	0.9955	2.327	0.858	0.9708	34.96
MIMO-SST [5]	57.27	0.9965	0.780	1.419	0.9922	20.90	49.69	0.9967	1.961	0.670	0.9741	26.48
OTPNet [12]	57.06	0.9965	0.786	1.448	0.9922	21.39	50.08	0.9968	1.951	0.653	0.9737	25.29
SMGU-Net [15]	55.14	0.9957	0.949	1.705	0.9898	26.60	48.03	0.9958	2.331	0.783	0.9723	32.05
CLSNet [19]	56.65	0.9963	0.811	1.514	0.9917	22.35	49.50	0.9967	1.993	0.684	0.9735	27.09
CaFormer [16]	53.56	0.9949	1.054	2.138	0.9880	32.07	48.43	0.9961	2.145	0.757	0.9724	30.56
M2DTN [23]	54.87	0.9952	1.003	1.977	0.9877	27.44	48.65	0.9963	2.156	0.729	0.9727	29.76
ExactFuse	**58.05**	**0.9969**	**0.724**	**1.334**	**0.9930**	**19.21**	**50.63**	**0.9970**	**1.892**	**0.619**	**0.9753**	**23.70**

TABLE VI

WHERE THE MARGIN OVER THE RETRAINED NETWORKS COMES FROM (4 H ON ONE T4, SINGLE PASS). THE TWO STRONGEST BASELINES ARE RETRAINED WITH THEIR AUTHORS' OWN RECIPE AND WITH EXACTFUSE'S; EXACTFUSE IS REDUCED TO ITS PRIOR ALONE, TO ONE STAGE, TO THREE STAGES WITHOUT THE DATA STEP, AND TRAINED WITH THE BASELINES' RECIPE. Δ: PSNR AGAINST THE FULL EXACTFUSE OF THE SAME DATASET; IT.: TRAINING ITERATIONS COMPLETED.

	Chikusei ×4				Pavia Centre ×4			
Model and training	PSNR	Δ	SAM	it.	PSNR	Δ	SAM	it.
MIMFormer								
pipeline recipe (batch 8, AdamW, no decay, cosine)	56.96	-1.10	0.801	100k	49.63	-1.00	1.982	97k
authors' recipe (batch 16, Adam, step decay)	56.00	-2.05	0.874	68k	49.32	-1.30	2.022	68k
ExactFuse's recipe (lr 3×10^{-4}, decay, clipping, SAM+SSIM)	57.52	-0.53	0.760	81k	50.23	-0.40	1.954	84k
PSRT								
pipeline recipe (batch 8)	53.99	-4.06	1.057	64k	47.87	-2.75	2.212	65k
authors' recipe (batch 32, AdamW decay, clipping, constant lr)	49.51	-8.54	1.551	17k	45.44	-5.18	2.597	17k
ExactFuse								
prior alone: one U-Net, no data step	57.67	-0.38	0.747	113k	50.39	-0.23	1.934	115k
one stage with data step	57.75	-0.31	0.743	97k	50.60	-0.02	1.909	109k
three stages, no data step	57.65	-0.41	0.752	41k	50.46	-0.17	1.928	42k
baselines' recipe (lr 10^{-4}, no decay, no clipping, L1)	57.67	-0.38	0.749	32k	50.07	-0.55	1.921	35k
full model	58.05	–	0.724	30k	50.63	–	1.892	34k

The Q2n shortfall on this dataset, and part of the SAM shortfall, come from one test image that contains a stretch of the river Ticino (63% water pixels; Section S-IV of the supplementary material). Water reflects little light, so its spectra are dominated by sensor noise that neither observation carries: over water ExactFuse's SAM is 2.95°, over all other pixels 1.56°, and this image alone carries 97% of the Q2n deficit of the four test images. Whether the test images of [6] contain water is not stated, so on this dataset SAM and Q2n are the least comparable of the eight metrics and RMSE the most. As on Chikusei, our GSA implementation scores below the value reported for it (34.78 against 35.35 dB), so the rebuilt data do not favor our method.

G. Robustness of the long-trained models

The long-trained base models behave like the ablation models under a blur mismatch (Table S-II and Fig. S3 of the supplementary material). With the training operator, a blur with σ between 1.5 and 3.0 costs 3.9–7.5 dB on Chikusei and 3.3–6.8 dB on Pavia Centre; with the true blur swapped in, it costs at most 0.26 and 0.11 dB. A spectral response shifted by 8 nm costs about 5 dB on Chikusei and is not compensated by swapping in the shifted response, because the priors rely on the training response beyond the data step; on Pavia Centre,

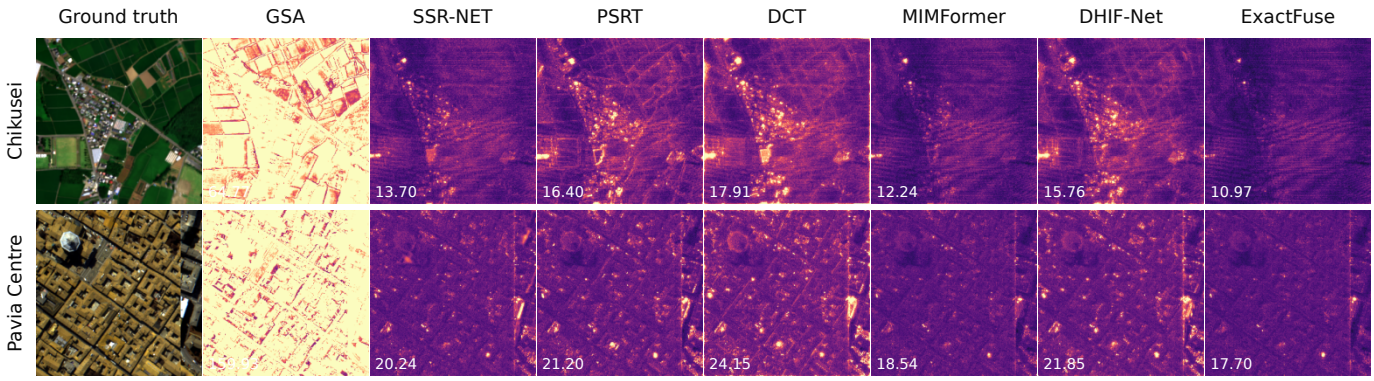


Fig. 7. Equal-budget benchmark: ground truth in false color and the mean absolute error over all bands (DN) of every method on the test image on which ExactFuse has its lowest PSNR. The color scale is set by the retrained networks (clipped at the 99.5th percentile of their errors), so GSA saturates. The number in each map is its mean.

TABLE VII

WHERE THE ERROR LIES (THEOREM 2). OBS. MSE: MEAN SQUARED ERROR IN THE OBSERVED SUBSPACE $\mathrm{row}(\mathcal{D}) + \mathrm{row}(\mathcal{R})$, THE PART THAT CONTRADICTS THE OBSERVATIONS ($\times 10^{-6}$, NORMALIZED INTENSITIES). NULL: SHARE OF THE SQUARED ERROR IN THE FUSION NULL SPACE, WHICH ONLY A PRIOR CAN REDUCE. $\mathrm{PSNR}_{\mathcal{N}}$: THE PSNR THE METHOD WOULD REACH IF ITS OBSERVED PART WERE EXACT. EQUAL-BUDGET MODELS, SINGLE PASS.

	Chikusei $\times 4$				Pavia Centre $\times 4$			
Method	PSNR	Obs. MSE	Null (%)	$\mathrm{PSNR}_{\mathcal{N}}$	PSNR	Obs. MSE	Null (%)	$\mathrm{PSNR}_{\mathcal{N}}$
Bicubic	33.63	431.001	8.6	44.27	25.52	2745.840	11.5	34.91
GSA	43.20	46.899	13.7	51.92	34.78	369.778	18.7	42.39
SSR-NET	55.55	0.703	74.9	56.84	48.45	2.049	85.8	49.11
PSRT	53.99	1.645	64.0	56.01	47.87	2.483	85.4	48.56
DCT	53.18	2.240	58.2	55.60	46.94	4.312	79.5	47.94
MIMFormer	56.95	0.186	90.6	57.38	49.63	0.430	96.3	49.79
DHIF-Net	54.39	1.546	66.1	56.32	47.67	3.980	78.1	48.75
ExactFuse	58.05	0.002	99.9	58.06	50.62	0.028	99.7	50.64

whose IKONOS bands are 66–96 nm wide, the shift costs about 1 dB. Noise is the main weakness of models trained on noise-free data: at 30 dB SNR, ExactFuse exceeds GSA by only 2.8 dB on Chikusei and by 6.0 dB on Pavia Centre.

H. One network for a family of observation models

Table XII tests the networks of Section IV-D on conditions that leave the training protocol: blur widths from $0.5\sigma_0$ to $2\sigma_0$, spectral responses shifted by up to 12 nm, noise from 40 to 30 dB SNR, all three at once and, outside the operator family, anisotropic and defocus blurs and an HR-MSI shifted against the LR-HSI by a fraction of a pixel (every condition is listed in the supplementary material). The operators are estimated from each test pair with Theorem 1. Within the family the estimate is accurate: on every test image and condition, 30 dB SNR included, the largest errors are 0.05 px in blur width and 0.0 nm in spectral shift on Chikusei, and 0.00 px and 0.0 nm on Pavia Centre (search steps 0.05 px and 1 nm).

Three results stand out. First, training over the family is what makes the network robust to noise and to combined changes. At 40–30 dB SNR the mean PSNR rises from 47.60 to 52.64 dB on Chikusei and from 43.15 to 46.55 dB on Pavia Centre, and under combined changes from 46.71 to 52.59 dB and from 42.32 to 46.53 dB. The price is a lower PSNR at the training operator, 57.77 against 58.05 dB on Chikusei and 50.43 against 50.63 dB on Pavia Centre.

Second, in the base model, giving the data steps the estimated operator adds to randomized training alone. With observation-model conditioning it is the most accurate of the base-model variants in every group, for instance 52.83 against 52.30 dB for the same network trained on randomized data with the nominal operator under combined changes on Chikusei, and 46.80 against 46.35 dB on Pavia Centre. Observation-model conditioning adds little to operator conditioning alone (52.92 against 52.79 dB under noise on Chikusei), and the exact operator stays ahead of the zero border (52.70 against 52.45 dB).

Third, outside the family the conditioned network is fragile. With an anisotropic or a defocus blur, the estimated isotropic operator is wrong; the conditioned network, which trusts its

TABLE VIII

EQUAL-BUDGET BENCHMARK ON THE OTHER DATASETS AND SCALE FACTORS: PSNR (DB) AND SAM (DEGREES) OF EVERY METHOD, RETRAINED UNDER THE SAME PROTOCOL AND 4 H BUDGET ON ONE T4 (SINGLE PASS). AT $\times s$ THE BLUR IS A $(2s-1)\times(2s-1)$ GAUSSIAN WITH $\sigma = s/2$; $\times 16$ TRAINS ON 128×128 PATCHES (BATCH 4). BEST LEARNED METHOD IN BOLD; –: RUN NOT AVAILABLE.

	CAVE $\times 4$		Harvard $\times 4$		CAVE $\times 8$		Harvard $\times 8$		Chikusei $\times 8$		Pavia Centre $\times 8$		Chikusei $\times 16$		Pavia Centre $\times 16$	
Method	PSNR↑	SAM↓	PSNR↑	SAM↓	PSNR↑	SAM↓	PSNR↑	SAM↓	PSNR↑	SAM↓	PSNR↑	SAM↓	PSNR↑	SAM↓	PSNR↑	SAM↓
GSA [1]	42.30	4.651	46.30	2.379	38.01	7.273	43.13	2.700	38.69	3.972	30.24	7.163	35.20	5.303	27.07	8.960
SSR-NET [3]	46.51	2.922	49.18	2.148	44.25	3.657	48.61	2.262	54.57	1.030	47.71	2.648	53.98	1.086	47.38	2.662
PSRT [4]	48.51	2.394	49.29	2.127	45.21	3.229	48.61	2.259	53.05	1.172	47.13	2.394	49.90	1.599	45.47	2.755
DCT [20]	46.36	2.794	48.92	2.188	43.17	3.929	47.85	2.410	51.77	1.318	45.51	2.739	47.88	1.929	43.43	3.177
MIMFormer [10]	50.93	1.955	49.72	2.061	48.10	2.413	49.12	2.163	55.57	0.919	49.09	2.064	54.21	1.029	48.06	2.206
DHIF-Net [8]	49.82	2.047	49.58	2.084	47.19	2.511	49.07	2.177	53.24	1.173	47.33	2.351	45.25	2.572	41.99	3.492
ExactFuse	**52.46**	**1.659**	**50.03**	**2.009**	**50.39**	**1.934**	**49.45**	**2.116**	**56.95**	**0.817**	**50.05**	**1.977**	**56.11**	**0.882**	**49.25**	**2.101**

TABLE IX

OPERATOR ABLATION ACROSS DATASETS AND SCALE FACTORS (ONE RUN EACH, SEED 0, 4 H ON ONE T4; TEST PSNR IN DB). TRAINING BLUR: PSNR AT THE BLUR USED FOR TRAINING. CHANGED BLUR: WORST PSNR OVER THE THREE TEST BLURS $\sigma \in \{0.75, 1.25, 1.5\}\,\sigma_0$, WITH THE TRUE BLUR SUPPLIED (SWAP) OR THE TRAINING OPERATOR KEPT (FIXED). –: RUN NOT AVAILABLE.

	Training blur				Changed blur (worst case)				
Setting	Exact	Zero border	Learned	No data step	Exact, swap	Exact, fixed	Zero border, swap	Learned	No data step
Chikusei $\times 4$	57.95	57.79	57.73	57.65	57.70	51.41	54.38	54.28	–
Pavia Centre $\times 4$	50.63	50.52	50.57	50.46	50.50	43.13	49.64	49.33	–
CAVE $\times 4$	52.57	52.19	52.22	51.69	52.29	48.66	51.82	51.08	51.38
Harvard $\times 4$	50.01	49.97	49.98	49.84	49.95	48.41	49.29	49.69	49.74
Chikusei $\times 8$	56.82	56.53	56.39	55.96	56.69	50.00	55.27	55.31	55.88
Pavia Centre $\times 8$	49.94	49.65	49.56	49.54	49.88	42.69	49.38	48.43	49.52

operator, falls to 55.34 dB on Chikusei and 46.76 dB on Pavia Centre, whereas the network trained on randomized data with the nominal operator keeps 57.18 and 50.01 dB; given the true kernel, the conditioned network recovers (57.59 and 50.33 dB). A sensitivity test shows why. With noise-free nominal data and a blur of $0.9\,\sigma_0$ in its data steps, a relative operator residual $\|(\tilde{\mathcal{D}} - \mathcal{D})\mathbf{X}\|/\|\mathcal{D}\mathbf{X}\|$ of only 0.7%, the conditioned network drops from 57.77 to 57.32 dB on Chikusei and the nominally trained one from 58.05 to 55.40 dB, while the randomized-only network is unaffected (57.42 dB): as Theorem 3 states, an exact data step trusts its operator. Exactness thus makes the network precise when the operator family contains the test conditions and fragile when it does not, so the family used for training and estimation has to cover the sensors it will meet. Finally, an HR-MSI misregistered by half a pixel costs every method trained on registered pairs more than 15 dB (38.09 dB on Chikusei, with GSA at a similar loss): none of them models registration errors. Registering the HR-MSI before fusion, with the global offset estimated from the pair by the criterion of Theorem 1 on a 1/8-pixel grid, recovers most of the loss. Under a misregistration of 0.5–1 px, ExactFuse rises from 37.88 to 49.74 dB on Chikusei and from 29.72 to 40.66 dB on Pavia Centre; MIMFormer registered the same way reaches 49.38 and 40.38 dB, and GSA with the same registration 42.42 and 33.78 dB. Putting the offset into ExactFuse's operator instead and shifting the result back is worse (40.60 dB on Chikusei), mainly at a full-pixel offset. With an anisotropic blur and 35 dB noise added to a half-pixel misregistration, the network trained on noise-free data falls behind MIMFormer (43.48 against 45.49 dB on Chikusei), as it does under noise alone. Trained over the wider family, which includes noise, anisotropic blurs and grid offsets, ExactFuse reaches 45.72 dB there and 36.90 dB on Pavia Centre, against 36.92 dB for MIMFormer.

I. Stability of the operator estimate

Table XIII checks Theorem 1(c) on the test images. The Jacobian of the residual has full rank for all five parameters (two blur widths, two grid offsets, response shift), with condition numbers of 1.28 on Chikusei and 1.47 on Pavia Centre after scaling its columns, so the parameters are nearly uncorrelated. For the grid offsets and the response shift, the first-order prediction of Theorem 1(c) matches the measured error of the estimate (at 30 dB SNR, 0.0013 px predicted and 0.0013 px measured for the offsets on Chikusei). For the blur widths the measured error exceeds the prediction as the SNR falls (0.0017 against 0.0042 px at 30 dB), because noise also enters the Jacobian and biases the blur estimate upward, a second-order effect. All errors stay far below the grid step of 0.05 px. For scenes in the null space of $C_{\theta\phi}$ the corresponding columns of $\mathbf{G}$ vanish, as Theorem 1(b) predicts.

a) Confidence intervals and per-image results.: The 95% paired bootstrap interval of ExactFuse's mean per-image gain over the strongest retrained network excludes zero in every setting (Table XIV); the smallest gain, 0.21 dB on Harvard, has the interval [0.11, 0.33] dB, and pooled over all test images the gain is 0.91 dB [0.74, 1.09]. ExactFuse has the lower SAM and ERGAS on every test image of every setting and the higher

TABLE X

CHIKUSEI ×4 COMPARISON WITH PUBLISHED NUMBERS. ALL ROWS EXCEPT EXACTFUSE ARE COPIED FROM TABLE IV OF [6], NOT RETRAINED BY US; THEY COME FROM ANOTHER IMPLEMENTATION OF THE PROTOCOL AND GIVE CONTEXT, NOT A RANKING (TABLE V IS THE CONTROLLED COMPARISON). BEST VALUE PER COLUMN IN BOLD. SSIM OF EXACTFUSE USES THE PSRT DEFINITION (DATA RANGE 1); WITH DATA RANGE EQUAL TO EACH IMAGE'S MAXIMUM IT IS 0.9969 (SINGLE PASS AND SELF-ENSEMBLE).

Method	PSNR↑	SSIM↑	SAM↓	ERGAS↓	Q2n↑	CC↑	SCC↑	RMSE (DN)↓
GSA [1]	47.5769	0.9933	1.7216	2.1667	0.9658	0.9894	0.9773	70.54
CNMF	44.5291	0.9891	1.7259	2.8676	0.9662	0.9828	0.9645	101.64
FUSE [2]	41.1550	0.9787	2.5814	3.4710	0.9418	0.9736	0.9486	148.10
U2Net [11]	50.5455	0.9956	1.2389	1.8740	0.9928	0.9920	0.9798	49.97
PSRT [4]	50.6722	0.9952	1.2733	2.0107	0.9862	0.9906	0.9778	49.00
DCT [20]	52.1998	0.9977	0.9239	1.7448	0.9832	0.9935	0.9873	41.45
DDIF	53.4421	0.9972	0.9387	1.8087	0.9963	0.9929	0.9840	35.58
MIMFormer [10]	53.3578	0.9973	0.9094	1.7611	0.9948	0.9931	0.9850	35.90
MIMO-SST [5]	53.6073	0.9972	0.9315	1.8119	0.9944	0.9926	0.9834	34.93
OTPNet [12]	52.8243	0.9969	1.0201	1.9332	0.9940	0.9916	0.9841	38.49
SMGU [15]	54.5244	0.9977	0.8719	1.6405	0.9952	0.9939	0.9875	31.29
AMSF	49.3827	0.9944	1.4852	2.3943	0.9878	0.9879	0.9790	56.78
SSDT	52.4788	0.9975	0.8768	1.7992	0.9939	0.9928	0.9847	40.00
ASMNet	53.4528	0.9972	0.9472	1.9153	0.9874	0.9917	0.9848	35.89
DPFormer	53.1632	0.9970	0.9483	1.8795	0.9949	0.9923	0.9824	37.18
CLSNet [19]	55.2007	0.9979	0.7836	1.6318	0.9943	0.9941	0.9869	29.06
Two-stage diffusion [6]	56.1871	0.9982	0.7203	1.5158	**0.9969**	0.9948	0.9891	25.73
ExactFuse, 8 h (single pass)	58.1840	**0.9989**	0.7146	1.3284	0.9931	**0.9967**	**0.9992**	18.96
ExactFuse, 8 h (self-ensemble)	**58.2129**	**0.9989**	**0.7129**	**1.3260**	0.9931	**0.9967**	**0.9992**	**18.91**

TABLE XI

PAVIA CENTRE ×4 COMPARISON. ALL ROWS EXCEPT THE BASE MODEL ARE REPORTED IN [6] ON THE BENCHMARK DATA RELEASED WITH PSRT [4]; THE BASE MODEL IS EVALUATED ON OUR SIMULATION OF THE SAME PROTOCOL FROM THE ORIGINAL SCENE (SECTION V). BEST VALUE PER COLUMN IN BOLD. SSIM OF THE BASE MODEL USES THE PSRT DEFINITION (DATA RANGE 1); WITH DATA RANGE EQUAL TO EACH IMAGE'S MAXIMUM IT IS 0.9970 (SINGLE PASS) AND 0.9970 (SELF-ENSEMBLE).

Method	PSNR↑	SSIM↑	SAM↓	ERGAS↓	Q2n↑	CC↑	SCC↑	RMSE (DN)↓
GSA	35.3495	0.9736	5.0812	2.7260	0.9818	0.9832	0.9768	137.78
CNMF	32.1693	0.9469	3.8494	3.8446	0.9489	0.9683	0.9294	198.26
FUSE	34.7418	0.9612	5.5628	3.2042	0.9103	0.9779	0.9469	146.90
U2Net	44.3439	0.9963	1.6542	1.0924	0.9964	0.9973	0.9963	49.10
PSRT	45.8875	0.9970	1.5230	0.9750	0.9974	0.9978	0.9966	41.85
DCT	46.6745	0.9971	1.5169	0.8753	0.9973	0.9983	0.9972	37.66
DDIF	46.2530	0.9970	1.5253	0.9125	0.9960	0.9982	0.9970	39.56
MIMFormer	44.8497	0.9959	1.8225	1.0683	0.9957	0.9975	0.9954	46.19
MIMO-SST	46.6334	0.9969	1.5587	0.8607	0.9975	0.9984	0.9972	37.68
OTPNet	45.8081	0.9967	1.6573	0.9465	**0.9987**	0.9980	0.9968	41.51
SMGU	47.1239	0.9972	1.4690	0.8547	0.9981	0.9984	0.9974	37.09
AMSF	43.4465	0.9950	2.0533	1.2728	0.9929	0.9965	0.9940	54.25
SSDT	46.3173	0.9965	1.6439	0.8918	0.9968	0.9982	0.9969	39.01
ASMNet	46.2072	0.9966	1.6517	0.9019	0.9940	0.9982	0.9969	39.50
DPFormer	42.4280	0.9948	1.9367	1.3986	0.9945	0.9957	0.9917	60.96
CLSNet	46.5277	0.9969	1.5708	0.8412	0.9983	0.9983	0.9975	36.21
Two-stage diffusion	47.1597	0.9972	**1.4542**	0.8262	0.9980	0.9985	0.9976	35.71
Base model (single pass)	50.6299	**0.9974**	1.8955	0.6206	0.9751	**0.9994**	**0.9995**	23.68
Base model (self-ensemble)	**50.6915**	**0.9974**	1.8833	**0.6170**	0.9753	**0.9994**	**0.9995**	**23.52**

PSNR on all but one Harvard image (Table XV lists Chikusei and Pavia Centre).

b) Training curves.: Is the budget long enough? Fig. 8 shows the validation PSNR of every network. The four strongest baselines gain at most 0.11 dB on Chikusei and 0.08 dB on Pavia Centre in the last quarter of their budget (MIMO-SST 0.02 dB, ExactFuse 0.03 dB on Chikusei), so they are close to converged. DCT and DHIF-Net, the slowest per iteration, still gain 1.37 and 1.13 dB and are undertrained, but end more than 3 dB below the strongest networks; the 8 h runs of Section VI-B test the effect of a longer budget.

c) Spectral regions without HR-MSI coverage.: The WorldView-2 HR-MSI does not respond to 13 of the 128 Chikusei bands (Fig. 9); there only the LR-HSI and the prior inform the result. ExactFuse's mean RMSE in these bands is 20.9 DN, against 25.2 DN for MIMFormer, and 12.7 against 15.6 DN in the covered bands.

d) Blur of any shape.: The identifiability residual $\mathcal{D}_k\mathbf{Y}_\mathrm{M} - \mathcal{R}_\phi\mathbf{Y}_\mathrm{H}$ is linear in the kernel k, so a kernel of any shape is the least-squares solution over all bands and low-resolution pixels, with taps that sum to one (5 ms per image). On noise-free pairs its relative error against the true kernel is 0.00% on Chikusei and 0.79% on Pavia Centre over the out-of-family blurs, and the network given it reaches the PSNR it reaches with the true kernel (57.60 and 50.30 dB under the anisotropic and defocus blurs, against 57.59 and 50.33 dB with

TABLE XII

ONE NETWORK FOR A FAMILY OF SENSORS: MEAN TEST PSNR (DB) OVER GROUPS OF TEST CONDITIONS THAT DIFFER FROM THE TRAINING PROTOCOL. RANDOMIZED: TRAINED WITH A BLUR, SRF SHIFT AND NOISE LEVEL DRAWN ANEW FOR EVERY BATCH; + OPERATOR: THE BATCH'S OPERATOR IS ALSO GIVEN TO THE DATA STEPS. FIXED: THE TRAINING-PROTOCOL OPERATOR IS KEPT AT TEST TIME; EST.: BLUR AND SRF SHIFT ESTIMATED FROM THE TEST PAIR ALONE; TRUE: THE TRUE TEST OPERATOR (AN UPPER REFERENCE, NOT MARKED). †ANISOTROPIC AND DEFOCUS BLURS, WHICH THE OPERATOR FAMILY DOES NOT CONTAIN; ‡AN HR-MSI SHIFTED BY 0.5–1 PIXEL AGAINST THE LR-HSI. BEST PER ROW IN BOLD.

	MIMFormer, nominal training	MIMFormer, randomized	ExactFuse, nominal training		base model, randomized data only		base model, randomized, zero border	base model, randomized + operator			+ observation-model conditioning			ExactFuse, randomized + conditioning		
Test condition	fixed	fixed	fixed	est.	fixed	est.	est.	fixed	est.	true	fixed	est.	true	fixed	est.	true
Chikusei ×4																
Training operator	56.96	56.50	**58.05**	**58.05**	57.43	57.43	57.46	57.59	57.59	57.59	57.61	57.61	57.61	57.77	57.77	57.77
Blur ×0.5–2	54.18	56.38	48.00	**57.79**	57.33	57.13	57.36	53.69	57.49	57.49	53.68	57.51	57.51	54.32	57.68	57.68
SRF shift ±4–12 nm	53.61	55.66	52.72	55.13	57.35	55.13	57.39	51.45	57.51	57.51	49.72	57.54	57.54	50.20	**57.73**	57.73
Noise 40–30 dB	50.33	51.67	47.60	47.60	52.35	52.35	52.52	52.79	52.79	52.79	**52.92**	**52.92**	52.92	52.64	52.64	52.64
Blur + SRF + noise	48.98	51.30	44.31	46.71	52.30	51.12	52.45	48.34	52.70	52.70	46.53	**52.83**	52.83	47.23	52.59	52.59
Worst of the above	45.78	48.03	40.84	41.89	49.25	48.12	49.49	45.99	49.80	49.80	44.09	**49.98**	49.98	44.94	49.57	49.57
Blur outside the family†	54.23	56.43	47.38	50.79	**57.18**	57.07	56.39	53.59	54.79	57.41	53.57	54.78	57.31	54.24	55.34	57.59
LR-HSI grid offset	50.60	54.90	44.87	44.82	55.40	**55.41**	–	50.78	50.88	55.98	50.76	50.87	57.43	51.80	51.90	56.08
HR-MSI misregistered‡	38.00	37.46	**38.43**	38.38	37.48	37.48	37.82	38.15	38.12	–	38.15	38.12	–	38.09	38.06	–
Corrupted observations	–	–	**39.13**	38.86	–	–	–	–	–	–	–	–	–	37.23	37.18	37.23
Locally misregistered	–	–	**50.63**	**50.63**	–	–	–	–	–	–	–	–	–	–	–	–
Band-dependent noise	39.17	38.70	**44.43**	44.40	–	–	–	–	–	–	–	–	–	–	–	–
Pavia Centre ×4																
Training operator	49.63	49.37	**50.63**	**50.63**	50.40	50.40	50.41	50.45	50.45	50.45	50.42	50.42	50.42	50.43	50.43	50.43
Blur ×0.5–2	47.05	49.28	40.36	**50.50**	50.33	50.22	50.36	45.85	50.40	50.40	46.19	50.35	50.35	45.92	50.38	50.38
SRF shift ±4–12 nm	48.66	48.31	49.61	49.02	50.32	49.39	50.35	49.59	**50.39**	50.39	49.00	50.38	50.38	48.96	**50.39**	50.39
Noise 40–30 dB	45.24	45.87	43.15	43.15	46.40	46.40	46.51	46.76	46.76	46.76	**46.84**	**46.84**	46.84	46.55	46.55	46.55
Blur + SRF + noise	44.15	45.22	38.96	42.32	46.35	45.90	46.49	44.57	46.72	46.72	44.64	**46.80**	46.80	44.52	46.53	46.53
Worst of the above	41.35	42.63	32.71	37.95	43.69	43.29	43.91	40.48	44.27	44.27	41.38	**44.39**	44.39	40.82	43.97	43.97
Blur outside the family†	46.89	49.32	39.20	42.37	**50.01**	49.89	48.65	45.45	46.77	50.41	45.87	47.02	50.29	45.58	46.76	50.33
LR-HSI grid offset	43.83	**48.21**	37.48	37.40	47.79	47.69	–	42.60	42.69	49.29	43.12	43.20	49.28	42.94	43.01	49.23
HR-MSI misregistered‡	29.86	29.33	**30.33**	30.31	29.50	29.52	29.75	30.09	30.07	–	30.03	30.01	–	30.08	30.06	–
Corrupted observations	–	–	33.40	33.01	–	–	–	–	–	–	33.93	33.73	33.93	**34.06**	33.87	34.06
Locally misregistered	–	–	43.62	**43.63**	–	–	–	–	–	–	43.09	43.11	43.09	43.16	43.17	43.16
Band-dependent noise	30.86	30.42	36.88	36.90	–	–	–	–	–	–	**43.30**	**43.30**	43.30	**43.30**	**43.30**	43.30

TABLE XIII

CONDITIONING AND STABILITY OF THE OPERATOR ESTIMATE (THEOREM 1(C)) ON THE TEST IMAGES. $\tilde{\kappa}$: CONDITION NUMBER OF THE JACOBIAN $\mathbf{G}$ OF THE RESIDUAL WITH RESPECT TO THE FIVE PARAMETERS (TWO BLUR WIDTHS, TWO GRID OFFSETS, RESPONSE SHIFT) AFTER SCALING ITS COLUMNS TO UNIT NORM. PRED.: FIRST-ORDER STANDARD DEVIATION $\mathbf{G}^{+}\boldsymbol{\Sigma}_{\mathbf{n}}\mathbf{G}^{+\top}$; MEAS.: RMSE OF THE ESTIMATE (GRID SEARCH REFINED BY LEVENBERG–MARQUARDT) OVER THE TEST IMAGES AND THREE NOISE DRAWS. BLUR WIDTH AND OFFSET IN PIXELS (MEAN OF THE TWO AXES), RESPONSE SHIFT IN NM.

			40 dB SNR						30 dB SNR						20 dB SNR					
			blur		offset		shift		blur		offset		shift		blur		offset		shift	
Dataset	Case	$\tilde{\kappa}$	pred.	meas.	pred.	meas.	pred.	meas.	pred.	meas.	pred.	meas.	pred.	meas.	pred.	meas.	pred.	meas.	pred.	meas.
Chikusei	nominal	1.28	0.0005	0.0008	0.0004	0.0005	0.003	0.003	0.0017	0.0042	0.0013	0.0013	0.010	0.013	0.0054	0.0363	0.0042	0.0038	0.030	0.033
	anisotropic	1.23	0.0007	0.0010	0.0005	0.0006	0.003	0.003	0.0024	0.0050	0.0016	0.0014	0.010	0.013	0.0075	0.0417	0.0049	0.0048	0.030	0.033
	offsets and shift	1.29	0.0006	0.0009	0.0004	0.0005	0.003	0.003	0.0017	0.0041	0.0013	0.0014	0.010	0.012	0.0055	0.0370	0.0042	0.0036	0.032	0.042
Pavia Centre	nominal	1.47	0.0004	0.0004	0.0003	0.0003	0.012	0.022	0.0014	0.0016	0.0010	0.0013	0.038	0.075	0.0043	0.0126	0.0031	0.0040	0.119	0.697
	anisotropic	1.38	0.0006	0.0005	0.0004	0.0003	0.012	0.022	0.0018	0.0020	0.0011	0.0014	0.038	0.078	0.0057	0.0139	0.0036	0.0050	0.119	0.725
	offsets and shift	1.47	0.0004	0.0004	0.0003	0.0003	0.013	0.020	0.0014	0.0016	0.0010	0.0013	0.042	0.066	0.0043	0.0133	0.0030	0.0035	0.132	0.560

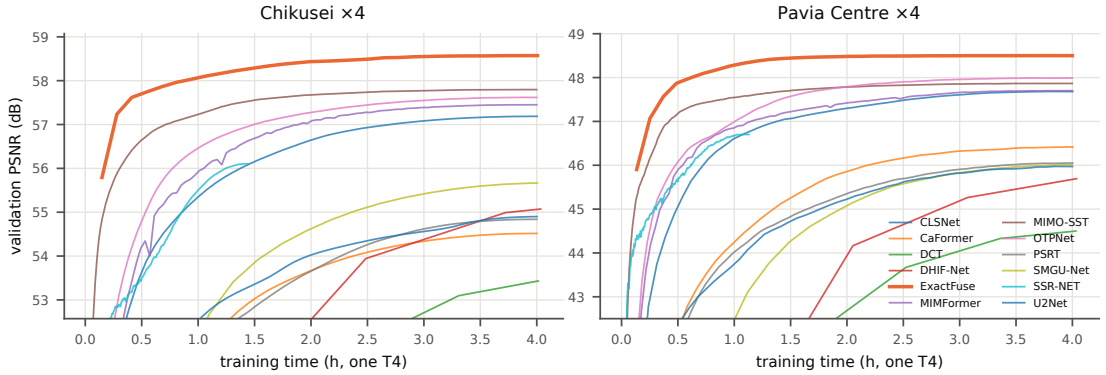


Fig. 8. Validation PSNR of every network during its training budget (4 h on one T4; SSR-NET 1.5 h), Chikusei and Pavia Centre ×4.

the true kernel and 56.43 and 49.32 dB for MIMFormer trained on the same randomized data). Its price is noise sensitivity: at 30 dB SNR the kernel error grows to 4.6% and 2.8%, and the PSNR is 0.10 and 0.16 dB below that with the isotropic estimate; at the training operator the two agree within 0.05 dB. A non-separable kernel loses the closed form, so the network takes its conjugate-gradient data step with the estimated kernel.

e) Misregistration.: Table XVI compares the strategies on six conditions. The bilinear shift of the simulation also blurs the HR-MSI, which caps every registered result (ExactFuse 49.74 dB on Chikusei); with a sharp shift, as a sensor records it, registration brings ExactFuse to 54.56 dB, close to the 58.05 dB without misregistration. Noise is where the network trained on noise-free data falls behind MIMFormer, as in the robustness table; ExactFuse trained over the wider family, with noise and grid offsets, has the best registered result in every noisy condition but one, where it ties MIMFormer (36.90 against 36.92 dB). Putting the offset into ExactFuse's operator and shifting the result back is worse than registering the HR-MSI first (40.60 dB on Chikusei). The MIMFormer of this comparison is trained on noise-free data; M2DTN [23], a registration-aware network that warps the LR-HSI onto the

TABLE XIV
PER-IMAGE GAIN OF EXACTFUSE OVER THE STRONGEST RETRAINED NETWORK OF EACH SETTING (BY MEAN PSNR): MEAN PSNR GAIN (DB) WITH ITS 95% PAIRED BOOTSTRAP INTERVAL OVER THE TEST IMAGES (10,000 RESAMPLES), AND THE NUMBER OF TEST IMAGES ON WHICH EXACTFUSE HAS THE HIGHER PSNR AND THE LOWER SAM AND ERGAS.

Setting	Strongest	n	Gain	95% CI	PSNR	SAM	ERGAS
Chik. ×4	MIMO-SST	8	0.79	[0.52, 1.08]	8/8	8/8	8/8
Pavia ×4	OTPNet	4	0.55	[0.41, 0.69]	4/4	4/4	4/4
CAVE ×4	MIMFormer	12	1.53	[1.30, 1.75]	12/12	12/12	12/12
Harv. ×4	MIMO-SST	20	0.21	[0.11, 0.33]	19/20	20/20	20/20
CAVE ×8	MIMFormer	12	2.29	[1.82, 2.81]	12/12	12/12	12/12
Chik. ×8	OTPNet	8	0.84	[0.51, 1.22]	8/8	8/8	8/8
Pavia ×8	OTPNet	4	0.53	[0.42, 0.75]	4/4	4/4	4/4
Chik. ×16	MIMO-SST	8	1.64	[0.98, 2.41]	8/8	8/8	8/8
Harv. ×8	MIMFormer	20	0.33	[0.18, 0.52]	20/20	20/20	20/20
Pavia ×16	MIMO-SST	4	0.99	[0.70, 1.29]	4/4	4/4	4/4

TABLE XV
EVERY TEST IMAGE: PSNR (DB), SAM (DEGREES) AND ERGAS OF EXACTFUSE AND OF THE STRONGEST RETRAINED NETWORK; THE BETTER VALUE OF EACH PAIR IN BOLD.

Image	PSNR		SAM		ERGAS	
Chikusei ×4, ExactFuse / MIMO-SST						
1	**58.86**	57.69	**0.664**	0.737	**1.562**	1.682
2	**60.00**	58.62	**0.671**	0.747	**0.994**	1.088
3	**60.67**	59.35	**0.615**	0.686	**1.104**	1.235
4	**57.01**	56.46	**0.874**	0.935	**1.340**	1.423
5	**57.24**	56.82	**0.731**	0.768	**1.303**	1.348
6	**57.43**	56.99	**0.675**	0.711	**1.361**	1.421
7	**56.35**	55.78	**0.823**	0.876	**1.404**	1.501
8	**56.83**	56.41	**0.741**	0.780	**1.602**	1.655
Pavia Centre ×4, ExactFuse / OTPNet						
1	**50.37**	49.69	**1.560**	1.608	**0.559**	0.594
2	**52.28**	51.92	**2.714**	2.795	**0.815**	0.844
3	**49.43**	48.72	**1.690**	1.741	**0.562**	0.602
4	**50.43**	49.98	**1.602**	1.661	**0.541**	0.571

HR-MSI, is retrained on HR-MSIs shifted by up to 1 px with the reference shifted alike, and scored against the reference in the HR-MSI's frame: unregistered, it reaches 51.59 and 46.20 dB for the sharp half-pixel shift (Chikusei, Pavia Centre), against 54.56 and 46.36 dB for ExactFuse after registration. Trained with the reference on the LR-HSI grid, as ExactFuse is, it does not learn to compensate the shift (38.79 dB unregistered). On Pavia Centre M2DTN is the most accurate method in the four rows with a bilinear shift (41.95 against 40.66 dB for registered ExactFuse at 0.5–1 px), whereas registered ExactFuse leads for sharp shifts and in every row on Chikusei.

f) Numerical check of the null-space split.: On the test images, on random images and on bicubic errors, the round trip through the eigenbasis, the idempotence and symmetry of the projector, and the sum of the observed and null energies deviate from exactness by at most 6e-7 (single precision); the null-space component passes through the actual $\mathcal{D}$ and $\mathcal{R}$, reflect border included, with a relative norm of at most 7e-8; and the null-space fraction of the unknowns is 87.89% on Chikusei and 89.67% on Pavia Centre, as Theorem 2(c) predicts.

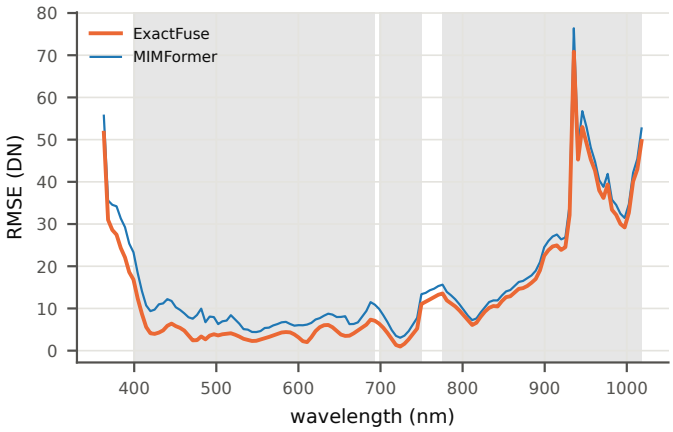


Fig. 9. RMSE (DN) of every band on the Chikusei test images. Shaded: wavelengths to which the WorldView-2 HR-MSI responds.

g) Corrupted observations.: An exact data step trusts its observations, so corrupted pixels pass into the reconstruction. With local scene changes in the HR-MSI, dead LR-HSI columns and 35 dB noise, ExactFuse trained over the operator family reaches 37.23 dB on Chikusei; trained on corrupted pairs with a confidence-weighted data step, which replaces each observation by $w\mathbf{Y} + (1-w)\mathbf{AX}$ with a learned per-pixel weight w, it reaches 47.74 dB (GSA 34.41 dB). The gain comes from training on corrupted pairs: the same network without the confidence step reaches 47.76 dB. What the confidence step adds is a map of the corrupted pixels. The map $1-w$ is a detector of the corrupted pixels: its ROC AUC against the corruption masks is 1.00 for the dead columns and 1.00 for the changed regions on Chikusei. It fails for changes in the four-band HR-MSI of Pavia Centre (AUROC 0.32), where the network still tolerates them (45.66 dB against 35.74 dB without confidence weighting) but does not localize them; dead columns are found there as well (1.00). A detector without learning localizes them. For consistent observations $\mathcal{D}\mathbf{Y}_\mathrm{M} = \mathcal{R}\mathbf{Y}_\mathrm{H}$ (Theorem 1(a)), so the energy of $\mathcal{D}\mathbf{Y}_\mathrm{M} - \mathcal{R}\mathbf{Y}_\mathrm{H}$ over the bands, with the estimated operators and normalized by its median, marks the low-resolution pixels where the two observations disagree. Its AUROC is 0.996 for the changed regions of Pavia Centre (0.980 with dead columns and 35 dB noise) and 0.965 for locally misregistered regions, where the confidence map reaches 0.47; on Chikusei it is 0.997 and 0.985. The residual localizes at the resolution of the LR-HSI and cannot tell which of the two observations is wrong, so it complements the learned confidence map.

J. One network for any sensor and scale factor

Every input of the prior lies in the hyperspectral bands, so one ExactFuse network can be trained on random multispectral responses (3–10 bands of random position and width, a new response for every batch) and applied to any sensor.

TABLE XVI
HR-MSI MISREGISTERED AGAINST THE LR-HSI AND THE REFERENCE (TEST PSNR, DB). REGISTERED: THE HR-MSI IS SHIFTED BACK BY THE OFFSET ESTIMATED FROM THE PAIR BEFORE FUSION (ANY METHOD). BILINEAR: THE SHIFT OF THE SIMULATION ALSO BLURS THE HR-MSI; SHARP: BAND-LIMITED SHIFT, AS A SENSOR SAMPLING A SHIFTED SCENE RECORDS IT, REGISTERED WITH THE SAME INTERPOLATION. EXACTFUSE IS TRAINED ON NOISE-FREE DATA, THE WIDER FAMILY ADDS NOISE, ANISOTROPIC BLURS AND GRID OFFSETS; MIMFORMER IS TRAINED ON NOISE-FREE DATA. M2DTN [23], A REGISTRATION-AWARE NETWORK, WARPS THE LR-HSI ONTO THE HR-MSI; IT IS TRAINED ON HR-MSIS SHIFTED BY UP TO 1 PX WITH THE REFERENCE SHIFTED ALIKE, AND SCORED AGAINST THE REFERENCE IN THE HR-MSI'S FRAME, WITHOUT REGISTRATION. BEST RESULT IN BOLD.

	Chikusei ×4						Pavia Centre ×4					
Misregistration	ExactFuse, unreg.	ExactFuse	ExactFuse, wider family	MIMFormer	GSA	M2DTN, unreg.	ExactFuse, unreg.	ExactFuse	ExactFuse, wider family	MIMFormer	GSA	M2DTN, unreg.
0.5–1 px (bilinear)	37.88	**49.74**	49.59	49.38	42.42	48.87	29.72	40.66	40.58	40.38	33.78	**41.95**
0.5 px + 35 dB	39.91	44.39	**45.87**	45.61	41.67	45.37	32.14	36.67	37.14	37.05	33.16	**37.37**
1 px + 35 dB	34.83	46.89	**50.46**	49.18	42.43	50.22	27.01	41.12	42.76	42.10	34.12	**44.80**
0.5 px, sharp	39.87	**54.56**	54.41	54.05	43.06	51.59	31.43	**46.36**	46.34	45.97	34.59	46.20
0.5 px, sharp + 35 dB	39.20	47.26	**51.57**	49.98	42.59	49.79	31.25	42.24	**44.79**	43.83	34.40	44.73
0.5 px + aniso. blur + 35 dB	39.53	43.48	**45.72**	45.49	41.59	45.39	31.55	35.12	36.90	36.92	33.11	**37.37**

Table XVII tests such a network on the responses of seven real sensors, none of them seen in training. It exceeds GSA by 9.9–13.5 dB on Chikusei and 10.5–15.6 dB on Pavia Centre. MIMFormer trained on the same random responses with the same input $\mathcal{R}^{+}\mathbf{Y}_{\mathrm{M}}$ stays 3.2–11.3 dB (Chikusei) and 3.6–7.2 dB (Pavia Centre) below it; unlike MIMFormer, ExactFuse uses the response of the sensor at hand in its data steps. Not retraining has a price: on WorldView-2 the sensor-independent network is 1.79 dB below ExactFuse trained for that sensor (1.23 dB on IKONOS for Pavia Centre) and 0.69–0.72 dB below MIMFormer trained for WorldView-2 or Sentinel-2. Trained for one sensor, ExactFuse stays ahead of MIMFormer trained for the same sensor, also with a three-band RGB camera, where the multispectral image measures the least (47.15 against 46.03 dB on Chikusei). The data steps also carry the scale factor, so ExactFuse can be trained over several of them, each batch with its own decimation and blur. One network trained at $\times 4$ and $\times 8$ is at most 0.14 and 0.15 dB below the networks trained for each scale (Chikusei, Pavia Centre) and still above the best retrained network of every scale; at $\times 16$, which it never saw, it leads that network by 1.49 and 1.12 dB.

The response models of Table XVII are built from published band edges. Table XVIII repeats the test with the measured responses of six other sensors (Landsat 8 OLI, Sentinel-2A and 2B, Sentinel-3 OLCI, PROBA-V and MODIS, from the filter functions of the 6S radiative-transfer code), none of them seen in training: the sensor-independent ExactFuse network exceeds GSA by 15.5–19.3 dB on Chikusei and MIMFormer trained the same way by 3.1–12.1 dB.

K. Complexity

Table XIX compares the parameters, operations, inference time and memory of all retrained networks on the same input. ExactFuse has 10.45 M parameters and needs 412 GFLOPs for a 256×256 Chikusei image, which it reconstructs in 221 ms (the base model: 296 GFLOPs, 181 ms) on one Tesla T4 with mixed precision (92 bands: 0.18 s); the self-ensemble doubles this. It is slower than most of the retrained networks, which have no data steps, but about ten times faster than DCT and DHIF-Net. Because the attention is computed across channels, memory and time grow linearly with the number of pixels.

The whole pipeline (Table XX) takes 0.39 s for a 256×256 Chikusei pair on a T4: 0.08 s for the operator estimate, 0.07 s for the eigendecompositions, which are computed once per image size, and 0.24 s for the network. The closed-form data step costs about as much as a gradient step (0.21 s), which reproduces the LR-HSI 6.7 dB worse, and less than a conjugate-gradient step (0.33 and 0.45 s with 5 and 10 iterations). The time grows about linearly with the number of pixels (5.30 s at 1024×1024, against 1.93 s for MIMFormer) and slowly with the number of bands.

VII. DISCUSSION AND LIMITATIONS

What the exact operators buy. The exact operators matter less for accuracy at the training operator than for what they make possible. Their separable structure gives the data step in closed form and the null space of fusion explicitly, which turns error analysis into a split between what contradicts the observations and what only a prior can supply; ExactFuse's error in the observed subspace is several times smaller than that of every retrained network. Because the operator enters only the data steps, it can be supplied, estimated from the test pair and replaced, and a single network covers a family of blurs, responses and noise levels. The same property makes the network depend on its operator: inside the family this is a strength, outside it a weakness, and the family used for training and estimation has to cover the sensors met in practice. The prior, in contrast, is a building block rather than a contribution: most of the margin over the retrained networks is already reached by the residual U-Net alone, and a convolutional mixer performs as well as transposed attention.

Limitations. All quantitative results with a reference rest on Wald's protocol, in which the observations are simulated from the reference, as in every benchmark of this field; in real data the border rule, the blur and the spectral response are only approximately known, and Theorems 1 and 2 assume that the operator family contains the true operator and that the blur is separable. Section VI-H shows the consequence: outside the family, an operator-conditioned network that trusts a wrong operator loses several decibels, more than a network trained on randomized data without an operator, and a misregistration of half a pixel defeats every method trained on registered pairs. Wider operator families, learned with the network or covered by the training distribution, and an explicit registration step are needed for real sensor pairs; registering the HR-MSI with the estimated offset recovers most of a global misregistration, but registration remains a separate step (a retrained registration-aware network, M2DTN [23], beats registered ExactFuse for

TABLE XVII

ONE NETWORK FOR ANY MULTISPECTRAL SENSOR AND FOR SEVERAL SCALE FACTORS (TEST PSNR, DB; CHIKUSEI AND PAVIA CENTRE). TOP: NETWORKS WITH THE SENSOR-INDEPENDENT INPUT $\mathcal{R}^{+}\mathbf{Y}_{\mathrm{M}}$ TRAINED ON RANDOM RESPONSES (3–10 BANDS OF RANDOM POSITION AND WIDTH, DRAWN ANEW FOR EVERY BATCH) AND TESTED ON THE RESPONSES OF SEVEN REAL SENSORS, NONE OF WHICH WAS SEEN IN TRAINING (THE PAIRS ARE SIMULATED FROM THE REFERENCE WITH THESE RESPONSES, NOT ACQUIRED BY THE SENSORS; M: BANDS INSIDE THE HYPERSPECTRAL RANGE), AGAINST GSA AND AGAINST NETWORKS TRAINED FOR THAT ONE SENSOR (RETR.). BOTTOM: EXACTFUSE TRAINED OVER $\times4$ AND $\times8$ AND TESTED AT $\times4$, $\times8$ AND $\times16$, AGAINST EXACTFUSE TRAINED AT EACH SCALE AND THE BEST RETRAINED NETWORK OF THAT SCALE. –: NOT RUN. BEST OF THE TWO SENSOR-INDEPENDENT NETWORKS IN BOLD.

	Chikusei $\times4$						Pavia Centre $\times4$					
			One network		Retr. for sensor				One network		Retr. for sensor	
Sensor	M	GSA	MIMFormer	ExactFuse	MIMFormer	ExactFuse	M	GSA	MIMFormer	ExactFuse	MIMFormer	ExactFuse
WorldView-2	8	43.20	46.17	**56.26**	56.96	58.05	6	34.69	46.37	**50.28**	–	–
IKONOS	4	42.64	50.33	**53.57**	–	–	4	34.78	45.55	**49.39**	49.63	50.63
QuickBird	4	42.70	50.74	**53.93**	–	–	4	34.66	45.80	**49.35**	–	–
PlanetScope	4	40.64	44.33	**53.98**	–	–	3	29.82	33.76	**40.95**	–	–
Landsat 8	5	40.67	42.85	**54.12**	–	–	3	29.82	34.00	**40.77**	–	–
Sentinel-2	10	43.43	49.47	**56.06**	56.79	58.22	7	35.24	43.45	**50.52**	–	–
RGB camera	3	34.49	36.92	**44.41**	46.03	47.15	3	29.77	35.15	**40.30**	–	–
Scale	GSA		ExactFuse, $\times4$+$\times8$		ExactFuse	best retr.	GSA		ExactFuse, $\times4$+$\times8$		ExactFuse	best retr.
$\times4$	43.20		57.96		58.05	57.27	34.78		50.47		50.63	50.08
$\times8$	38.69		56.89		56.95	56.11	30.24		49.96		50.05	49.51
$\times16$	35.20		55.97		56.11	54.48	27.07		49.38		49.25	48.25

TABLE XVIII

THE SENSOR-INDEPENDENT NETWORKS OF TAB. XVII (TRAINED ON RANDOM RESPONSES) ON SIX SENSORS WITH MEASURED SPECTRAL RESPONSES (6S FILTER FUNCTIONS), NONE SEEN IN TRAINING; THE PAIRS ARE SIMULATED FROM THE REFERENCE WITH THESE RESPONSES, NOT ACQUIRED BY THE SENSORS; M: BANDS INSIDE THE HYPERSPECTRAL RANGE. TEST PSNR (DB).

	Chikusei				Pavia Centre			
Sensor	M	GSA	MIMF.	ExactFuse	M	GSA	MIMF.	ExactFuse
Landsat 8 OLI	5	35.55	42.01	**54.07**	3	28.31	33.47	**40.73**
Sentinel-2A MSI	10	38.82	52.56	**55.89**	7	32.78	39.28	**50.19**
Sentinel-2B MSI	10	38.79	52.24	**55.87**	7	32.82	39.68	**50.23**
Sentinel-3 OLCI	21	37.07	49.67	**56.39**	13	29.50	35.61	**47.91**
PROBA-V	3	38.21	50.57	**53.69**	2	33.22	40.75	**44.74**
MODIS (Aqua)	11	36.15	44.75	**54.34**	9	29.84	36.31	**48.88**

TABLE XIX

COMPLEXITY FOR ONE 256×256 CHIKUSEI IMAGE (128 BANDS, $\times4$) ON ONE TESLA T4.

Method	Params (M)	GFLOPs	Time (ms)	Memory (GB)
ExactFuse	10.45	412	221	0.54
Base model	10.13	296	181	0.46
SSR-NET	0.44	58	11	0.11
PSRT	0.30	43	122	0.42
DCT	8.46	2640	2255	1.10
MIMFormer	16.46	224	66	0.46
DHIF-Net	48.67	14930	2807	2.09
SMGU-Net	9.50	2286	511	0.98
CLSNet	4.17	142	170	0.36
OTPNet	2.51	328	155	0.34
MIMO-SST	5.32	81	29	0.33
CaFormer	7.83	222	171	0.57

TABLE XX

END-TO-END COST ON ONE TESLA T4 (BATCH 1, CHIKUSEI BANDS AND A 256×256 HR-MSI UNLESS STATED; SECONDS PER IMAGE AND PEAK GPU MEMORY). THE DATA-STEP ROWS USE THE SAME NETWORK WITH ANOTHER DATA STEP.

Step	Time (s)	Memory (GB)
Operator estimate, isotropic grid	0.081	–
Operator estimate, wide family	0.600	–
Operator estimate, nonparametric kernel	0.0048	–
Eigendecompositions (once per image size)	0.069	–
Network, closed-form step with learned gains	0.241	0.54
Network, closed form	0.242	0.54
Network, gradient step	0.206	0.60
Network, conjugate gradient, 5 iterations	0.333	0.54
Network, conjugate gradient, 10 iterations	0.447	0.54
128×128: ExactFuse / MIMFormer / eigendecompositions	0.090 / 0.053 / 0.007	0.17 / 0.17
256×256: ExactFuse / MIMFormer / eigendecompositions	0.242 / 0.115 / 0.022	0.54 / 0.46
512×512: ExactFuse / MIMFormer / eigendecompositions	1.192 / 0.445 / 0.071	2.02 / 1.58
1024×1024: ExactFuse / MIMFormer / eigendecompositions	5.301 / 1.931 / 0.374	7.92 / 6.11
31 bands (256×256, $M = 8$): ExactFuse	0.187	0.29
64 bands (256×256, $M = 8$): ExactFuse	0.197	0.38
128 bands (256×256, $M = 8$): ExactFuse	0.244	0.56
224 bands (256×256, $M = 8$): ExactFuse	0.384	0.92

bilinear shifts on Pavia Centre), and the only real acquisitions we test are those of NEON, at two sites; trained on one, ExactFuse leads by 1.11 dB on the other, and the sensor tests use real responses on simulated pairs. The confidence-weighted step finds corrupted pixels in the LR-HSI and in the eight-band HR-MSI of Chikusei but does not localize changes in the four-band HR-MSI of Pavia Centre, which it only tolerates; a model-free residual detector localizes them, at the resolution of the LR-HSI. The equal-budget benchmark gives every method 4 h on one GPU; the authors' own recipes do not improve the baselines under this budget, and the strongest baselines are close to converged at the end of the budget (they gain at most 0.11 dB in its last quarter), two slow ones (DCT, DHIF-Net) are not, and an 8 h check narrows but keeps the gap. The 17-method comparison of Tables X and XI uses copied values on test images that follow the published protocol but are not the same pixels; our retrained baselines exceed their reported values by up to 3.6 dB, so these tables give context rather than a ranking, and the strongest reported methods could not be retrained because their code is not available. The test sets of Chikusei and Pavia Centre are small (8 and 4 images); significance is assessed on the pooled settings, and the 95% bootstrap intervals of the gain exclude zero in every setting.

VIII. Conclusion

We formulated hyperspectral and multispectral image fusion as maximum a posteriori estimation with parametric operators and unrolled it, and we proved three properties of the operators inside such a network: their parameters are identified by the observations, their separable structure gives the data step in closed form together with the null space of fusion, and an exact data step leaves the ground truth fixed while an inexact one moves it by a bounded amount. ExactFuse, built on these results, is the most accurate of six networks retrained under one budget on Chikusei and Pavia Centre, its error lies almost entirely in the null space that no observation constrains, and trained over a family of observation models with estimated operators it keeps its accuracy under changed blurs, spectral responses and noise. Operator families that cover non-separable blurs and registration errors, and real sensor pairs, are the main directions for future work.

Acknowledgment

The authors thank N. Yokoya and A. Iwasaki for making the Chikusei data available, P. Gamba of the University of Pavia for the Pavia Centre data, and Kaggle for the GPU time used to train and evaluate the models.

References

[1] B. Aiazzi, S. Baronti, and M. Selva, "Improving component substitution pansharpening through multivariate regression of MS +Pan data," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 45, no. 10, pp. 3230–3239, Oct. 2007.

[2] Q. Wei, N. Dobigeon, and J.-Y. Tourneret, "Fast fusion of multi-band images based on solving a Sylvester equation," *IEEE Transactions on Image Processing*, vol. 24, no. 11, pp. 4109–4121, Nov. 2015.

[3] X. Zhang, W. Huang, Q. Wang, and X. Li, "SSR-NET: Spatial–spectral reconstruction network for hyperspectral and multispectral image fusion," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 59, no. 7, pp. 5953–5965, 2021.

[4] S.-Q. Deng, L.-J. Deng, X. Wu, R. Ran, D. Hong, and G. Vivone, "PSRT: Pyramid shuffle-and-reshuffle transformer for multispectral and hyperspectral image fusion," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 61, pp. 1–15, 2023.

[5] J. Fang, J. Yang, A. Khader, and L. Xiao, "MIMO-SST: Multi-input multi-output spatial-spectral transformer for hyperspectral and multispectral image fusion," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 62, pp. 1–20, 2024.

[6] Y. Chen, R. Liu, W. K. Wong, and J. Wen, "A two-stage conditional diffusion model with differential attention for hyperspectral and multispectral image fusion," *IEEE Transactions on Image Processing*, vol. 35, pp. 8980–8995, 2026.

[7] Q. Xie, M. Zhou, Q. Zhao, Z. Xu, and D. Meng, "MHF-Net: An interpretable deep network for multispectral and hyperspectral image fusion," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 44, no. 3, pp. 1457–1473, Mar. 2022.

[8] T. Huang, W. Dong, J. Wu, L. Li, X. Li, and G. Shi, "Deep hyperspectral image fusion network with iterative spatio-spectral regularization," *IEEE Transactions on Computational Imaging*, vol. 8, pp. 201–214, 2022.

[9] D. Pineau, F. Orieux, and A. Abergel, "Exact solution for multispectral and hyperspectral fusion via hessian inversion," in *13th Workshop on Hyperspectral Imaging and Signal Processing: Evolution in Remote Sensing (WHISPERS)*, 2023, pp. 1–5.

[10] R. Li, L. Zhang, Z. Wang, and X. Li, "MIMFormer: Multiscale inception mixer transformer for hyperspectral and multispectral image fusion," *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, vol. 17, pp. 15 122–15 135, 2024.

[11] S. Peng, C. Guo, X. Wu, and L.-J. Deng, "U2Net: A general framework with spatial-spectral-integrated double U-Net for image fusion," in *Proceedings of the 31st ACM International Conference on Multimedia*, ser. MM '23, Oct. 2023, pp. 3219–3227.

[12] W. Yu, Z. Li, Q. Liu, and X. Sun, "OTPNet: ODE-inspired tuning-free proximal network for remote sensing image fusion," *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 39, no. 9, pp. 9671–9679, Apr. 2025.

[13] S. Wang, Q. Xie, Q. Zhao, and D. Meng, "AMHF-Net: A multispectral and hyperspectral image fusion network for arbitrary-band hyperspectral images," *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, vol. 19, pp. 12 367–12 383, 2026.

[14] Q. Zhang, J. Li, W. Zhu, M. Li, Q. Shen, and Y. Peng, "Structure-aware Riemannian flow matching for registration and fusion of hyperspectral and multispectral images," in *International Conference on Machine Learning (ICML)*, 2026.

[15] J. Yan, K. Zhang, Q. Sun, C. Ge, W. Wan, J. Sun, and H. Zhang, "Spatial–spectral unfolding network with mutual guidance for multispectral and hyperspectral image fusion," *Pattern Recognition*, vol. 161, p. 111277, May 2025.

[16] X. Cao, X. Wang, X. Dun, Y. Lian, X. Cheng, and X. Hao, "Cross-domain-aware deep unfolding transformer for hyperspectral image super-resolution," *Pattern Recognition*, vol. 172, p. 112374, 2026.

[17] S. W. Zamir, A. Arora, S. Khan, M. Hayat, F. S. Khan, and M.-H. Yang, "Restormer: Efficient transformer for high-resolution image restoration," in *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022, pp. 5718–5729.

[18] K. Zhang, L. Van Gool, and R. Timofte, "Deep unfolding network for image super-resolution," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2020, pp. 3214–3223.

[19] H. Yang, J. Liu, R. Dian, and S. Li, "Correlation-aware learning self-guiding network for hyperspectral fusion imaging," *Information Fusion*, vol. 133, p. 104263, 2026.

[20] Q. Ma, J. Jiang, X. Liu, and J. Ma, "Reciprocal transformer for hyperspectral and multispectral image fusion," *Information Fusion*, vol. 104, p. 102148, Apr. 2024.

[21] N. Yokoya and A. Iwasaki, "Airborne hyperspectral data over Chikusei," Space Application Laboratory, University of Tokyo, Tokyo, Japan, Tech. Rep. SAL-2016-05-27, May 2016.

[22] P. Gamba, "A collection of data for urban area characterization," in *Proc. IEEE Int. Geosci. Remote Sens. Symp. (IGARSS)*, vol. 1, 2004, pp. 69–72.

[23] W. Dong, Y. Xu, J. Qu, and S. Hou, "Learning multi-modal cross-scale deformable transformer network for unregistered hyperspectral image super-resolution," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 38, no. 2, 2024, pp. 1573–1581.

[24] F. Yasuma, T. Mitsunaga, D. Iso, and S. K. Nayar, "Generalized assorted pixel camera: Postcapture control of resolution, dynamic range, and spectrum," *IEEE Transactions on Image Processing*, vol. 19, no. 9, pp. 2241–2253, 2010.

[25] A. Chakrabarti and T. Zickler, "Statistics of real-world hyperspectral images," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2011, pp. 193–200.