

Exact Observation Operators in Deep Unfolding for Hyperspectral and Multispectral Image Fusion

Abstract—Deep unfolding networks for hyperspectral and multispectral image fusion alternate data-consistency steps with learned priors. Inside these steps the observation operators are usually either learned or implemented with the zero padding of convolution layers, whereas benchmark data are simulated with a different border rule. We examine what changes when the network’s operators are exactly those of the simulation, image border included. With exact operators the ground truth is a fixed point of every data-consistency step; with a zero-border or a learned operator it is not, and the step pushes it away by an amount that grows with the share of border pixels. We build PUFormer, an unfolded network that uses the exact blur–decimation and spectral operators and their exact adjoints, with a transposed-attention U-Net as prior. In an equal-budget benchmark, in which five published networks are retrained from their authors’ code under one pipeline on Chikusei, Pavia Centre, CAVE and Harvard at scale factors from 4 to 16, PUFormer has the highest PSNR in 2 of 2 settings. A module-by-module ablation with repeated seeds shows where the accuracy comes from: removing the data-consistency step costs most, whereas cross-stage memory, the type of attention and the number of stages change PSNR by at most 0.05 dB, so a one-stage model with a third of the parameters performs almost as well. At the training blur the exact operator gains only 0.08 dB over a zero border, but it raises the agreement of the reconstruction with its inputs by about 3 dB and is the only operator model that can be replaced at test time: given the true blur, PUFormer loses at most 0.13 dB for blur widths it was not trained on, against 0.88 dB for a zero border and 1.24 dB for a learned operator, and a blur estimated from the test pair itself performs almost as well as the true one. Under the protocol of a recent study comparing 17 methods, PUFormer reaches 58.20 dB on Chikusei $\times 4$ and 50.69 dB on Pavia Centre $\times 4$. Code, training logs and all result files are public.

Index Terms—Hyperspectral image fusion, multispectral image, deep unfolding, observation model, boundary conditions, super-resolution.

I. INTRODUCTION

HYPERSPECTRAL imagers record hundreds of narrow spectral bands but, for a fixed photon budget, at a coarse spatial resolution; multispectral imagers do the opposite. Hyperspectral and multispectral image fusion recovers a high-resolution hyperspectral image (HR-HSI) from one observation of each kind [1]–[3]. Classical methods formulate fusion as a regularized inverse problem, using component substitution [4], coupled matrix factorization [5], subspace regularization [6] or fast Sylvester-equation solvers [7]. Deep networks have since become the dominant approach, from convolutional reconstruction networks [8], [9] to transformers [10]–[13] and diffusion models [14]–[17].

E-mail: amarnath10chinu@gmail.com. Code, training logs and all result files are available at <https://github.com/Amarnath10i/Research-On-Hyperspectral-ImageFusion>.

Model-guided and unfolded networks [18]–[23] bring the observation model back into the network: each stage takes a step that reduces the mismatch between the current estimate and the two observations, and a learned module then refines the estimate. How faithfully the network reproduces the observation model has received less attention than the design of the learned module. In most of these networks the blur–decimation and spectral operators are learned convolution layers, or are implemented with the zero padding that convolution layers apply at the image border. The benchmark data, however, are simulated with a fixed blur and a different border rule. The network’s operator is then not the one that produced the data, and even the ground truth fails the network’s own consistency check. In simulated benchmarks this mismatch is avoidable, because the forward model is known exactly. A second difficulty concerns evaluation: reported numbers on the same dataset vary by tens of decibels across papers because of differences in sensor response, blur, data split, normalization and even metric definitions [3], so improvements are only meaningful within one fully specified protocol in which all methods are run.

This paper studies the operator question in a controlled way. We use an unfolded network, PUFormer, whose data-consistency steps apply exactly the simulation operators and their exact adjoints, border included, and compare it with variants that use a zero-border or a learned operator, or no data step at all, all trained with the same budget. We then compare PUFormer with published networks that we retrain ourselves under one protocol on four datasets and three scale factors. Our contributions are:

- 1) **The role of exact operators in unfolding.** We show that with exact operators the ground truth is a fixed point of every data-consistency step, and we quantify how far a zero-border or a learned operator moves it (Proposition 1). Experimentally, the exact operator yields a small but consistent accuracy gain, a reconstruction that agrees better with its inputs, and the ability to change the blur at test time, with a known or an estimated kernel, at almost no cost.
- 2) **A module-by-module ablation,** with repeated seeds for the key rows, that identifies which parts of an unfolded network matter: the data-consistency step and the exact operator do, whereas cross-stage memory, the type of attention and the number of stages hardly do. A one-stage variant with a third of the parameters is almost as accurate.
- 3) **An equal-budget benchmark** in which SSR-NET, PSRT, DCT, MIMFormer and the unfolded DHIF-Net are retrained from their authors’ code alongside PU-

Former on Chikusei, Pavia Centre, CAVE and Harvard at $\times 4$, $\times 8$ and $\times 16$, with one data pipeline, one training budget and one metric implementation, complemented by a comparison with the 17 methods reported in [15].

- 4) **An analysis of the remaining errors**, showing that the Q2n shortfall on Chikusei comes from bands that the multispectral sensor does not cover and the one on Pavia Centre from dark water pixels.

Code, training logs and all result files are publicly available.

II. RELATED WORK

A. Model-based fusion

Component substitution methods such as GSA [4] inject multispectral detail into the upsampled hyperspectral bands. Unmixing-based methods such as CNMF [5] factor both observations with shared endmembers, and subspace methods such as HySure [6] regularize the HR-HSI within a low-dimensional spectral subspace, whose dimension can be estimated with HySime [24]. FUSE [7] solves the resulting quadratic problem in closed form through a Sylvester equation. These methods use the forward model explicitly but rely on hand-crafted priors.

B. Supervised deep fusion

SSR-NET [8] and GuidedNet [9] learn a direct mapping from the two observations. Transformers model long-range spatial and spectral dependencies: PSRT [10] uses pyramid shuffle-and-reshuffle attention, MIMO-SST [11] multi-input multi-output supervision, LGCT [12] local-global collaboration, DSPNet [25] a dual spatial-spectral pyramid, SSDT [13] dilated attention, and MIMFormer [26], AMSF-Net [27], ASMNet [28] and DPFormer [29] multiscale spectral-spatial designs. Recent work also studies state-space models [30], [31], mixture-of-experts routing [32] and modality interaction [33]. General image-fusion frameworks such as U2Net [34] and OTPNet [35] have also been applied to this task.

C. Model-guided and unfolded networks

Algorithm unrolling [36]–[38] turns the iterations of an optimization method into network stages. MHF-Net [18], MoG-DCN [19] and DHIF-Net [20] unfold fusion-specific solvers, MHF-Net and DHIF-Net with operators learned as convolution layers, DUNET [21] adds transformer priors for unregistered inputs, HMPNet [22] extends unfolding to three modalities, and an information-aware transformer-based unfolding network (ITU-Net) [23], a mutual-guidance unfolding network (SMGU) [39], the perceptive spectral transformer unfolding network (PSTUN) [40], which also handles arbitrary scales, and the cross-domain-aware unfolding transformer CaFormer [41] strengthen the learned proximal steps with transformer priors. PUFormer belongs to this family, and none of its prior's components is new: passing features from one stage to the next follows memory-augmented unfolding [42], spectral-wise attention inside an unfolded network appears in DAUHST [43], and the transposed attention is that of

Restormer [44]. A single unfolded network that takes the blur kernel as an input, and therefore handles kernels it was not trained on, was shown for single-image super-resolution by USRNet [45].

D. Boundary conditions

In classical deconvolution, an image border handled differently by the model and by the data causes ringing that spreads inward, and reflexive (Neumann) boundary conditions, i.e., reflect padding, are a standard remedy [46], [47]. In unfolded fusion networks the operators are either learned or implemented with the zero padding of convolution layers, while benchmark data are simulated with other border rules; to our knowledge, the effect of this mismatch has not been examined. Our contribution is on this side: an unfolded network whose operators, border included, are exactly those of the simulation, and an ablation that separates the exact operator from a zero-border and a learned operator under the same budget.

E. Generative fusion

Diffusion models have been conditioned on the two observations (DDIF [14]), constrained by the observation model during sampling (CDPS [16]), restricted to a learned subspace (ARGS-Diff [17]) or combined with differential attention in a two-stage conditioning scheme [15]. Null-space formulations [48] make the data-consistency part of the reconstruction exact. The two-stage conditional diffusion model of Chen *et al.* [15] holds the best published Chikusei $\times 4$ results in our survey and is our main reference; its comparison also includes the correlation-aware network CLSNet [49] and the dual cross transformer (DCTransformer) [50], denoted DCT in their table.

III. OBSERVATION MODEL AND EXACT OPERATORS

Let $\mathbf{X} \in \mathbb{R}^{B \times H \times W}$ be the HR-HSI with B bands ($B = 128$ for Chikusei, 92 for Pavia Centre and 31 for CAVE and Harvard). The two observations are

$$\mathbf{Y}_H = \mathcal{D}\mathbf{X}, \quad \mathbf{Y}_M = \mathcal{R}\mathbf{X}, \quad (1)$$

where $\mathcal{D} = \mathcal{SKP}$ applies, band by band, reflect padding $\mathcal{P}$ of width p , a valid correlation $\mathcal{K}$ with a normalized $(2p+1) \times (2p+1)$ Gaussian kernel, and decimation $\mathcal{S}$ that keeps every s -th pixel starting at index zero, so $\mathbf{Y}_H \in \mathbb{R}^{B \times H/s \times W/s}$. For a scale factor s we use $p = s-1$ and a standard deviation of $s/2$, which at $s = 4$ is the 7×7 , $\sigma = 2$ blur of [15]; the kernel thus grows with the factor, so that the LR-HSI stays band-limited at $\times 8$ and $\times 16$. The spectral operator $\mathcal{R} \in \mathbb{R}^{M \times B}$ integrates the spectrum with the $M = 8$ WorldView-2 band responses [51] for Chikusei, the $M = 4$ IKONOS band responses [52] for Pavia Centre and the $M = 3$ Nikon D700 responses for CAVE and Harvard (Section V). Because the manufacturers publish the responses only as plots together with band edges, we model each band as a flat top with logistic flanks (4 nm roll-off) that pass exactly through the published edges, at 5% response for WorldView-2, whose edges are tabulated at that level, and at 50% for the IKONOS band limits (blue 445–516, green 506–595, red 632–698 and near infrared 757–853 nm), and normalize every band to unit sum.

Exact adjoint. Every gradient step on $\|\mathcal{D}\mathbf{X} - \mathbf{Y}_H\|^2$ requires $\mathcal{D}^\top = \mathcal{P}^\top \mathcal{K}^\top \mathcal{S}^\top$. $\mathcal{S}^\top$ inserts zeros and $\mathcal{K}^\top$ is a full convolution; together they are a strided transposed convolution with an output padding of $s-1$. The adjoint of reflect padding folds the padded margins back onto the pixels they were copied from. For a padded row index $j \in [0, H+2p]$ the source row is $p-j$ for $j < p$, $j-p$ in the interior and $2H-2+p-j$ in the bottom margin, hence

$$(\mathcal{P}^\top \mathbf{g})_i = \mathbf{g}_{i+p} + [1 \leq i \leq p] \mathbf{g}_{p-i} + [H-1-p \leq i \leq H-2] \mathbf{g}_{2H-2+p-i}, \quad (2)$$

applied to rows and then to columns. With (2), $\mathcal{D}$ in the network is exactly the simulation operator, and the inner-product test $\langle \mathcal{D}\mathbf{X}, \mathbf{Y} \rangle = \langle \mathbf{X}, \mathcal{D}^\top \mathbf{Y} \rangle$ holds to a relative error below 10^{-13} in double precision at every scale factor. $\mathcal{R}^\top$ is the transpose of the response matrix.

Why the operator matters. Let $\tilde{\mathcal{D}}$ denote the blur-decimation operator used inside a network, and consider the data-consistency step on the LR-HSI term,

$$\mathbf{Z} = \mathbf{X}' - \alpha \tilde{\mathcal{D}}^\top (\tilde{\mathcal{D}}\mathbf{X}' - \mathbf{Y}_H), \quad \alpha > 0, \quad (3)$$

applied to an estimate $\mathbf{X}'$. A step of this kind should leave a correct estimate unchanged.

Proposition 1: Let $\mathbf{Y}_H = \mathcal{D}\mathbf{X}$. If $\tilde{\mathcal{D}} = \mathcal{D}$, the ground truth $\mathbf{X}$ is a fixed point of (3) for every α . For any other $\tilde{\mathcal{D}}$, the step moves $\mathbf{X}$ by $\mathbf{b} = \alpha \tilde{\mathcal{D}}^\top (\tilde{\mathcal{D}} - \mathcal{D})\mathbf{X}$, with $\|\mathbf{b}\| \leq \alpha \|\tilde{\mathcal{D}}\| \|(\tilde{\mathcal{D}} - \mathcal{D})\mathbf{X}\|$. If $\tilde{\mathcal{D}}$ differs from $\mathcal{D}$ only in its border rule (e.g., zero instead of reflect padding), $(\tilde{\mathcal{D}} - \mathcal{D})\mathbf{X}$ is nonzero only on the LR pixels whose kernel support reaches into the padding, i.e., within $\lceil p/s \rceil$ pixels of the LR border.

Proof: With $\mathbf{X}' = \mathbf{X}$ the step gives $\mathbf{Z} - \mathbf{X} = -\alpha \tilde{\mathcal{D}}^\top (\tilde{\mathcal{D}}\mathbf{X} - \mathcal{D}\mathbf{X}) = -\mathbf{b}$, which vanishes when $\tilde{\mathcal{D}} = \mathcal{D}$; the bound follows from the submultiplicativity of the operator norm. Both operators compute the same valid correlation of the padded image followed by the same decimation; they differ only in the padded values, which enter an output pixel only if its $(2p+1) \times (2p+1)$ support overlaps the padding, i.e., if the output pixel lies within $\lceil p/s \rceil$ pixels of the LR border. ■

The fraction of affected LR pixels in an $h \times w$ LR image is about $2\lceil p/s \rceil (h+w)/(hw)$, so a mismatched border is a small perturbation for a 64×64 LR image but not for the 16×16 LR patches on which networks are trained. A learned operator, in contrast, differs from $\mathcal{D}$ everywhere. In both cases the learned prior has to cancel $\mathbf{b}$, a bias that depends on the image content, in addition to its actual task. The same argument explains why an exact operator can be replaced at test time: if the data were produced by a different blur $\mathcal{D}'$, using $\tilde{\mathcal{D}} = \mathcal{D}'$ restores the fixed-point property for the new data, whereas a network that has absorbed a mismatched or learned operator into its prior cannot be corrected in this way. Section VI-A measures $\mathbf{b}$ and its consequences.

IV. PUFORMER

A. Unfolded iterations

PUFormer (Fig. 1) starts from the bicubic upsampling $\mathbf{X}_0$ of $\mathbf{Y}_H$ and performs $K = 3$ stages. Stage k first takes a gradient

Algorithm 1 PUFormer inference

Require: $\mathbf{Y}_H, \mathbf{Y}_M$, exact operators $\mathcal{D}, \mathcal{D}^\top, \mathcal{R}, \mathcal{R}^\top$

- 1: $\mathbf{X}_0 \leftarrow \text{bicubic}_{\times s}(\mathbf{Y}_H), \mathbf{F}_0 \leftarrow \emptyset$
- 2: **for** $k = 1, \dots, K$ **do**
- 3: $\mathbf{Z}_k \leftarrow \text{data step (4)}$ ▷ fp32
- 4: $\mathbf{r}_k \leftarrow s^2 \mathcal{D}^\top (\mathcal{D}\mathbf{Z}_k - \mathbf{Y}_H)$
- 5: $(\Delta_k, \mathbf{F}_k) \leftarrow \mathcal{P}_k([\mathbf{Z}_k, \mathbf{Y}_M, \mathbf{r}_k], \mathbf{F}_{k-1})$ ▷ fp16
- 6: $\mathbf{X}_k \leftarrow \mathbf{Z}_k + \Delta_k$
- 7: **end for**
- 8: **return** $\mathbf{X}_K$

step on both data terms,

$$\mathbf{Z}_k = \mathbf{X}_{k-1} - s^2 \alpha_k \mathcal{D}^\top (\mathcal{D}\mathbf{X}_{k-1} - \mathbf{Y}_H) - \beta_k \mathcal{R}^\top (\mathcal{R}\mathbf{X}_{k-1} - \mathbf{Y}_M), \quad (4)$$

with learnable step sizes $\alpha_k = \text{softplus}(a_k)$ and $\beta_k = \text{softplus}(b_k)$. The factor s^2 compensates the $1/s^2$ gain of $\mathcal{D}^\top \mathcal{D}$ on smooth content under $\times s$ decimation, so that both steps start at a comparable scale. The prior then predicts a residual,

$$(\Delta_k, \mathbf{F}_k) = \mathcal{P}_k([\mathbf{Z}_k, \mathbf{Y}_M, s^2 \mathcal{D}^\top (\mathcal{D}\mathbf{Z}_k - \mathbf{Y}_H)], \mathbf{F}_{k-1}), \quad (5)$$

$$\mathbf{X}_k = \mathbf{Z}_k + \Delta_k,$$

where the third input tells the prior where the updated estimate still disagrees with $\mathbf{Y}_H$, and $\mathbf{F}_{k-1}$ are the final decoder features of the previous stage (cross-stage memory, merged by a 1×1 convolution). The last convolution of every prior is initialized to zero, so training starts from the pure physics iteration. The image path, including all operator applications, stays in 32-bit floating point; only the priors run in 16-bit mixed precision [53].

B. Transposed-attention prior

Each prior is a three-level U-Net [54] with widths 48, 96 and 192, two, three and four transformer blocks on the encoder levels and bottleneck, mirrored decoder levels, and two refinement blocks. A block follows Restormer [44]: multi-Dconv-head transposed attention computes queries, keys and values with 1×1 and depth-wise 3×3 convolutions, normalizes each feature map over space, and forms a $c' \times c'$ attention matrix across channels per head (1, 2 and 4 heads on the three levels) with a learned temperature. Attention is thus computed between spectral feature maps rather than between spatial positions, its cost is linear in the number of pixels, and the receptive field of every block covers the whole image. A gated depth-wise feed-forward network with expansion 2.66 follows. The three priors hold 3.37–3.38 M parameters each; the whole network has 10.13 M parameters for Chikusei and 9.99 M for Pavia Centre, whose input and output convolutions are smaller.

C. Training objective

With ground truth $\mathbf{X}$, the loss is

$$\mathcal{L} = \|\mathbf{X}_K - \mathbf{X}\|_1 + 0.1 \sum_{k < K} \|\mathbf{X}_k - \mathbf{X}\|_1 + \lambda_{\text{SAM}} \text{SAM}(\mathbf{X}_K, \mathbf{X}) + \lambda_S (1 - \text{SSIM}(\mathbf{X}_K, \mathbf{X})), \quad (6)$$




Fig. 1. PUFormer. Top: three unfolded stages, each an exact data-consistency step on both observations followed by a learned residual prior; decoder features are passed to the next stage’s prior (memory). Bottom: the prior of stage k , a three-level U-Net of transposed-attention transformer blocks (channel counts for Chikusei; for Pavia Centre the input has 188 and the output 92 channels).

with mean absolute errors, deep supervision of the intermediate stages, the mean spectral angle in radians, and a band-wise Gaussian SSIM [55] whose data range is the maximum of each training patch. We use $\lambda_{\text{SAM}} = 0.05$ and $\lambda_S = 0.1$; both terms are computed in 32-bit precision. An exponential moving average of the weights (decay 0.999) [56] is used for validation and testing.

D. Self-ensemble

At test time we average the prediction for the input with the transposed prediction for the transposed input [57]. Transposition preserves the decimation phase, so both passes see observations generated by exactly the same operator. Flips and rotations do not: flipping a 256-pixel row maps sampled position $4m$ to $255 - 4m$, which has phase three, so we do not use them.

V. EXPERIMENTAL SETUP

A. Data and protocol

Chikusei [58] is an airborne scene of 2517×2335 pixels with 128 bands from 363 to 1018 nm. The split sizes follow [15]: eight 256×256 test images, 64 validation patches and 832 training patches of 64×64 . Because their locations are not specified there, we fix them as follows. We take the central 2048×2048 crop, which contains no no-data border, and divide it by its global maximum (15,133 digital numbers, DN). The test images are the eight 256×256 tiles of the top 256 rows, the validation patches tile rows 272–400, and training draws random 64×64 crops with dihedral augmentation from rows 416–2048 (800 non-overlapping patch equivalents). Sixteen-row gaps keep the three regions disjoint.

Pavia Centre [59] was recorded by the ROSIS sensor over the centre of Pavia, Italy; the distributed scene has 1096×715 pixels at 1.3 m and 102 bands between 430 and 860 nm. For this dataset, [15] uses the benchmark files released with PSRT [10]: 92 bands, a four-band multispectral image, four

256×256 test images and 142 validation and 1284 training patches of 64×64 . These files are distributed through a file-sharing service that we could not access, so we rebuild the benchmark from the original scene to the same specification. We keep bands 11–102 (473–860 nm if the 102 bands are taken as evenly spaced over 430–860 nm), divide by the maximum of the retained bands (8,000 DN), apply the same blur and decimation as for Chikusei, and simulate the HR-MSI with the IKONOS responses. The test images are the four 256×256 tiles of the top-left 512×512 region. The 63 validation patches lie on a 32-pixel grid in rows 528–656, and training draws random 64×64 crops with dihedral augmentation from rows 672–1096 and from the strip to the right of the test region (rows 0–512, columns 528–715), 97 non-overlapping patch equivalents in total. Gaps of 16 pixels keep the regions disjoint. As a check on the rebuilt data, our GSA implementation reaches 34.78 dB on our test images, against 35.35 dB reported for GSA in [15].

For all datasets we follow Wald’s protocol [60]: the original image serves as ground truth and both observations are simulated from it, and every image, including every training crop, is degraded on its own with (1). Table I summarizes the two remote-sensing protocols.

CAVE and Harvard. To test the method beyond airborne scenes we add the two indoor/outdoor benchmarks most used in the fusion literature. CAVE [61] has 32 scenes of 512×512 pixels with 31 bands (400–700 nm). We use the common 20/12 split [18], [50]: the first 20 scenes in alphabetical order train, and the last 12, as full 512×512 images, test. Harvard [62] has 50 images of 1040×1392 pixels with 31 bands (420–720 nm), split into 30 training and 20 test images; each image is divided by its own maximum, and the central 512×512 crop of each test image is tested. In both cases the bottom-right corner of every training image is held out for validation (80 and 60 patches of 64×64) and the HR-MSI is simulated with the Nikon D700 response used throughout this literature [18], [50]. RMSE is reported on the 8-bit scale for these two datasets.

Scales. Besides $\times 4$ we evaluate $\times 8$ on all four datasets and

TABLE I
THE TWO $\times 4$ PROTOCOLS (DEGRADATION AND SPLIT SIZES OF [15]; IMAGE LOCATIONS OURS)

Item	Chikusei	Pavia Centre
HR-HSI	central 2048 ² crop, 128 bands	1096 $\times$ 715, bands 11–102 (92)
Normalization	$\div 15,133$ DN	$\div 8,000$ DN
LR-HSI	Gaussian 7×7 , $\sigma = 2$; $\times 4$ decimation; reflect border	
HR-MSI	WorldView-2, 8 bands	IKONOS, 4 bands
Test	8 of 256 ² (rows 0–256)	4 of 256 ² (rows, columns 0–512)
Validation	64 of 64 ² (rows 272–400)	63 of 64 ² (rows 528–656)
Training	random 64 ² crops, rows 416–2048	random 64 ² crops, see text
Selection	validation PSNR only	

$\times 16$ on Chikusei and Pavia Centre, with the blur defined in Section III ($(2s - 1) \times (2s - 1)$ Gaussian, standard deviation $s/2$). With a fixed 7×7 kernel the LR-HSI would be strongly aliased at $\times 16$, and the normalization of the data step in (4) would no longer hold, because most HR pixels would not lie under any kernel. At $\times 16$ a 64×64 patch leaves a 4×4 LR-HSI, so all methods train on 128×128 patches (batch 4) there.

Equal-budget benchmark. Reported numbers come from different splits, normalizations and metric code (Section V-B), so for the main comparison we retrain every learned method ourselves under one pipeline: SSR-NET [8], PSRT [10], DCT [50], MIMFormer [26] and DHIF-Net [20], built from their authors’ code¹. They cover a CNN, a pyramid window transformer, the dual-cross transformer that is among the strongest entries of [15], a multiscale mixer, and a model-guided unfolding network whose blur, decimation and spectral operators are learned convolution layers. Every method, PUFormer included, is trained for the same 4 h on one Tesla T4, with the same random patches and augmentation, batch size, EMA weights and validation-only model selection. The published networks keep their official optimizer settings (Adam, L1 loss, learning rate 10^{-4} , or 3×10^{-4} for DHIF-Net, no weight decay or gradient clipping) and are trained in 32-bit precision as released; DCT and MIMFormer use gradient checkpointing to fit into the 16 GB of a T4, which does not change what they compute. The test metrics are computed by the same code for all methods.

B. Metrics

We report the eight metrics of [15], averaged over the test images: PSNR; SSIM [55]; the spectral angle mapper (SAM, degrees, averaged over the pixels whose ground-truth and estimated spectra are both nonzero) [63]; ERGAS with the resolution ratio s [64]; Q2n [65] on 32×32 blocks, computed on the image in DN (on the 16-bit scale for CAVE and Harvard) exactly as in the reference implementation of the pansharpening toolbox [66], [67]; the band-averaged correlation coefficient (CC); the spatial correlation coefficient

¹PSRT: github.com/Deng-shangqi/PSRT (commit 9c2be90);
DCT: github.com/qingma2016/DCTransformer (bd75b15); MIM-
Former: github.com/meiruni/MIMFormer (0904158); DHIF-Net:
github.com/TaoHuang95/DHIF-Net (6434366). Only the numbers of
input and output bands and the scale factor are changed; DHIF-Net’s code
covers $\times 8$ and $\times 16$, and at $\times 4$ we use one of the two stride-4 layers of its
 $\times 16$ operator.

TABLE II
TRAINING RUNS

	Chikusei		Pavia Centre
	First run	v3 (final)	final
Initialization	random	first run	random
Border of $\mathcal{D}$	zero padding	reflect (exact)	reflect (exact)
Loss	L1 + 0.1 deep sup.	(6)	(6)
GPUs $\times$ batch	1×8	2×8	2×8
Peak learning rate	3×10^{-4}	1.5×10^{-4}	3×10^{-4}
Iterations	88,946	94,450	93,850
Training patches	0.71 M	1.51 M	1.50 M
Wall-clock time	9.5 h	10.9 h	10.2 h
Selected at	best val. (58.60 dB)	it. 44,000 (58.72 dB)	it. 43,000 (48.58 dB)

(SCC) [68] of Sobel gradient magnitudes as in the same toolbox; and RMSE in DN (on the 8-bit scale for CAVE and Harvard). The headline PSNR is $10 \log_{10}(1/\text{MSE})$ over the whole normalized image, i.e., its peak is the dataset maximum. The PSNR values of [15] also use a fixed peak: for every row of their Chikusei table, $\text{RMSE} \cdot 10^{\text{PSNR}/20}$ is between 16.6k and 17.1k DN. Since our peak is lower (15,133 DN), our headline PSNR is conservative; with a 16.7k DN peak it would be 0.86 dB higher. For Pavia Centre the same product lies between 7.7k and 8.4k DN (median 8.07k), close to our peak of 8,000 DN, so there the PSNR values are on the same scale. The SSIM definition of [15] is not stated. We report SSIM as in the code released with PSRT [10], whose benchmark data [15] adopts (Gaussian 11×11 window, $\sigma = 1.5$, data range 1, i.e., the dataset peak), and additionally a stricter variant whose data range is each image’s own maximum. We also report observation consistency, $\text{PSNR}(\mathcal{D}\hat{\mathbf{X}}, \mathbf{Y}_H)$ and $\text{PSNR}(\mathcal{R}\hat{\mathbf{X}}, \mathbf{Y}_M)$.

C. Implementation

PUFormer is implemented in PyTorch [69] and trained on Kaggle NVIDIA Tesla T4 GPUs (Table II). On Chikusei we train in two runs. The first run trains from scratch on one GPU with zero-padded operators and the L1 losses only. The final model (v3) fine-tunes it with the exact reflect-border operators and the full loss (6) on two GPUs with distributed data parallelism [70]. On Pavia Centre, a single run trains the model from scratch with the exact operators and the full loss on two GPUs, within one 12-hour session that also runs all tests. All runs use AdamW [71] ($\beta = (0.9, 0.99)$), weight decay 10^{-4} , a 1000-iteration linear warm-up followed by a cosine decay [72] to 10^{-3} of the peak rate that ends at the time budget, and gradient-norm clipping at 0.1. The weights with the best validation PSNR are tested once. The ablation and benchmark runs (Sections VI-A and VI-B) train from scratch with these settings, the full loss and a single GPU for 4 h each; their models are evaluated without self-ensemble.

VI. RESULTS

A. Ablation study

Table III changes one module of the full model at a time and retrains each variant from scratch with the same budget (4 h on one T4, about 40,000 iterations for the full model). On Pavia Centre the full model and the zero-border variant




Fig. 2. Test PSNR when the blur of the test data differs from the training blur ($\sigma = 2$, dotted line), for the exact operator with and without the true blur supplied at test time, the zero-border operator with the true blur supplied, the learned operator and the model without data step.

were trained with three seeds; their PSNR varies by only 0.02 and 0.03 dB across seeds, which sets the scale for the other differences.

Data step and operator. Removing the data-consistency step together with the residual input (row (b)) costs 0.17 dB and lowers the agreement of the reconstruction with the LR-HSI from 77.6 to 68.9 dB. Keeping the exact kernel but replacing the reflect border by zero padding (row (d)) costs 0.08 dB on average over the three seeds, and each zero-border seed is below each full-model seed. A blur kernel and spectral response learned with the network (row (e)), the operator model of DHIF-Net and MHF-Net, cost 0.06 dB. Both variants reproduce the LR-HSI about 3 dB worse than the exact model (74.5 and 73.6 against 77.6 dB). This is what Proposition 1 predicts: their data steps move the ground truth, and the prior has to compensate. On the test images the displacement $\|\mathbf{b}\|$, relative to $\|\alpha \hat{\mathcal{D}}^\top \hat{\mathbf{D}} \mathbf{X}\|$, is ??% for the zero border and ??% for the learned operator, and on 64×64 crops, where border pixels are a larger share, ??% and ??%.

The accuracy gap at the training blur is small; it becomes large when the blur changes (last column of Table III and Fig. 2). Given the true blur at test time, the exact model loses at most 0.13 dB for $\sigma \in \{1.5, 2.5, 3.0\}$, the zero-border model with the same swap 0.88 dB, and the learned operator, which cannot be replaced, 1.24 dB. Without the swap the exact model loses 7.50 dB: it is the most sensitive variant to a wrong operator, and the only one that fully recovers when given the right one.

Prior. Cross-stage memory (row (f)), the residual input (row (c)), the type of token mixer (rows (j) and (k)) and the number of stages (rows (g)–(i)) change the PSNR on Pavia Centre by at most 0.05 dB, comparable to the spread across seeds. A one-stage model with 3.33 M parameters is within 0.02 dB of the full model with 9.99 M; part of this parity comes from the fixed time budget, in which the smaller model completes almost three times as many iterations. The SAM and SSIM loss terms (row (l)) do not improve the PSNR, whereas deep supervision of the intermediate stages does (row (m); 0.12 dB lower without it). We keep the full model for all other experiments, but the ablation shows that PUFormer’s accuracy comes from its exact data-consistency steps rather than from its transformer prior.

B. Equal-budget benchmark

Tables IV–VII compare PUFormer with the five retrained networks, and Fig. 3 shows the PSNR against the scale factor;

the remaining metrics are listed in Table XV. Across the 2 dataset and scale combinations, PUFormer has the highest PSNR in 2, the lowest SAM in 2 and the lowest ERGAS in 2. On Chikusei $\times 4$ it reaches 57.95 dB, 0.99 dB above the best retrained network, and on Pavia Centre $\times 4$ 50.63 dB (+1.00 dB), with 10% and 11% lower RMSE.

Two checks address the fairness of a fixed budget. First, the retrained baselines reach or exceed the values reported for them in [15] although our data and splits differ: on Chikusei $\times 4$, for example, MIMFormer reaches 56.96 dB against 53.36 dB reported and PSRT 53.99 dB against 50.67 dB, so our pipeline does not weaken them. Second, the budget favors fast networks: DCT completes only 4,800 iterations in 4 h on Pavia Centre, against 37,450 for PUFormer and 97,400 for MIMFormer. Doubling the budget to 8 h for PUFormer and the strongest competitor, MIMFormer, gives ?? and ?? dB on Pavia Centre $\times 4$, a gap of ?? dB. On Chikusei $\times 4$, PUFormer is more accurate than the best retrained network on all test images (one-sided Wilcoxon signed-rank test, $p = 3.9e - 03$).

C. Comparison with reported results on Chikusei

The comparisons in this and the next subsection place PUFormer against the 17 methods tabulated in [15]. They are less controlled than Section VI-B: the other rows are copied from [15], our test images follow its protocol but are not necessarily the same pixels, and the PUFormer models here were trained longer (Table II). Table VIII lists all Chikusei results of [15] together with PUFormer. The compared methods are GSA [4], CNMF [5], FUSE [7], U2Net [34], PSRT [10], DCT [50], DDIF [14], MIMFormer [26], MIMO-SST [11], OTPNet [35], SMGU [39], AMSF [27], SSdT [13], ASMNet [28], DPFormer [29], CLSNet [49] and the two-stage conditional diffusion model [15]; their values are copied from [15].

PUFormer ranks first of 18 on PSNR, SSIM (PSRT definition), SAM, ERGAS, CC, SCC and RMSE, with or without the self-ensemble. Relative to the best reported method, PSNR improves by 2.01 dB, RMSE by 26.4%, ERGAS by 12.2% and SAM by 1.1%. It ranks tenth on Q2n, which Section VI-I analyzes. Because the RMSE is expressed in sensor digital numbers, it does not depend on any choice of peak, and it gives the most direct comparison across papers. Table IX lists our own runs of reference methods under the identical pipeline and all PSNR and SSIM variants. Our GSA implementation scores 4.4 dB below the value reported in [15], so our protocol does not favor classical methods.

D. Comparison with reported results on Pavia Centre

Table X compares PUFormer with the 17 methods of [15] on Pavia Centre. Trained from scratch in one session, PUFormer reaches 50.69 dB PSNR, 3.53 dB above the best reported value, and lowers the RMSE by 34.1% and ERGAS by 25.3%. It ranks first on six of the eight metrics (PSNR, SSIM in the PSRT definition, ERGAS, CC, SCC and RMSE), but 13th on SAM (1.883° against 1.4542°) and 16th on Q2n (0.9753 against 0.9987). With the stricter per-image data range its SSIM is 0.9970, below the reported 0.9972. The PSNR gain

TABLE III

ABLATION STUDY: EACH ROW CHANGES ONE MODULE OF THE FULL MODEL (BOLD; ✓: USED, -: REMOVED) AND IS TRAINED FROM SCRATCH WITH THE SAME 4 H BUDGET ON ONE T4 (TEST SET, SINGLE PASS). Δ : PSNR CHANGE AGAINST ROW (A) OF THE SAME DATASET. CONS.: $\text{PSNR}(\mathcal{D}\hat{\mathbf{X}}, \mathbf{Y}_H)/\text{PSNR}(\mathcal{R}\hat{\mathbf{X}}, \mathbf{Y}_M)$ IN dB, WITH THE SIMULATION OPERATORS. SWAP: WORST TEST PSNR OVER BLUR $\sigma \in \{1.5, 2.5, 3.0\}$ WHEN THE TRUE BLUR IS GIVEN TO THE MODEL AT TEST TIME, FOR VARIANTS WITH AN EXPLICIT OPERATOR. ROWS WITH $\pm$: MEAN AND STANDARD DEVIATION OVER THREE SEEDS.

#	Data step	Res. input	Operator $\mathcal{D}, \mathcal{R}$	Memory	Mixer	K	SAM+SSIM	Deep sup.	Params (M)	PSNR	Δ	SAM	ERGAS	Cons. (dB)	Swap
<i>Pavia Centre</i> $\times 4$															
(a)	✓	✓	exact	✓	MDTA	3	✓	✓	9.99	50.62 $\pm$ 0.02	–	1.897	0.621	77.6/74.8	50.50
(b)	–	–	exact	✓	MDTA	3	✓	✓	9.99	50.46	-0.17	1.928	0.631	68.9/73.3	–
(c)	✓	–	exact	✓	MDTA	3	✓	✓	9.99	50.62	-0.01	1.900	0.621	75.7/73.8	–
(d)	✓	✓	zero border	✓	MDTA	3	✓	✓	9.99	50.54 $\pm$ 0.03	-0.08	1.915	0.627	74.5/73.0	49.64
(e)	✓	✓	learned	✓	MDTA	3	✓	✓	9.99	50.57	-0.06	1.908	0.626	73.6/72.6	–
(f)	✓	✓	exact	–	MDTA	3	✓	✓	9.98	50.67	+0.05	1.898	0.618	77.2/74.8	–
(g)	✓	✓	exact	–	MDTA	1	✓	✓	3.33	50.60	-0.02	1.909	0.624	71.3/70.8	–
(h)	✓	✓	exact	✓	MDTA	2	✓	✓	6.66	50.64	+0.02	1.901	0.619	75.0/73.4	–
(i)	✓	✓	exact	✓	MDTA	4	✓	✓	13.32	50.68	+0.06	1.894	0.617	80.2/76.3	–
(j)	✓	✓	exact	✓	conv	3	✓	✓	9.94	50.60	-0.02	1.909	0.624	78.5/73.9	–
(k)	✓	✓	exact	✓	window	3	✓	✓	9.99	50.63	+0.01	1.898	0.622	78.6/75.5	–
(l)	✓	✓	exact	✓	MDTA	3	–	✓	9.99	50.66	+0.03	1.897	0.619	77.5/75.7	–
(m)	✓	✓	exact	✓	MDTA	3	✓	–	9.99	50.50	-0.12	1.900	0.629	77.7/74.6	–
<i>Chikusei</i> $\times 4$															
(n)	✓	✓	exact	✓	MDTA	3	✓	✓	10.13	57.95	–	0.731	1.344	83.9/83.4	57.70
(o)	–	–	exact	✓	MDTA	3	✓	✓	–	–	–	–	–	–	–
(p)	✓	✓	zero border	✓	MDTA	3	✓	✓	–	–	–	–	–	–	–
(q)	✓	✓	learned	✓	MDTA	3	✓	✓	–	–	–	–	–	–	–
(r)	✓	✓	exact	–	MDTA	3	✓	✓	–	–	–	–	–	–	–
(s)	✓	✓	exact	–	MDTA	1	✓	✓	–	–	–	–	–	–	–
(t)	✓	✓	exact	✓	conv	3	✓	✓	–	–	–	–	–	–	–
<i>CAVE</i> $\times 4$															
(u)	✓	✓	exact	✓	MDTA	3	✓	✓	–	–	–	–	–	–	–
(v)	–	–	exact	✓	MDTA	3	✓	✓	–	–	–	–	–	–	–
(w)	✓	✓	zero border	✓	MDTA	3	✓	✓	–	–	–	–	–	–	–
(x)	✓	✓	learned	✓	MDTA	3	✓	✓	–	–	–	–	–	–	–

TABLE IV

CHIKUSEI: ALL METHODS RETRAINED UNDER THE IDENTICAL PROTOCOL AND BUDGET (4 H ON ONE T4, SAME DATA, PATCHES, BATCH, EMA AND VALIDATION-ONLY MODEL SELECTION; SINGLE PASS). BEST LEARNED METHOD IN BOLD.

Method	$\times 4$				$\times 8$				$\times 16$			
	PSNR $\uparrow$	SSIM $\uparrow$	SAM $\downarrow$	ERGAS $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	SAM $\downarrow$	ERGAS $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	SAM $\downarrow$	ERGAS $\downarrow$
Bicubic	33.63	0.7404	4.406	9.011	–	–	–	–	–	–	–	–
GSA	43.20	0.9778	2.730	3.279	–	–	–	–	–	–	–	–
SSR-NET [8]	–	–	–	–	–	–	–	–	–	–	–	–
PSRT [10]	53.99	0.9948	1.057	1.794	–	–	–	–	–	–	–	–
DCT [50]	53.18	0.9945	1.133	2.066	–	–	–	–	–	–	–	–
MIMFormer [26]	56.96	0.9964	0.801	1.488	–	–	–	–	–	–	–	–
DHIF-Net [20]	–	–	–	–	–	–	–	–	–	–	–	–
PUFormer	57.95	0.9969	0.731	1.344	–	–	–	–	–	–	–	–

is larger than on Chikusei although the training region is eight times smaller (97 against 800 non-overlapping patch equivalents).

The Q2n shortfall comes from one test image, which also has the largest SAM (Table XI, Fig. 4). On images 1, 3 and 4, Q2n is 0.9987–0.9991, at or above the best reported value, but SAM is still 1.55–1.68° (mean 1.60°), above the best reported 1.454°: the water image explains the Q2n gap, but only part of the SAM gap. Image 2 contains a stretch of the river Ticino: 63% of its pixels have a positive normalized difference water index (green against near infrared), against 11, 7 and 13% in the other images, and its SAM is 2.72°

and its Q2n 0.904. Water reflects little light, so its spectra are dominated by sensor noise that neither observation carries. The per-band noise that we estimate from the ground truth alone makes an angle of 3.8° with the noise-free spectrum in the darkest tenth of the pixels, against 0.23° in the brightest tenth. Over water PUFormer’s SAM is 2.95°, over all other pixels 1.56°, and the darkest 10% of the pixels account for 34% of its mean SAM. The same pixels make image 2 nearly flat within many 32×32 blocks, where Q2n divides the error by a very small local variance; image 2 alone carries 97% of the Q2n deficit $\sum_i (1 - Q2n_i)$ of the four images. Whether the test images of [15] contain water is not stated, so on this

TABLE V
PAVIA CENTRE: ALL METHODS RETRAINED UNDER THE IDENTICAL PROTOCOL AND BUDGET (4 H ON ONE T4, SAME DATA, PATCHES, BATCH, EMA AND VALIDATION-ONLY MODEL SELECTION; SINGLE PASS). BEST LEARNED METHOD IN BOLD.

Method	$\times 4$				$\times 8$				$\times 16$			
	PSNR $\uparrow$	SSIM $\uparrow$	SAM $\downarrow$	ERGAS $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	SAM $\downarrow$	ERGAS $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	SAM $\downarrow$	ERGAS $\downarrow$
Bicubic	25.52	0.6282	6.554	9.852	—	—	—	—	—	—	—	—
GSA	34.78	0.9438	5.229	3.494	—	—	—	—	—	—	—	—
SSR-NET [8]	—	—	—	—	—	—	—	—	—	—	—	—
PSRT [10]	47.87	0.9958	2.212	0.801	—	—	—	—	—	—	—	—
DCT [50]	46.94	0.9953	2.442	0.888	—	—	—	—	—	—	—	—
MIMFormer [26]	49.63	0.9967	1.982	0.674	—	—	—	—	—	—	—	—
DHIF-Net [20]	—	—	—	—	—	—	—	—	—	—	—	—
PUFormer	50.63	0.9970	1.897	0.621	—	—	—	—	—	—	—	—

TABLE VI
CAVE: ALL METHODS RETRAINED UNDER THE IDENTICAL PROTOCOL AND BUDGET (4 H ON ONE T4, SAME DATA, PATCHES, BATCH, EMA AND VALIDATION-ONLY MODEL SELECTION; SINGLE PASS). BEST LEARNED METHOD IN BOLD.

Method	$\times 4$				$\times 8$			
	PSNR $\uparrow$	SSIM $\uparrow$	SAM $\downarrow$	ERGAS $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	SAM $\downarrow$	ERGAS $\downarrow$
Bicubic	—	—	—	—	—	—	—	—
GSA	—	—	—	—	—	—	—	—
SSR-NET [8]	—	—	—	—	—	—	—	—
PSRT [10]	—	—	—	—	—	—	—	—
DCT [50]	—	—	—	—	—	—	—	—
MIMFormer [26]	—	—	—	—	—	—	—	—
DHIF-Net [20]	—	—	—	—	—	—	—	—
PUFormer	—	—	—	—	—	—	—	—

dataset SAM and Q2n are the least comparable of the eight metrics and RMSE the most.

As on Chikusei, our GSA scores below the value reported for it (34.78 against 35.35 dB), so the rebuilt data do not favor our method. The reconstruction reproduces its inputs to within 79.9 dB (LR-HSI) and 77.6 dB (HR-MSI). Validation PSNR stops improving after 43,000 iterations (Fig. 5), and those weights are tested; the test PSNR, logged but never used, peaks at 50.66 dB at 25,000 iterations, against 50.63 dB for the selected weights.

E. Per-image results and visual quality

Table XI gives the results of every test image. Chikusei images 5–8 are darker (their maximum is 0.36–0.47 of the dataset maximum), which lowers their strict SSIM but not their RMSE. Fig. 6 shows Chikusei image 7, the hardest test image, and Fig. 7 Pavia Centre image 3, the hardest there. The errors of bicubic interpolation and GSA follow every field boundary and road (mean absolute errors of 164.7 and 64.8 DN), whereas both PUFormer models keep the error

small and nearly uniform over the image (11.5 and 11.0 DN), apart from a few isolated pixels.

Fig. 8 shows the error of every band. PUFormer is more accurate than GSA in all 128 bands, and the final model improves on the first run in every band, most at the red edge (700–760 nm), where the mean band RMSE falls from 7.9 to 4.2 DN (from 9.8 to 1.1 DN at 724 nm). The error is largest in the first eight bands (363–404 nm), which lie below the first WorldView-2 band and therefore receive no high-resolution detail from the multispectral image, and in the last bands near 1000 nm, where the sensor’s signal-to-noise ratio drops.

F. Training stages and the exact border (Chikusei)

Table XII traces the path from the first run to the final model. Evaluating the first-run weights with the exact reflect-border operator, without retraining, drops the validation PSNR from 58.60 to 49.52 dB. The prior had learned to undo the error that the zero-padded operator introduces at the border, so with the exact operator it over-corrects there; the effect is strongest on the small 64×64 validation patches, where

TABLE VII

HARVARD: ALL METHODS RETRAINED UNDER THE IDENTICAL PROTOCOL AND BUDGET (4 H ON ONE T4, SAME DATA, PATCHES, BATCH, EMA AND VALIDATION-ONLY MODEL SELECTION; SINGLE PASS). BEST LEARNED METHOD IN BOLD.

Method	$\times 4$				$\times 8$			
	PSNR $\uparrow$	SSIM $\uparrow$	SAM $\downarrow$	ERGAS $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	SAM $\downarrow$	ERGAS $\downarrow$
Bicubic	—	—	—	—	—	—	—	—
GSA	—	—	—	—	—	—	—	—
SSR-NET [8]	—	—	—	—	—	—	—	—
PSRT [10]	—	—	—	—	—	—	—	—
DCT [50]	—	—	—	—	—	—	—	—
MIMFormer [26]	—	—	—	—	—	—	—	—
DHIF-Net [20]	—	—	—	—	—	—	—	—
PUFormer	—	—	—	—	—	—	—	—




Fig. 3. Test PSNR of the retrained methods against the scale factor (equal budget, single pass). GSA is shown for reference.

border pixels are a large fraction of the image. Fine-tuning removes this within 2000 iterations (58.60 dB validation) and ends at 58.72 dB. The consistency of the reconstruction with its inputs rises from 76.7 to 85.5 dB for the LR-HSI and from 75.0 to 89.0 dB for the HR-MSI. Fig. 9 shows the validation curves; validation PSNR stops improving after about 44,000 fine-tuning iterations, and the test PSNR stays between 58.13 and 58.15 dB from iteration 26,000 onward.

G. Robustness and test-time operator swap

Published methods are usually trained and tested with the same degradation. We test the final model on inputs made with a different blur or a shifted spectral response, and with Gaussian noise added to both observations (Table XIII, Fig. 10). Because the operators are explicit, PUFormer can also be given the true test-time operator without retraining (“swap”). With the training operator, a blur mismatch costs 3.9–7.5 dB. With the true blur swapped in, the loss is at most 0.26 dB for σ between 1.5 and 3.0; with the zero-border operators of the first run, the same swap recovered only to 55.2–57.1 dB. A shifted spectral response is not compensated by the swap, which lowers the PSNR by a further 0.5 dB: the priors, trained with one response, rely on it beyond the

data step. Noise is the main weakness of a model trained on noise-free data; PUFormer still exceeds GSA by 10.0 dB at 40 dB SNR and by 2.8 dB at 30 dB SNR. On Pavia Centre the blur behaves the same way: the training operator loses 3.3–6.8 dB and the swapped-in true operator at most 0.11 dB. Unlike on Chikusei, swapping in a shifted IKONOS response raises the PSNR slightly (by 0.2–0.3 dB), perhaps because an 8 nm shift is small relative to the 66–96 nm IKONOS bands, and the margins over GSA under noise are larger (12.8, 9.8 and 6.0 dB at 40, 35 and 30 dB SNR).

H. Operators estimated from the test pair

The swap above assumes that the test blur is known. In practice it has to be estimated. Because $\mathcal{D}$ acts on space and $\mathcal{R}$ on wavelength, they commute, and for the true blur $\mathcal{D}\mathbf{Y}_M = \mathcal{R}\mathbf{Y}_H$ holds exactly. We therefore estimate the blur width of each test image as $\hat{\sigma} = \arg \min_{\sigma} \|\mathcal{D}_{\sigma}\mathbf{Y}_M - \mathcal{R}\mathbf{Y}_H\|^2$ by a grid search with step 0.05, using only the two observations and the training spectral response, and give $\mathcal{D}_{\hat{\sigma}}$ to the network. Table ?? compares this with the training operator and the true one, without and with 35 dB noise on both inputs. On Pavia Centre the estimate is within ?? of the true width at $\sigma = 3$, and with it the exact model’s worst PSNR over all test blurs

TABLE VIII

CHIKUSEI $\times 4$ COMPARISON. ALL ROWS EXCEPT PUFORMER ARE REPORTED IN [15]. BEST VALUE PER COLUMN IN BOLD. SSIM OF PUFORMER USES THE PSRT DEFINITION (DATA RANGE 1); WITH DATA RANGE EQUAL TO EACH IMAGE'S MAXIMUM IT IS 0.9969 (SINGLE PASS) AND 0.9970 (SELF-ENSEMBLE).

Method	PSNR $\uparrow$	SSIM $\uparrow$	SAM $\downarrow$	ERGAS $\downarrow$	Q2n $\uparrow$	CC $\uparrow$	SCC $\uparrow$	RMSE (DN) $\downarrow$
GSA [4]	47.5769	0.9933	1.7216	2.1667	0.9658	0.9894	0.9773	70.54
CNMF [5]	44.5291	0.9891	1.7259	2.8676	0.9662	0.9828	0.9645	101.64
FUSE [7]	41.1550	0.9787	2.5814	3.4710	0.9418	0.9736	0.9486	148.10
U2Net [34]	50.5455	0.9956	1.2389	1.8740	0.9928	0.9920	0.9798	49.97
PSRT [10]	50.6722	0.9952	1.2733	2.0107	0.9862	0.9906	0.9778	49.00
DCT [50]	52.1998	0.9977	0.9239	1.7448	0.9832	0.9935	0.9873	41.45
DDIF [14]	53.4421	0.9972	0.9387	1.8087	0.9963	0.9929	0.9840	35.58
MIMFormer [26]	53.3578	0.9973	0.9094	1.7611	0.9948	0.9931	0.9850	35.90
MIMO-SST [11]	53.6073	0.9972	0.9315	1.8119	0.9944	0.9926	0.9834	34.93
OTPNNet [35]	52.8243	0.9969	1.0201	1.9332	0.9940	0.9916	0.9841	38.49
SMGU [39]	54.5244	0.9977	0.8719	1.6405	0.9952	0.9939	0.9875	31.29
AMSF [27]	49.3827	0.9944	1.4852	2.3943	0.9878	0.9879	0.9790	56.78
SSDT [13]	52.4788	0.9975	0.8768	1.7992	0.9939	0.9928	0.9847	40.00
ASMNet [28]	53.4528	0.9972	0.9472	1.9153	0.9874	0.9917	0.9848	35.89
DPFormer [29]	53.1632	0.9970	0.9483	1.8795	0.9949	0.9923	0.9824	37.18
CLSNet [49]	55.2007	0.9979	0.7836	1.6318	0.9943	0.9941	0.9869	29.06
Two-stage diffusion [15]	56.1871	0.9982	0.7203	1.5158	0.9969	0.9948	0.9891	25.73
PUFormer (single pass)	58.1467	0.9989	0.7171	1.3367	0.9930	0.9967	0.9992	19.06
PUFormer (self-ensemble)	58.2016	0.9989	0.7126	1.3304	0.9931	0.9967	0.9992	18.94

TABLE IX

METHODS RUN UNDER OUR PIPELINE, WITH ALL PSNR AND SSIM DEFINITIONS. P_g : PEAK = DATASET MAXIMUM (HEADLINE); P_i : PEAK = IMAGE MAXIMUM; P_b : BAND-WISE MEAN, PEAK 1. S_1 : SSIM, DATA RANGE 1; S_i : DATA RANGE = IMAGE MAXIMUM.

Method	P_g	P_i	P_b	S_1	S_i	SAM	RMSE
<i>Chikusei</i>							
Bicubic	33.63	28.20	35.25	0.8500	0.7404	4.406	321.8
GSA [4]	43.20	37.76	44.82	0.9859	0.9778	2.730	107.7
SSR-NET [8]	49.98	44.55	53.83	—	0.9922	1.761	48.8
PUFormer, first run	57.99	52.56	65.29	0.9989	0.9969	0.728	19.3
PUFormer v3	58.15	52.71	66.26	0.9989	0.9969	0.717	19.1
PUFormer v3, ensemble	58.20	52.77	66.33	0.9989	0.9970	0.713	18.9
<i>Pavia Centre</i>							
Bicubic	25.52	25.11	25.87	0.6427	0.6282	6.555	436.3
GSA [4]	34.78	34.37	35.19	0.9450	0.9438	5.229	157.9
PUFormer	50.63	50.22	51.11	0.9974	0.9970	1.895	23.7
PUFormer, ensemble	50.69	50.28	51.17	0.9974	0.9970	1.883	23.5

is ?? dB, against ?? dB with the training operator, ?? dB for the zero-border model with the same estimate, and ?? dB for the learned operator. The spectral response is not estimated: a swapped response did not help even when it was the true one (Table XIII), because the priors depend on the training response beyond the data step.

I. Where the Q2n gap on Chikusei comes from

Q2n normalizes every band of every 32×32 block by the ground truth's own mean and standard deviation in that block, so an error that is small in absolute terms counts heavily in a band that is nearly flat within a block. Table XIV and Fig. 11 decompose the Q2n of the final model (0.9931). If the eight ultraviolet bands (363–404 nm) were reconstructed exactly, Q2n would rise to 0.9988; if the last 24 bands were exact as well, to 0.9998. In the worst 10% of blocks, the error energy of these eight bands equals 0.41 of the ground truth's local variance, against at most 0.040 for any other group of eight bands. The fine detail of these bands cannot be predicted from

the inputs: a least-squares injection of the multispectral detail, fitted on the test ground truth itself, reduces the band-1 error only from 61.6 to 57.5 DN, while PUFormer reaches 51.8 DN, and PUFormer is more accurate than this oracle in every band we checked. Q2n is therefore nearly insensitive to overall accuracy on these test images: blending our reconstruction toward bicubic interpolation lowers the PSNR from 58.20 to 52.15 dB but Q2n only from 0.9931 to 0.9904 (Fig. 11, right). The values reported in [15] do not follow PSNR in this way. Between 52.2 and 53.6 dB they range from 0.9832 (DCT) to 0.9963 (DDIF); of the 16 methods within the PSNR range of our curve, 11 lie above it and 5 below. U2Net, for example, is listed at 0.9928 at 50.55 dB, a value our implementation reaches only near 57.6 dB. We conclude that the Q2n values of [15] were most likely computed with a different implementation, normalization or test region, and that they are not directly comparable with ours.

J. Complexity

PUFormer has 10.13 M parameters and needs 18.6 GFLOPs for a 64×64 patch and 297 GFLOPs for a 256×256 image. On one Tesla T4 it reconstructs a 256×256 image with 128 bands in 0.20 s with mixed precision (92 bands: 0.18 s), and the self-ensemble doubles this. Because the attention is computed across channels, memory and time grow linearly with the number of pixels. Table ?? compares the parameters, operations, time and memory of all retrained networks on the same input.

VII. DISCUSSION AND LIMITATIONS

What the exact operator buys. At the blur used for training, the exact operator improves PSNR only slightly over a zero-border or a learned operator, by less than a tenth of a decibel on Pavia Centre. Its clearer effects are a reconstruction that reproduces its inputs about 3 dB more closely, and a network whose operator can be exchanged at test time, which

TABLE X

PAVIA CENTRE $\times 4$ COMPARISON. ALL ROWS EXCEPT PUFORMER ARE REPORTED IN [15] ON THE BENCHMARK DATA RELEASED WITH PSRT [10]; PUFORMER IS EVALUATED ON OUR SIMULATION OF THE SAME PROTOCOL FROM THE ORIGINAL SCENE (SECTION V). BEST VALUE PER COLUMN IN BOLD. SSIM OF PUFORMER USES THE PSRT DEFINITION (DATA RANGE 1); WITH DATA RANGE EQUAL TO EACH IMAGE'S MAXIMUM IT IS 0.9970 (SINGLE PASS) AND 0.9970 (SELF-ENSEMBLE).

Method	PSNR $\uparrow$	SSIM $\uparrow$	SAM $\downarrow$	ERGAS $\downarrow$	Q2n $\uparrow$	CC $\uparrow$	SCC $\uparrow$	RMSE (DN) $\downarrow$
GSA [4]	35.3495	0.9736	5.0812	2.7260	0.9818	0.9832	0.9768	137.78
CNMF [5]	32.1693	0.9469	3.8494	3.8446	0.9489	0.9683	0.9294	198.26
FUSE [7]	34.7418	0.9612	5.5628	3.2042	0.9103	0.9779	0.9469	146.90
U2Net [34]	44.3439	0.9963	1.6542	1.0924	0.9964	0.9973	0.9963	49.10
PSRT [10]	45.8875	0.9970	1.5230	0.9750	0.9974	0.9978	0.9966	41.85
DCT [50]	46.6745	0.9971	1.5169	0.8753	0.9973	0.9983	0.9972	37.66
DDIF [14]	46.2530	0.9970	1.5253	0.9125	0.9960	0.9982	0.9970	39.56
MIMFormer [26]	44.8497	0.9959	1.8225	1.0683	0.9957	0.9975	0.9954	46.19
MIMO-SST [11]	46.6334	0.9969	1.5587	0.8607	0.9975	0.9984	0.9972	37.68
OTPNNet [35]	45.8081	0.9967	1.6573	0.9465	0.9987	0.9980	0.9968	41.51
SMGU [39]	47.1239	0.9972	1.4690	0.8547	0.9981	0.9984	0.9974	37.09
AMSF [27]	43.4465	0.9950	2.0533	1.2728	0.9929	0.9965	0.9940	54.25
SSDT [13]	46.3173	0.9965	1.6439	0.8918	0.9968	0.9982	0.9969	39.01
ASMNNet [28]	46.2072	0.9966	1.6517	0.9019	0.9940	0.9982	0.9969	39.50
DPFormer [29]	42.4280	0.9948	1.9367	1.3986	0.9945	0.9957	0.9917	60.96
CLSNet [49]	46.5277	0.9969	1.5708	0.8412	0.9983	0.9983	0.9975	36.21
Two-stage diffusion [15]	47.1597	0.9972	1.4542	0.8262	0.9980	0.9985	0.9976	35.71
PUFormer (single pass)	50.6299	0.9974	1.8955	0.6206	0.9751	0.9994	0.9995	23.68
PUFormer (self-ensemble)	50.6915	0.9974	1.8833	0.6170	0.9753	0.9994	0.9995	23.52

TABLE XI

PER-IMAGE RESULTS OF THE FINAL MODELS (SELF-ENSEMBLE)

Image	PSNR	S_i	S_1	SAM	ERGAS	Q2n
<i>Chikusei, PUFormer v3</i>						
1	59.15	0.9994	0.9994	0.644	1.550	0.9957
2	60.45	0.9985	0.9994	0.643	0.972	0.9967
3	60.97	0.9987	0.9995	0.597	1.093	0.9962
4	57.10	0.9981	0.9988	0.862	1.334	0.9917
5	57.29	0.9951	0.9986	0.726	1.304	0.9929
6	57.51	0.9957	0.9987	0.669	1.357	0.9913
7	56.24	0.9953	0.9984	0.830	1.449	0.9913
8	56.90	0.9947	0.9986	0.729	1.584	0.9889
Mean	58.20	0.9970	0.9989	0.713	1.330	0.9931
<i>Pavia Centre, PUFormer</i>						
1	50.48	0.9977	0.9977	1.546	0.555	0.9987
2	52.30	0.9949	0.9964	2.722	0.813	0.9044
3	49.53	0.9979	0.9979	1.677	0.557	0.9991
4	50.45	0.9975	0.9976	1.588	0.543	0.9991
Mean	50.69	0.9970	0.9974	1.883	0.617	0.9753

TABLE XII

FROM THE FIRST RUN TO THE FINAL MODEL (TEST SET UNLESS MARKED). CONS.: PSNR($\mathcal{D}\hat{\mathbf{X}}, \mathbf{Y}_H$) / PSNR($\mathcal{R}\hat{\mathbf{X}}, \mathbf{Y}_M$) IN dB

Model	Val.	PSNR	SAM	ERGAS	Cons.
GSA	–	43.20	2.730	3.279	45.4 / 45.0
First run (zero border)	58.60	57.99	0.728	1.356	76.7 / 75.0
+ self-ensemble	–	58.02	0.726	1.353	76.8 / 75.0
First-run weights, exact border	49.52	54.24	0.773	1.474	–
v3 (exact border, full loss)	58.72	58.15	0.717	1.337	85.5 / 89.0
+ self-ensemble	–	58.20	0.713	1.330	86.0 / 89.5

turns a blur mismatch from a loss of several decibels into one of about a tenth of a decibel, with the true or an estimated blur. The ablation also shows that the transformer prior is not where the accuracy comes from; a simpler prior, or a single stage, performs almost as well. We therefore see the exact operator

as a property worth enforcing in any unfolded fusion network rather than as a feature of PUFormer's architecture.

Limitations. All experiments use simulated data, in which the operators are known up to the parameters we vary; on real sensor pairs the blur, the spectral response and the co-registration are only approximately known. The blur estimate of Section VI-H handles the first, but a mismatched spectral response is not compensated by swapping it, and training with randomized responses would be needed. The models are trained without noise, which costs accuracy at low signal-to-noise ratios (Table XIII). The equal-budget benchmark gives every method the same 4 h on one GPU; this is fair in compute but favors fast networks, and although the longer-budget check and the comparison with reported values suggest that the ranking holds, longer training could narrow some gaps. The comparison with reported results in Tables VIII and X rests on test images that follow the published protocol but are not necessarily the same pixels, and on Pavia Centre data rebuilt from the original scene. There RMSE in digital numbers is the most robust basis, whereas SSIM with a per-image data range, Q2n and, on Pavia Centre, SAM are below the best reported values. Finally, the test sets are small (4 to 20 images), which limits per-image significance tests.

VIII. CONCLUSION

We studied how exactly the observation operators inside a deep unfolding network should match the operators that generated the data. With exact operators, border included, the ground truth is a fixed point of every data-consistency step; with a zero border or a learned operator it is not. In an ablation with repeated seeds, the exact operator gives a small but consistent accuracy gain, a reconstruction that agrees better with its inputs, and a network whose blur can be replaced at test time with almost no loss, also when the blur is estimated from the test images. Most components of the transformer




Fig. 4. The four Pavia Centre test images in false color (bands 57, 34 and 3; top), with water pixels (NDWI > 0) outlined, and the spectral angle of PUFormer's reconstruction at every pixel (bottom; color scale clipped at 5°).




Fig. 5. Pavia Centre: PSNR and SAM during training from scratch, on the validation patches (used to select the weights) and on the test images (logged only). The dashed line is the best value reported in [15] on the PSRT benchmark data.

prior, in contrast, can be removed with little effect. In an equal-budget benchmark against five published networks retrained from their authors' code on four datasets and scale factors from 4 to 16, the resulting network, PUFormer, has the highest PSNR in 2 of 2 settings, and under the protocol of [15] it reaches 58.20 dB on Chikusei $\times 4$ and 50.69 dB on Pavia Centre $\times 4$. Applying exact operators to real sensor pairs with estimated blur and spectral response is the main direction for future work. All code, training logs and result files are publicly available.

APPENDIX A COMPLETE BENCHMARK METRICS

Table XV lists all eight metrics of every method in the equal-budget benchmark of Section VI-B; the per-job result files, including per-image values and training curves, are in the public repository.

ACKNOWLEDGMENT

The author thanks N. Yokoya and A. Iwasaki for making the Chikusei data available, P. Gamba of the University of Pavia




Fig. 6. Chikusei test image 7 in false color (bands 60, 40 and 20; top) and the mean absolute error over all bands in DN (bottom; color scale clipped at 60 DN).




Fig. 7. Pavia Centre test image 3 in false color (bands 57, 34 and 3; top) and the mean absolute error over all bands in DN (bottom; color scale clipped at 60 DN).

for the Pavia Centre data, and Kaggle for the GPU time used to train and evaluate the models.

REFERENCES

- [1] N. Yokoya, C. Grohnfeldt, and J. Chanussot, "Hyperspectral and multispectral data fusion: A comparative review of the recent literature," *IEEE Geoscience and Remote Sensing Magazine*, vol. 5, no. 2, pp. 29–56, 2017.
- [2] L. Loncan, L. B. de Almeida, J. M. Bioucas-Dias, X. Briottet, J. Chanussot, N. Dobigeon, S. Fabre, W. Liao, G. A. Licciardi, M. Simoes, J.-Y. Tournet, M. A. Vezanones, G. Vivone, Q. Wei, and N. Yokoya, "Hyperspectral pansharpening: A review," *IEEE Geoscience and Remote Sensing Magazine*, vol. 3, no. 3, pp. 27–46, 2015.
- [3] G. Vivone, L.-J. Deng, S. Deng, D. Hong, M. Jiang, C. Li, W. Li, H. Shen, X. Wu, J.-L. Xiao, J. Yao, M. Zhang, J. Chanussot, S. García, and A. Plaza, "Deep learning in remote sensing image fusion: Methods, protocols, data, and future perspectives," *IEEE Geoscience and Remote Sensing Magazine*, vol. 13, no. 1, pp. 269–310, Mar. 2025.
- [4] B. Aiazzi, S. Baronti, and M. Selva, "Improving component substitution pansharpening through multivariate regression of MS + Pan data," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 45, no. 10, pp. 3230–3239, Oct. 2007.
- [5] N. Yokoya, T. Yairi, and A. Iwasaki, "Coupled nonnegative matrix factorization unmixing for hyperspectral and multispectral data fusion," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 50, no. 2, pp. 528–537, Feb. 2012.
- [6] M. Simoes, J. Bioucas-Dias, L. B. Almeida, and J. Chanussot, "A convex formulation for hyperspectral image superresolution via subspace-based regularization," *IEEE Transactions on Geoscience and Remote Sensing*,




Fig. 8. Band-wise PSNR (left, peak 1) and RMSE in DN (right, log scale) over the test images of Chikusei (top) and Pavia Centre (bottom). Shaded ranges have no multispectral coverage: below the first WorldView-2 band for Chikusei, and between and beyond the IKONOS bands (50% edges) for Pavia Centre.

TABLE XIII

ROBUSTNESS OF THE FINAL MODELS (TEST PSNR in dB / SAM in DEGREES). FIXED: TRAINING OPERATOR; SWAP: TRUE OPERATOR SUPPLIED AT TEST TIME. THE BETTER PUFORMER ENTRY OF EACH CONDITION IS IN BOLD.

Condition	Chikusei (PUFormer v3)			Pavia Centre (PUFormer)		
	GSA	PUFormer fixed	PUFormer swap	GSA	PUFormer fixed	PUFormer swap
Nominal	43.20 / 2.73	58.15 / 0.72	–	34.78 / 5.23	50.63 / 1.90	–
Blur $\sigma = 1.5$	44.05 / 2.57	50.68 / 1.15	57.97 / 0.73	35.59 / 4.99	43.80 / 2.75	50.56 / 1.90
Blur $\sigma = 2.5$	42.77 / 2.82	54.26 / 0.88	58.05 / 0.72	34.36 / 5.36	47.30 / 2.14	50.59 / 1.90
Blur $\sigma = 3.0$	42.54 / 2.87	51.69 / 1.03	57.89 / 0.74	34.13 / 5.43	44.87 / 2.37	50.52 / 1.91
SRF shift -8 nm	42.86 / 2.91	53.01 / 1.20	52.54 / 1.31	34.58 / 5.36	49.46 / 1.99	49.75 / 2.03
SRF shift $+8$ nm	43.29 / 2.68	53.03 / 1.19	52.47 / 1.28	34.79 / 5.20	49.62 / 1.99	49.82 / 2.02
Noise, 40 dB SNR	43.04 / 2.85	53.03 / 1.39	–	34.71 / 5.37	47.55 / 2.89	–
Noise, 35 dB SNR	42.71 / 3.09	49.11 / 2.20	–	34.57 / 5.64	44.36 / 4.19	–
Noise, 30 dB SNR	41.84 / 3.67	44.62 / 3.69	–	34.16 / 6.37	40.18 / 6.60	–

TABLE XIV

Q2N ANALYSIS OF PUFORMER V3 (SELF-ENSEMBLE, TEST SET)

Setting	Q2n
PUFormer v3	0.9931
bands 1–8 (363–404 nm) set to ground truth	0.9988
bands 1–8 and 105–128 set to ground truth	0.9998
Deficit share, worst 10% / 50% of blocks	25% / 75%

vol. 53, no. 6, pp. 3373–3388, 2015.

- [7] Q. Wei, N. Dobigeon, and J.-Y. Tourneret, “Fast fusion of multi-band images based on solving a Sylvester equation,” *IEEE Transactions on*

Image Processing, vol. 24, no. 11, pp. 4109–4121, Nov. 2015.

- [8] X. Zhang, W. Huang, Q. Wang, and X. Li, “SSR-NET: Spatial-spectral reconstruction network for hyperspectral and multispectral image fusion,” *IEEE Transactions on Geoscience and Remote Sensing*, vol. 59, no. 7, pp. 5953–5965, 2021.
- [9] R. Ran, L.-J. Deng, T.-X. Jiang, J.-F. Hu, J. Chanussot, and G. Vivone, “GuidedNet: A general CNN fusion framework via high-resolution guidance for hyperspectral image super-resolution,” *IEEE Transactions on Cybernetics*, vol. 53, no. 7, pp. 4148–4161, 2023.
- [10] S.-Q. Deng, L.-J. Deng, X. Wu, R. Ran, D. Hong, and G. Vivone, “PSRT: Pyramid shuffle-and-resuffle transformer for multispectral and hyperspectral image fusion,” *IEEE Transactions on Geoscience and Remote Sensing*, vol. 61, pp. 1–15, 2023.
- [11] J. Fang, J. Yang, A. Khader, and L. Xiao, “MIMO-SST: Multi-input multi-output spatial-spectral transformer for hyperspectral and multi-




Fig. 9. Validation PSNR and SAM during training, plotted against the number of 64×64 training patches seen. The first run starts from random weights (its first two evaluations lie below the plotted range); v3 continues from its final weights with the exact border and the full loss.




Fig. 10. Test PSNR under degradation mismatch and noise on Chikusei (left) and Pavia Centre (right). The dashed line is the nominal PUFormer result.

- spectral image fusion,” *IEEE Transactions on Geoscience and Remote Sensing*, vol. 62, pp. 1–20, 2024.
- [12] W. He, X. Fu, N. Li, Q. Ren, and S. Jia, “LGCT: Local-global collaborative transformer for fusion of hyperspectral and multispectral images,” *IEEE Transactions on Geoscience and Remote Sensing*, vol. 62, pp. 1–14, 2024.
- [13] X. Fu, P. Liu, P. Han, and S. Jia, “SSDT: Multiscale spatial-spectral dilated transformer for hyperspectral and multispectral image fusion,” *IEEE Transactions on Geoscience and Remote Sensing*, vol. 63, pp. 1–15, 2025.
- [14] Z. Cao, S. Cao, L.-J. Deng, X. Wu, J. Hou, and G. Vivone, “Diffusion model with disentangled modulations for sharpening multispectral and hyperspectral images,” *Information Fusion*, vol. 104, p. 102158, Apr. 2024.
- [15] Y. Chen, R. Liu, W. K. Wong, and J. Wen, “A two-stage conditional diffusion model with differential attention for hyperspectral and multispectral image fusion,” *IEEE Transactions on Image Processing*, vol. 35, pp. 8980–8995, 2026.
- [16] Y. Xu, J. Zhu, D. Hong, Z. Wei, and Z. Wu, “Coupled diffusion posterior sampling for unsupervised hyperspectral and multispectral images fusion,” *IEEE Transactions on Image Processing*, vol. 35, pp. 69–84, 2026.
- [17] J. Zhu, H. Wang, Y. Xu, Z. Wu, and Z. Wei, “Self-learning hyperspectral and multispectral image fusion via adaptive residual guided subspace diffusion model,” in *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025, pp. 17 862–17 871.
- [18] Q. Xie, M. Zhou, Q. Zhao, Z. Xu, and D. Meng, “MHF-Net: An interpretable deep network for multispectral and hyperspectral image fusion,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 44, no. 3, pp. 1457–1473, Mar. 2022.
- [19] W. Dong, C. Zhou, F. Wu, J. Wu, G. Shi, and X. Li, “Model-guided deep hyperspectral image super-resolution,” *IEEE Transactions on Image Processing*, vol. 30, pp. 5754–5768, 2021.
- [20] T. Huang, W. Dong, J. Wu, L. Li, X. Li, and G. Shi, “Deep hyperspectral image fusion network with iterative spatio-spectral regularization,” *IEEE Transactions on Computational Imaging*, vol. 8, pp. 201–214, 2022.
- [21] J. Fang, J. Yang, A. Khader, and L. Xiao, “Deep unfolding network enhanced by transformer priors for unregistered hyperspectral and multispectral image fusion,” *IEEE Transactions on Geoscience and Remote Sensing*, vol. 62, pp. 1–16, 2024.
- [22] X. Tian, K. Li, W. Zhang, Z. Wang, and J. Ma, “Interpretable model-driven deep network for hyperspectral, multispectral, and panchromatic image fusion,” *IEEE Transactions on Neural Networks and Learning Systems*, vol. 35, no. 10, pp. 14 382–14 395, Oct. 2024.




Fig. 11. Left: error energy relative to the local ground-truth variance per group of eight bands in the worst 10% of 32×32 blocks. Right: Q2n against PSNR when PUFormer's output is blended toward bicubic interpolation (our implementation), and the (PSNR, Q2n) pairs reported in [15].

- [23] J. Sun, B. Chen, R. Lu, Z. Cheng, C. Qu, and X. Yuan, "Advancing hyperspectral and multispectral image fusion: An information-aware transformer-based unfolding network," *IEEE Transactions on Neural Networks and Learning Systems*, vol. 36, no. 4, pp. 7407–7421, Apr. 2025.
- [24] J. Bioucas-Dias and J. Nascimento, "Hyperspectral subspace identification," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 46, no. 8, pp. 2435–2445, Aug. 2008.
- [25] Y. Sun, H. Xu, Y. Ma, M. Wu, X. Mei, J. Huang, and J. Ma, "Dual spatial-spectral pyramid network with transformer for hyperspectral image fusion," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 61, pp. 1–16, 2023.
- [26] R. Li, L. Zhang, Z. Wang, and X. Li, "MIMFormer: Multiscale inception mixer transformer for hyperspectral and multispectral image fusion," *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, vol. 17, pp. 15 122–15 135, 2024.
- [27] S. Liu, T. Shao, S. Liu, B. Li, and Y.-D. Zhang, "An asymptotic multiscale symmetric fusion network for hyperspectral and multispectral image fusion," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 63, pp. 1–16, 2025.
- [28] J. Liu, S. Li, R. Dian, Z. Song, and L. Tan, "Asymptotic spectral mapping for hyperspectral image fusion," *IEEE Transactions on Circuits and Systems for Video Technology*, vol. 35, no. 4, pp. 3475–3485, Apr. 2025.
- [29] S. Liu, R. Cai, H. Yue, B. Li, Y. Zhang, and S. Liu, "Dual-phase transformer with spectral-spatial synergy for hyperspectral image fusion," *Pattern Recognition*, vol. 175, p. 113116, 2026.
- [30] S. Peng, X. Zhu, H. Deng, L.-J. Deng, and Z. Lei, "FusionMamba: Efficient remote sensing image fusion with state space model," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 62, pp. 1–16, 2024.
- [31] J. Huang, T. Luo, Z. He, H. Xu, Y. Song, and L. Fang, "HDGMamba: High-frequency dynamic guided mamba for robust multispectral-hyperspectral image fusion," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 64, pp. 5 514 016–5 514 016, 2026.
- [32] N. Xiao, X. Fu, Q. Ren, W. He, S. Wei, and S. Jia, "Region-aware MoE network for hyperspectral and multispectral image fusion," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 64, pp. 1–15, 2026.
- [33] X. Huo and H. Gan, "Leveraging modal interaction and window dilation in attention network for hyperspectral and multispectral remote sensing image fusion," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 64, pp. 1–15, 2026.
- [34] S. Peng, C. Guo, X. Wu, and L.-J. Deng, "U2Net: A general framework with spatial-spectral-integrated double U-Net for image fusion," in *Proceedings of the 31st ACM International Conference on Multimedia*, ser. MM '23, Oct. 2023, pp. 3219–3227.
- [35] W. Yu, Z. Li, Q. Liu, and X. Sun, "OTPNNet: ODE-inspired tuning-free proximal network for remote sensing image fusion," *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 39, no. 9, pp. 9671–9679, Apr. 2025.
- [36] K. Gregor and Y. LeCun, "Learning fast approximations of sparse coding," in *Proc. 27th Int. Conf. Mach. Learn. (ICML)*, 2010, pp. 399–406.
- [37] V. Monga, Y. Li, and Y. C. Eldar, "Algorithm unrolling: Interpretable, efficient deep learning for signal and image processing," *IEEE Signal Processing Magazine*, vol. 38, no. 2, pp. 18–44, Mar. 2021.
- [38] J. Zhang and B. Ghanem, "ISTA-Net: Interpretable optimization-inspired deep network for image compressive sensing," in *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2018, pp. 1828–1837.
- [39] J. Yan, K. Zhang, Q. Sun, C. Ge, W. Wan, J. Sun, and H. Zhang, "Spatial-spectral unfolding network with mutual guidance for multispectral and hyperspectral image fusion," *Pattern Recognition*, vol. 161, p. 111277, May 2025.
- [40] B. Wang, X. Xiong, Y. Lian, X. Cao, H. Zhou, and K. Yu, "Perceptive spectral transformer unfolding network with multiscale mixed training for arbitrary-scale hyperspectral and multispectral image fusion," *Information Fusion*, vol. 122, p. 103166, 2025.
- [41] X. Cao, X. Wang, X. Dun, Y. Lian, X. Cheng, and X. Hao, "Cross-domain-aware deep unfolding transformer for hyperspectral image super-resolution," *Pattern Recognition*, vol. 172, p. 112374, 2026.
- [42] J. Song, B. Chen, and J. Zhang, "Memory-augmented deep unfolding network for compressive sensing," in *Proc. 29th ACM Int. Conf. Multimedia*, 2021, pp. 4249–4258.
- [43] Y. Cai, J. Lin, H. Wang, X. Yuan, H. Ding, Y. Zhang, R. Timofte, and L. Van Gool, "Degradation-aware unfolding half-shuffle transformer for spectral compressive imaging," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 35, 2022, pp. 37 749–37 761.
- [44] S. W. Zamir, A. Arora, S. Khan, M. Hayat, F. S. Khan, and M.-H. Yang, "Restormer: Efficient transformer for high-resolution image restoration," in *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022, pp. 5718–5729.
- [45] K. Zhang, L. Van Gool, and R. Timofte, "Deep unfolding network for image super-resolution," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2020, pp. 3214–3223.
- [46] M. K. Ng, R. H. Chan, and W.-C. Tang, "A fast algorithm for deblurring models with Neumann boundary conditions," *SIAM Journal on Scientific Computing*, vol. 21, no. 3, pp. 851–866, 1999.
- [47] S. J. Reeves, "Fast image restoration without boundary artifacts," *IEEE Transactions on Image Processing*, vol. 14, no. 10, pp. 1448–1453, 2005.
- [48] Y. Wang, J. Yu, and J. Zhang, "Zero-shot image restoration using denoising diffusion null-space model," in *Proc. Int. Conf. Learn. Represent. (ICLR)*, 2023.

- [49] H. Yang, J. Liu, R. Dian, and S. Li, "Correlation-aware learning self-guiding network for hyperspectral fusion imaging," *Information Fusion*, vol. 133, p. 104263, 2026.
- [50] Q. Ma, J. Jiang, X. Liu, and J. Ma, "Reciprocal transformer for hyperspectral and multispectral image fusion," *Information Fusion*, vol. 104, p. 102148, Apr. 2024.
- [51] DigitalGlobe, "Spectral response for DigitalGlobe Earth imaging instruments," Technical note, 2014.
- [52] G. Dial, H. Bowen, F. Gerlach, J. Grodecki, and R. Oleszczuk, "IKONOS satellite, imagery, and products," *Remote Sensing of Environment*, vol. 88, no. 1-2, pp. 23–36, Nov. 2003.
- [53] P. Micikevicius, S. Narang, J. Alben, G. Diamos, E. Elsen, D. Garcia, B. Ginsburg, M. Houston, O. Kuchaiev, G. Venkatesh, and H. Wu, "Mixed precision training," in *Proc. Int. Conf. Learn. Represent. (ICLR)*, 2018.
- [54] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional networks for biomedical image segmentation," in *Proc. Med. Image Comput. Comput.-Assist. Interv. (MICCAI)*, ser. Lecture Notes in Computer Science, vol. 9351, 2015, pp. 234–241.
- [55] Z. Wang, A. Bovik, H. Sheikh, and E. Simoncelli, "Image quality assessment: from error visibility to structural similarity," *IEEE Transactions on Image Processing*, vol. 13, no. 4, pp. 600–612, Apr. 2004.
- [56] B. T. Polyak and A. B. Juditsky, "Acceleration of stochastic approximation by averaging," *SIAM Journal on Control and Optimization*, vol. 30, no. 4, pp. 838–855, 1992.
- [57] B. Lim, S. Son, H. Kim, S. Nah, and K. M. Lee, "Enhanced deep residual networks for single image super-resolution," in *2017 IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2017, pp. 1132–1140.
- [58] N. Yokoya and A. Iwasaki, "Airborne hyperspectral data over Chikusei," Space Application Laboratory, University of Tokyo, Tokyo, Japan, Tech. Rep. SAL-2016-05-27, May 2016.
- [59] P. Gamba, "A collection of data for urban area characterization," in *Proc. IEEE Int. Geosci. Remote Sens. Symp. (IGARSS)*, vol. 1, 2004, pp. 69–72.
- [60] L. Wald, T. Ranchin, and M. Mangolini, "Fusion of satellite images of different spatial resolutions: Assessing the quality of resulting images," *Photogrammetric Engineering and Remote Sensing*, vol. 63, no. 6, pp. 691–699, 1997.
- [61] F. Yasuma, T. Mitsunaga, D. Iso, and S. K. Nayar, "Generalized assorted pixel camera: Postcapture control of resolution, dynamic range, and spectrum," *IEEE Transactions on Image Processing*, vol. 19, no. 9, pp. 2241–2253, 2010.
- [62] A. Chakrabarti and T. Zickler, "Statistics of real-world hyperspectral images," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2011, pp. 193–200.
- [63] R. H. Yuhas, A. F. H. Goetz, and J. W. Boardman, "Discrimination among semi-arid landscape endmembers using the spectral angle mapper (SAM) algorithm," in *Summaries of the Third Annual JPL Airborne Geoscience Workshop*, vol. 1, 1992, pp. 147–149.
- [64] L. Wald, "Quality of high resolution synthesised images: Is there a simple criterion?" in *Proc. Third Conf. Fusion of Earth Data: Merging Point Measurements, Raster Maps and Remotely Sensed Images*, 2000, pp. 99–103.
- [65] A. Garzelli and F. Nencini, "Hypercomplex quality assessment of multi/hyperspectral images," *IEEE Geoscience and Remote Sensing Letters*, vol. 6, no. 4, pp. 662–665, Oct. 2009.
- [66] G. Vivone, L. Alparone, J. Chanussot, M. Dalla Mura, A. Garzelli, G. A. Licciardi, R. Restaino, and L. Wald, "A critical comparison among pansharpening algorithms," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 53, no. 5, pp. 2565–2586, May 2015.
- [67] L.-j. Deng, G. Vivone, M. E. Paoletti, G. Scarpa, J. He, Y. Zhang, J. Chanussot, and A. Plaza, "Machine learning in pansharpening: A benchmark, from shallow to deep networks," *IEEE Geoscience and Remote Sensing Magazine*, vol. 10, no. 3, pp. 279–315, 2022.
- [68] J. Zhou, D. L. Civco, and J. A. Silander, "A wavelet transform method to merge Landsat TM and SPOT panchromatic data," *International Journal of Remote Sensing*, vol. 19, no. 4, pp. 743–757, Jan. 1998.
- [69] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, A. Desmaison, A. Köpf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai, and S. Chintala, "PyTorch: An imperative style, high-performance deep learning library," in *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2019, pp. 8026–8037.
- [70] S. Li, Y. Zhao, R. Varma, O. Salpekar, P. Noordhuis, T. Li, A. Paszke, J. Smith, B. Vaughan, P. Damania, and S. Chintala, "PyTorch distributed: experiences on accelerating data parallel training," *Proceedings of the VLDB Endowment*, vol. 13, no. 12, pp. 3005–3018, Aug. 2020.
- [71] I. Loshchilov and F. Hutter, "Decoupled weight decay regularization," in *Proc. Int. Conf. Learn. Represent. (ICLR)*, 2019.
- [72] —, "SGDR: Stochastic gradient descent with warm restarts," in *Proc. Int. Conf. Learn. Represent. (ICLR)*, 2017.

TABLE XV
ALL METRICS OF THE EQUAL-BUDGET BENCHMARK (SINGLE PASS).

Method	PSNR	SSIM	SAM	ERGAS	CC	SCC	Q2n	RMSE
<i>Chikusei</i> $\times 4$								
Bicubic	33.63	0.7404	4.406	9.011	0.8408	0.7648	0.6914	321.81
GSA	43.20	0.9778	2.730	3.279	0.9905	0.9922	0.9596	107.69
SSR-NET	—	—	—	—	—	—	—	—
PSRT	53.99	0.9948	1.057	1.794	0.9942	0.9982	0.9887	30.57
DCT	53.18	0.9945	1.133	2.066	0.9930	0.9978	0.9871	33.58
MIMFormer	56.96	0.9964	0.801	1.488	0.9960	0.9989	0.9918	21.65
DHIF-Net	—	—	—	—	—	—	—	—
PUFormer	57.95	0.9969	0.731	1.344	0.9966	0.9991	0.9929	19.45
<i>Chikusei</i> $\times 8$								
Bicubic	—	—	—	—	—	—	—	—
GSA	—	—	—	—	—	—	—	—
SSR-NET	—	—	—	—	—	—	—	—
PSRT	—	—	—	—	—	—	—	—
DCT	—	—	—	—	—	—	—	—
MIMFormer	—	—	—	—	—	—	—	—
DHIF-Net	—	—	—	—	—	—	—	—
PUFormer	—	—	—	—	—	—	—	—
<i>Chikusei</i> $\times 16$								
Bicubic	—	—	—	—	—	—	—	—
GSA	—	—	—	—	—	—	—	—
SSR-NET	—	—	—	—	—	—	—	—
PSRT	—	—	—	—	—	—	—	—
DCT	—	—	—	—	—	—	—	—
MIMFormer	—	—	—	—	—	—	—	—
DHIF-Net	—	—	—	—	—	—	—	—
PUFormer	—	—	—	—	—	—	—	—
<i>Pavia Centre</i> $\times 4$								
Bicubic	25.52	0.6282	6.554	9.852	0.8174	0.7526	0.6806	436.32
GSA	34.78	0.9438	5.229	3.494	0.9958	0.9922	0.9276	157.93
SSR-NET	—	—	—	—	—	—	—	—
PSRT	47.87	0.9958	2.212	0.801	0.9990	0.9989	0.9716	32.60
DCT	46.94	0.9953	2.442	0.888	0.9988	0.9985	0.9712	36.26
MIMFormer	49.63	0.9967	1.982	0.674	0.9993	0.9993	0.9738	26.66
DHIF-Net	—	—	—	—	—	—	—	—
PUFormer	50.63	0.9970	1.897	0.621	0.9994	0.9995	0.9754	23.68
<i>Pavia Centre</i> $\times 8$								
Bicubic	—	—	—	—	—	—	—	—
GSA	—	—	—	—	—	—	—	—
SSR-NET	—	—	—	—	—	—	—	—
PSRT	—	—	—	—	—	—	—	—
DCT	—	—	—	—	—	—	—	—
MIMFormer	—	—	—	—	—	—	—	—
DHIF-Net	—	—	—	—	—	—	—	—
PUFormer	—	—	—	—	—	—	—	—
<i>Pavia Centre</i> $\times 16$								
Bicubic	—	—	—	—	—	—	—	—
GSA	—	—	—	—	—	—	—	—
SSR-NET	—	—	—	—	—	—	—	—
PSRT	—	—	—	—	—	—	—	—
DCT	—	—	—	—	—	—	—	—
MIMFormer	—	—	—	—	—	—	—	—
DHIF-Net	—	—	—	—	—	—	—	—
PUFormer	—	—	—	—	—	—	—	—
<i>CAVE</i> $\times 4$								
Bicubic	—	—	—	—	—	—	—	—
GSA	—	—	—	—	—	—	—	—
SSR-NET	—	—	—	—	—	—	—	—
PSRT	—	—	—	—	—	—	—	—
DCT	—	—	—	—	—	—	—	—
MIMFormer	—	—	—	—	—	—	—	—
DHIF-Net	—	—	—	—	—	—	—	—
PUFormer	—	—	—	—	—	—	—	—
<i>CAVE</i> $\times 8$								
Bicubic	—	—	—	—	—	—	—	—
GSA	—	—	—	—	—	—	—	—
SSR-NET	—	—	—	—	—	—	—	—
PSRT	—	—	—	—	—	—	—	—
DCT	—	—	—	—	—	—	—	—
MIMFormer	—	—	—	—	—	—	—	—
DHIF-Net	—	—	—	—	—	—	—	—
PUFormer	—	—	—	—	—	—	—	—
<i>Harvard</i> $\times 4$								
Bicubic	—	—	—	—	—	—	—	—
GSA	—	—	—	—	—	—	—	—
SSR-NET	—	—	—	—	—	—	—	—
PSRT	—	—	—	—	—	—	—	—
DCT	—	—	—	—	—	—	—	—
MIMFormer	—	—	—	—	—	—	—	—
DHIF-Net	—	—	—	—	—	—	—	—
PUFormer	—	—	—	—	—	—	—	—
<i>Harvard</i> $\times 8$								
Bicubic	—	—	—	—	—	—	—	—
GSA	—	—	—	—	—	—	—	—
SSR-NET	—	—	—	—	—	—	—	—
PSRT	—	—	—	—	—	—	—	—
DCT	—	—	—	—	—	—	—	—